

SCC.211 Operating Systems

Lecture 3 – Programs, Processes, Processors, and OS

Amit Chopra

School of Computing & Communications, Lancaster University, UK

amit.chopra@lancaster.ac.uk

The three central concepts in concurrent programming

Processor

Hardware device that executes machine instructions



Program

Instruction sequence defining potential execution path

Passive description of what you would like to happen

Stored on disk/secondary memory



Process

Active system entity executing associate program(s)

or; Program *in execution* on a processor

Resides in primary memory, removed on reboot

(We'll assume process, thread, and task are equivalent for now)

Image Name	CPU
AcroRd32.exe *32	00
explorer.exe	00
opera.exe *32	00
opera.exe *32	00
opera.exe *32	00

Why is a process not a program?


1) Program may be executed by multiple processes at the same time

- Wordpad opened twice as separate processes of the same program
- 1 program, multiple processes



2) Process might run one program, and then another

- gcc pre-processor is a different program (executable file) from the compilers syntax analyser and code generator
- The pre-processor, syntax and analyser, and code generator are executed one after the other by the same process
- 1 process, multiple programs




Processes abstract over...

1) Processors

- Can have multiple processes regardless number of processors

2) “Disturbed” sequential execution

- Interrupt processing: Handling external events (keypresses, clock ticks)
- **Context switch**: Execution jumps to another part of memory



Need to store context information for a switched-out thread so can later resume exactly where it left off

A thread's context (minimally) comprises

- **Program Counter (PC):** Instruction in the program (or function) that the thread will execute next
- **Stack Pointer (SP):** Address at the top of the thread's call stack. The call stack remembers the sequence of function calls the thread is making as it executes program (each thread needs own stack)

Context may also include other info

(Three-step context-switching sequence)

1

De-schedule currently-running thread

- Save PC and SP CPU registers of current running thread
- Required so that thread resumes execution exactly where left off

2

Scheduler selects 'best' ready thread to run next

- Time-slicing, priority, starvation
- Hardware architecture, etc.

3

Resume newly-selected thread


- Restore register contents back to PC & SP registers
- Thread resumes where it last left off (PC is loaded last)

- Processes are always concurrent
- ...but are not always parallel

Parallel

Multiple ($n > 1$) processes executing simultaneously.


All executing at a given instant.



Concurrent

$n > 1$ processes are underway simultaneously.

< n may execute at given instant



At each context switch, a scheduler decides which process is executed next

Which processes finishes first?

Process 1


```
while(i < 10)
{
    i = i + 1;
    println(i);
}
println("Process 1 completed.");
```

Process 2


```
while(i < 10)
{
    i = i + 1;
    println(i);
}
println("Process 2 completed.");
```


When do we not have to worry about concurrency?

1) No shared data or communication




2) Read only data (constant)






When should we worry about concurrency?

Threads use a **shared resource** without **synchronization**



One or more threads **modify** the **shared resource**




Commoditization

Massive improvement in technology



Multi-core processors: Each core can run a thread concurrently with other cores

Microprocessor Transistor Counts 1971-2011 & Moore's Law



Number of transistors in dense integrated circuits doubles every 2 years

...Not for much longer!

Multi-threaded Operating Systems (OS) and Virtual Machines (VM)



Windows

IEEE POSIX

Multi-threaded middleware



Multi-threaded Middleware

Multi-threaded OS & VM

Multi-core Processors

Commoditization


Knowing concurrency is a **requirement** for efficient multi-thread and multi-core systems

Limit to Performance: Amdahl's law

Program speedup by adding more processors is limited by the serial part of the program

e.g. If 95% of program (its runtime) can be parallelized, theoretical maximum speedup is x20

$1/(1-p)$ where p is proportion of the program that can be parallelized.



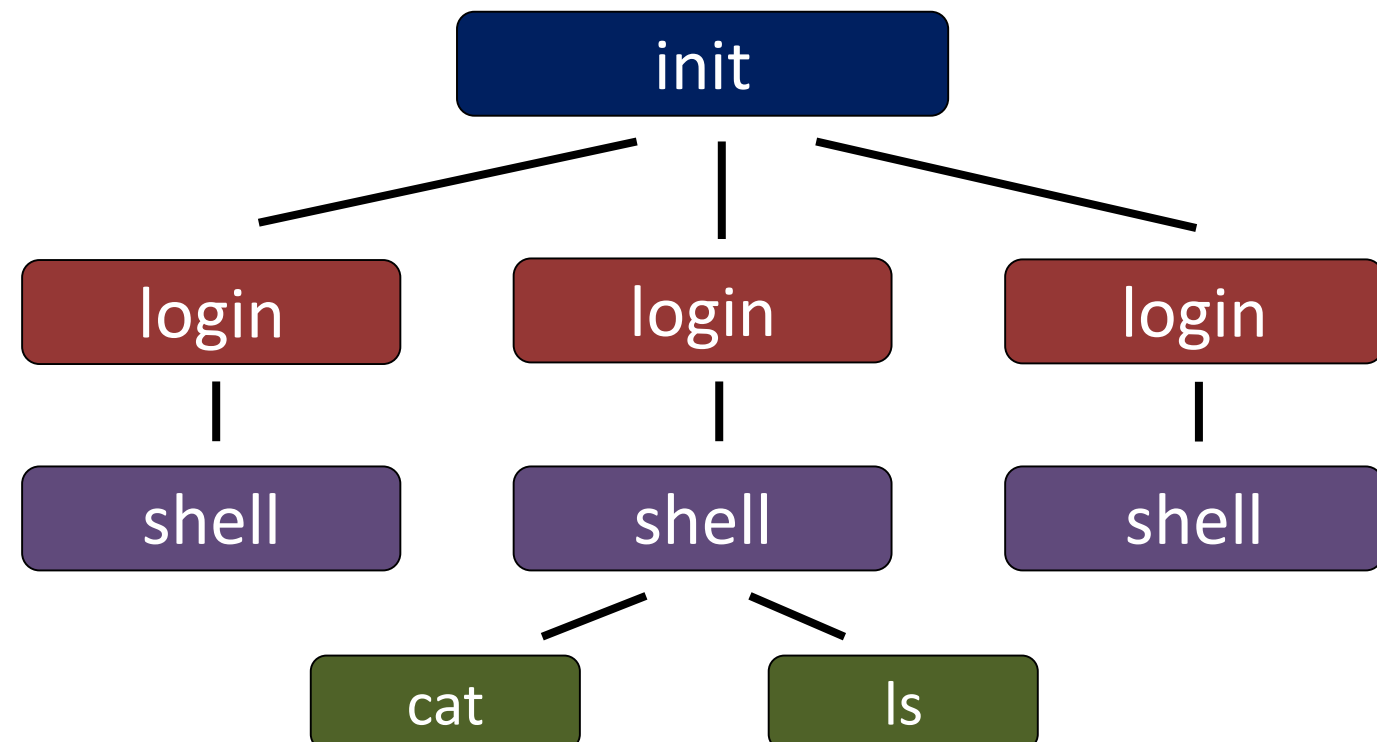
- **Accidental complexity**
 - Low-level APIs
 - Tedious, error-prone, time consuming
 - Non-portable
- **Limited debugging tools**
 - Actual behaviour vs. debug environment
 - Lack of tools to identify and rectify *race conditions*



Some System Calls for Unix Processes

fork()	Create a 'child' process
wait()	Block until child process dies
execve()	Make a calling process run a new program
kill()	Send a 'signal' to a process
signal()	Set a function to be called on arrival of 'signal'
pause()	Make calling process block until a 'signal' arrives
exit()	Commit suicide
getpid()	Get a process identifier of calling process

- Every Unix process (except one!) has a parent
- Processes may create multiple children (via `fork()`)
- Example: The (traditional) Unix boot procedure
 - Single parentless process runs (`init`)
 - `init` reads a file which lists all connected terminals
 - `init` forks a `login` process for each terminal
 - If a `login` process validates a user, it forks a `shell` process...

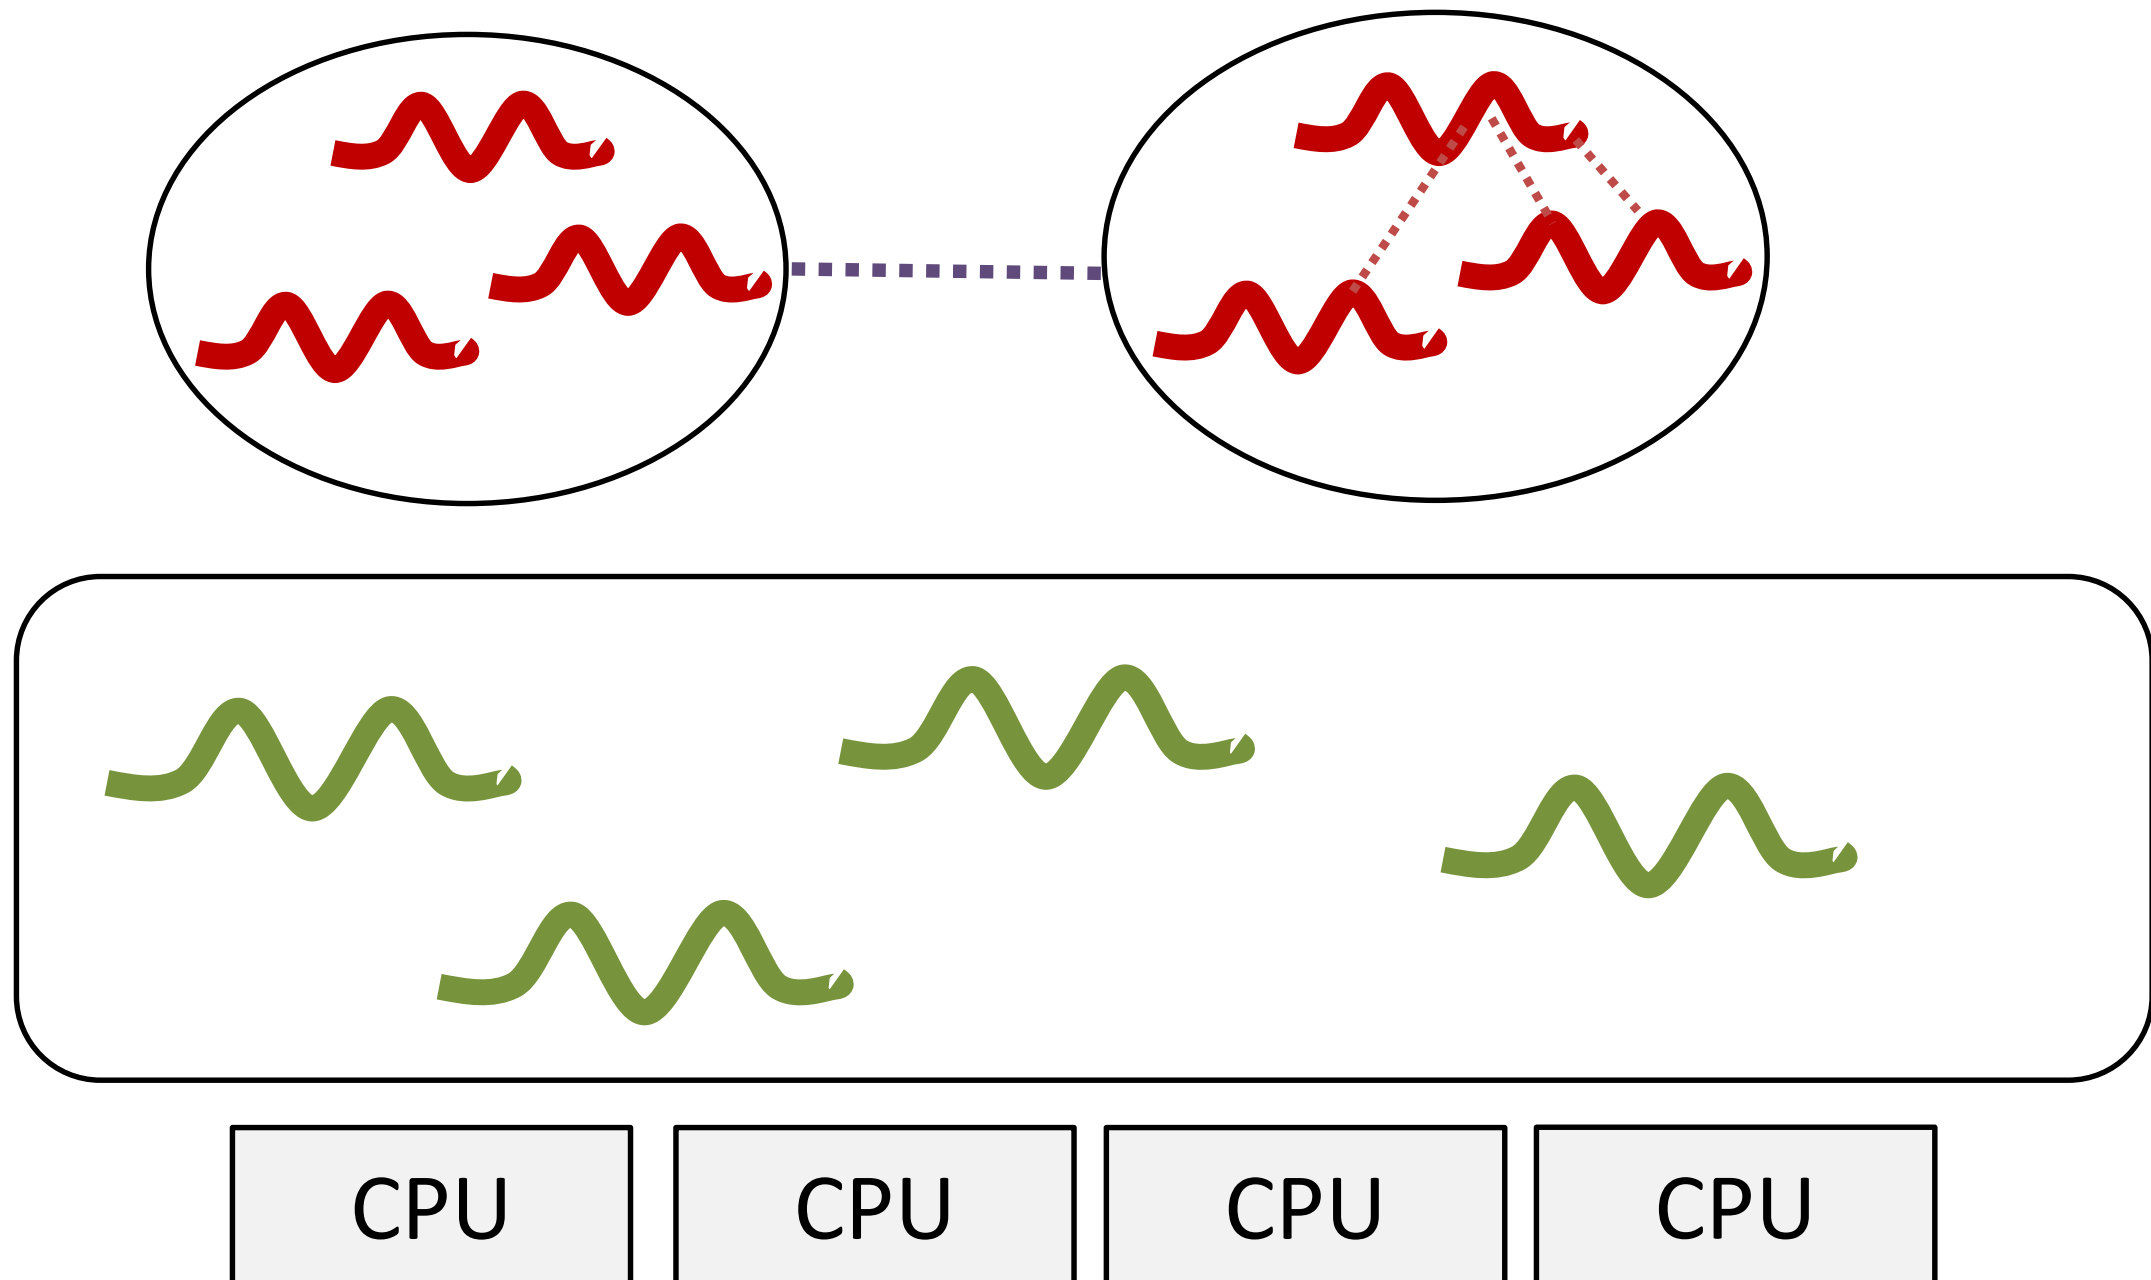


Unix processes combine a thread of execution with a ‘private virtual address space’

- Each process is like a ‘virtual machine’: CPU and memory combined in single abstraction
- User processes are protected from another
 - Good to avoid corruption and security problems
 - Cannot therefore straightforwardly share memory
 - Processes are heavyweight and costly to switch - as well as thread context (PC, SP, flags etc)

Threads share a common address space

- Can easily cooperate on common tasks
- Lighter weight switch; support fine-grained concurrency



User Threads

- Supported in user-level library
- Kernel only knows about processes
- Scheduled by a per-process, user-level scheduler

Kernel threads

- Supported by OS
 - Scheduled by OS scheduler
- Threads are preemptible by the OS
 - Potentially more concurrent, especially on multicore
 - More predictability useful for real-time applications
 - OS itself is multithreaded
- Supported in almost all modern OS

Threads in language runtime environments (Java)

- Available only in particular language environments
- Might be implemented either as user or kernel threads

User thread Advantages

- No OS support required
- Cheap context switch
 - Typically **order of magnitude faster** than kernel thread switch
 - (Which itself is an **order of magnitude faster** than process switch...)
 - Easy to offer per-application scheduling policies

User thread disadvantages

- If thread makes a system call that blocks, all other threads in that process also block (all reduce to the one underlying process)
- User threads in same process can't execute on separate CPUs on multicore – as the OS is not user-thread aware

User threads most useful for structuring **conceptually** concurrent applications
(GUIs, RPC servers, *some* web servers)

Each Java App. Launches a JVM

Inside JVM

Java Stack (created per thread)

Contains local variables, parameters and return values

Heap (one per JVM)

Contains objects, arrays (which are objects!)

- **Concurrency occurs in the physical world**
 - However computers need to be told explicitly what to do...
- **Programs, processes and processors** are different things
- Concurrency **does not** equal parallelism
- Concurrent programming is programming with >1 processes
- Threads can be implemented in user library or supported by OS
 - If library, then not visible to OS