

SCC.211 Operating Systems

Live Session 4

Dr Andrew Scott
a.scott@lancaster.ac.uk

Overview

- Topic 3: Memory Allocation (recap)
- Topic 4: Input and Output (introduction)
 - Device drivers
 - How can we create a device driver interface?
 - Disk storage

3. Memory Allocation

Recap

Videos

- Fixed and Variable length allocation
 - Internal vs. External fragmentation
- Buddy allocation
 - Allows allocation of contiguous memory blocks
 - Can be important for use by devices, sometimes in dedicated memory zone
 - Must track buddies; allocations cannot span buddies
- Slab allocation
 - Pre-allocate common sized memory areas, also avoids *allocate...free...allocate* cycle
 - Commonly used in kernel for frequently used structures

4. Input and Output (I/O)

Introduction

Includes diagrams from Silberschatz 10th ed.

Videos

- Overview of I/O mechanisms
 - How we attach devices and communicate with them
- Interrupt systems
 - How we work around unknown, and often long, duration of I/O operations
- Device drivers
 - Framework for device manufacturers to provide control software for devices

Device Drivers in Different OS Types

- Devices controlled by device drivers
 - In monolithic OSs
 - Tightly bound into OS *...failure will likely bring down whole system*
 - In modular OSs
 - Clean 'plug and play' interface *...more reliable and no need to reconfigure OS*
 - In microkernel OSs
 - Appear as unprivileged 'user' processes *...secure and easily restarted on failure*

An aside...

How can we create a device driver interface...

...or how to 'hook' arbitrary functions in C

C Function Pointers: *putting it together*

```
#include <stdio.h>

void input ( int * var ) {
    printf( "Enter starting value... " );
    scanf ( "%d", var );
    printf( "You input: %d\n", *var );
}

void addition ( int * var ) { (*var)++; }
void output ( int * var ) { printf( "After addition: %d\n", *var ); }

void ( * sequence[ ] )( int * ) = { input, addition, output };

int main ( int argc, char ** argv ) {
    int variable;

    for ( int i = 0; i < sizeof( sequence ) / sizeof( void * ); i++ )
        ( * sequence[ i ] )( &variable );
}
```

Enter starting value... 100
You input: 100
After addition: 101

Storage: Disk systems

- Overview
 - Structure
 - Two approaches for improving
 - Performance
 - Resilience

Hard Disk Structure

- Significant access latency
 - Seek time to move head in/ out
 - Rotational latency
- Sector (native block) size
 - Traditionally 512 bytes
 - Now typically 4K




Figure 11.1 HDD moving-head disk mechanism.

Hard Disk Structure II




Figure 11.1 HDD moving-head disk mechanism.

Hard Disk Structure II

- Group tracks into zones of n tracks
 - As move from spindle out, each zone will have more sectors per track
- A run of sectors is known as a cluster



Disk Addressing

- Traditionally used simple { Cylinder, Head, Sector } or CHS tuple
- Now disk controllers typically offer Logical Block Address (LBA)
 - Hides what can now be a complex internal zone structure/ geometry
 - Also allows controller to hide bad sectors
 - Controller may maintain and remap spare sectors in case any become unusable

Redundant Array of Independent Disks (RAID)

- Aims
 - Improve throughput with parallelism
 - Block or bit striping data across multiple disks
 - Improve Mean Time Between Failure (MTBF), recover from failed disk(s)
 - Mirroring ('backup' / 'shadow' drives)
 - Error correcting codes

Common RAID Levels


...there are many

- RAID 0: Block based **striping** across n disks to improve throughput
 - RAID 1: **Mirror** each drive onto a shadow drive to improve resilience
 - RAID 4: Block based **striping** across n disks + dedicated parity disk
 - RAID 5: As 4, but cycles **parity** through drives so one (parity) disk not a bottleneck
 - RAID 6: As 5, but adds another error correcting code to protect against multiple failures
 - RAID 0+1 and 1+0: combination of levels 0 and 1
 - Stripe then mirror, or mirror then stripe
- RAID 2 like 0 but bit based, RAID 3 like 4 but bit based (typically require complex drive synchronisation)

RAID Levels

- C: Mirror copy
 - Recover from single failure
- P: Parity
 - Recover from single failure
- Q: Second error code
 - Used in addition to parity enabling recovery from double failure

Figure 11.15 RAID levels.



Problems with RAID

- Mean Time To Repair can be long with some RAID levels
 - Significant time to rebuild lost drive
- Controller faults can cause whole array to fail
- Power failure can cause multiple (un-correctable) block writes to fail
 - A fast Solid-State Drive (SSD) based transaction cache can help

