

SCC.211 Operating Systems

Session 7

Dr Andrew Scott

a.scott@lancaster.ac.uk

1

Overview

- Summary/ recap of:
 - Topic 5: Multi-process Systems
 - Real-Time Systems and Real-Time Scheduling
 - Topic 6: File-systems
- Brief Introduction to
 - Topic 7: Memory Protection

2

5. Multi-process Systems

Summary/ Recap

3

Recap

- Types of multi-process system
- Difference between dispatchers and schedulers
 - Scheduler makes choice, dispatcher actions that choice
- Schedulers
 - Seen examples and comparison/ common metrics for
 - First Come First Served
 - Round Robin

4

Real-Time Systems

- Real-time process
 - Process that delivers results of processing in a given time-span
- Real-time system
 - System in which correctness of computation depends not only on obtaining the right result, but also upon providing the result on time
- Deadline
 - Deadline represents latest acceptable time for
 - Processing result to be considered correct, or...
 - Presentation of result to be valid

5

Real-Time Processes

- Soft Real-Time
 - Deadlines may be broken
e.g. High-end, professional multimedia
- Hard Real-Time
 - Deadlines critical
e.g. Embedded/ Control Systems
 - May require safety conformance certification (proof it will always meet deadlines)

6

Real-Time Scheduling

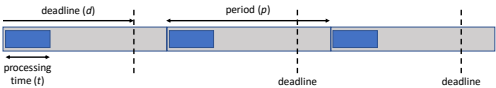
- Driven by events/ deadlines
 - Periodic events ...predictable
 - Aperiodic events ...unpredictable -- beware!
- Must ensure work can always be scheduled even in worst possible case
 - Never take on too much work ...requires **admission-control**
 - Utilisation of resources (including CPU) must never be greater than 100%
 - This must also consider context switching overheads

7

Periodic Tasks

- Rate = $1 / \text{period}$
- Need to check
 - Processing time for process must be < process deadline
 - Total processing time per cycle must be < period

...Much harder with multiple tasks, periods, and deadlines



8

Pre-emptive Real-Time Schedulers

- Rate Monotonic (RM)
 - Fixed, repeating schedule for worst possible case ...pre-emptive version of last slide
 - Priority inversely proportional to required work period ...more frequent = higher priority
 - Assume processes have same amount of work (CPU processing/ burst time) each cycle
- Earliest Deadline First (EDF)
 - Dynamic scheduling scheme
 - Each process has sequence of deadlines ...schedule process with closest deadline

Recall: we saw EDF with disk scheduling

9

6. File-Systems

Summary/ Recap

10

Overview

- Approaches to storing files
- Common file-systems
 - FAT, traditional Unix, e.g., ext2
- Journaling and Locking
- New approaches

11

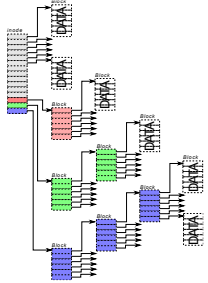
FAT Schemes

- Simple filesystems
- Common on small disks/ consumer products
- Limited capacity
 - Need to scan full sector/ block list to find any point in file
 - List held in File Allocation Table (typically multiple copies for resilience)
- To be efficient, whole FAT must be in-memory

12

Unix index-node (inode)

- Combined index scheme
 - Random access – can jump to any point in file
 - Index levels/ depth vary with file size
 - Supports very large files
 - Efficient for small files
- Supports sparse files
 - No need to store empty (all zero) blocks
 - Data blocks or index blocks



13

Creating a Sparse File

```
const char * message = "Hello World\n";

int
main( ) {
    FILE * stream = fopen( "testfile.txt", "w" );
    if( stream == NULL ) { /* Handle error */ }

    fwrite( message, sizeof( char ), strlen( message ), stream );

    fseek( stream, 1024 * 1024 * 1024, SEEK_SET ); // Make hole

    fwrite( message, sizeof( char ), strlen( message ), stream );


    fclose( stream );
}
```

14

Program Output

```
$ gcc hole.c -o hole
$ ./hole
$ ls -l
total 32
-rwxrwxr-x 1 acs acs 17440 Oct 13 21:35 hole
-rw-rw-r-- 1 acs acs 453 Oct 13 21:34 hole.c
-rw-rw-r-- 1 acs acs 1073741836 Oct 13 21:35 testfile.txt
$
```

~ 1GB in size



15

Program Output II

```
$ gcc hole.c -o hole
$ ./hole
$ ls -l
total 32
-rwxrwxr-x 1 acs acs 17920 Oct 13 21:35 hole
-rwxrwxr-x 1 acs acs 453 Oct 13 21:34 hole.c
-rwxrwxr-x 1 acs acs 1073741836 Oct 13 21:35 testfile.txt

$ du -k *
20 hole
4 hole.c
$ testfile.txt
$ od testfile.txt
0000000 H e l l o   W o r l d \n \0
0000020 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0 \0
*
10000000000 H e l l o   W o r l d \n
10000000014
```

~ 1GB in size

Only 8KB on disk
(2 x 4K blocks)

Notice the hole

16

Journaling File Systems

- Examples: NTFS, ext3, ...
- File-system updates written as *transactions* to journal/ log
 - Transactions periodically flushed to disk
 - Entries only removed from log when confirmed written/ flushed
- Fast as decouples file writes from disk head movement
 - Flush operation can do in-order sector (block) writes
 - Transaction log could be on faster disk/ media
- Resilient as log can be 'replayed' in event of system failure
 - Transactions must be **atomic** and **idempotent**

Atomic:

Idempotent:

All or no subparts of transaction must be completed

Can be repeated any number of times and still give same result

17

Journaling

- (Circular) Transaction log of filesystem updates
 - Can be replayed in event of system failure
 - If committed all updates must done, else none
 - Depending on filesystem, log can be held on same disk or fast SSD

Start (Metadata) Update Update Update Commit

```
graph LR
    Computer[Computer] <--> Cache[Cache (SSD)]
    Cache <--> Disk[Hard Disk]
```

18

POSIX File Locking: `fcntl()`, `lockf()`

- Often only way to coordinate unrelated applications
- Shared Lock
 - Any number of shared locks possible
 - Cannot include any data byte under an exclusive lock
- Exclusive Lock
 - Fails if any requested byte subject to existing lock
- API allows blocking and non-blocking lock requests
 - Blocking lock request hangs until bytes available (no other locks)

Note: a process can only have one type of lock on a file at any point

Early systems often locked at file descriptor (whole file) level, see (BSD style) `flock()`
-- now mainly used for process/ thread synchronisation, but note that threads must explicitly `open()` file to get new descriptor... i.e. not inherit via `fork()`, or via `dup()`.

19

POSIX File Locking

The diagram shows a horizontal array of 16 boxes representing bytes, indexed from 0 to 15. Bytes 4, 5, 6, and 7 are highlighted in orange. A blue bracket above these four boxes is labeled "Shared Lock A".

20

POSIX File Locking

The diagram shows a horizontal array of 16 boxes representing bytes, indexed from 0 to 15. Bytes 4, 5, 6, and 7 are highlighted in orange, and bytes 5, 6, 7, and 8 are highlighted in yellow. A blue bracket above bytes 4-7 is labeled "Shared Lock A", and a blue bracket below bytes 5-8 is labeled "Shared Lock B".

21

POSIX File Locking

22

POSIX File Locking

23

New filesystems: *extents and delayed writes*

- Portions of file stored as contiguous set of disk blocks
 - Makes file reads more efficient
 - Less meta-data (block pointers and indexes, etc.) to process
 - Less fragmentation
 - Less head movement
 - Delay (collate) writes, or applications can use *pre-allocation* with *fallocate()*
- File information contains list of extents:

Logical Block address (start of extent in file)
Number of contiguous disk blocks in extent
Block address of first disk block in extent

Used in ext4 and similar *Cluster Run* scheme used in NTFS

24

Other Improvements

- Fine grained timestamps
 - Granularity of 1 sec no longer acceptable... use *ns*
- Database technology in filesystems
 - Database style (H or B+ based) indexes for
 - Directories
 - Accessing tree of extents
 - Databases as a filesystem
 - Very powerful
 - Problematic due to semantics of/ expected result common operations
 - No longer simple tree of files and directories expected by users and applications

25

7. Memory Protection

Introduction

26

Memory Protection

- Basic approach
 - Memory Management Unit (MMU)
- Segmentation
 - Variable length scheme
- Paging
 - Fixed length scheme

27
