



SCC201: Databases

Week5: SQL and Advanced SQL Scripts

By
Uraz C turker

Week 1	Lecture 1 Lecture 2	Introduction to the module, Why do we need Databases? Entity Relationship Model Entity Relationship Model (ERM)
	Lab	NO LAB
Week 2	Lecture 1 Lecture 2	Relational Model (RM) ER to RM
	Lab	ER diagrams.
Week 3	Lecture 1 Lecture 2	Functional Dependencies 1st, 2nd, 3rd and Boyce-Cott Normal Forms
	Lab	ER to Relational Model.
Week 4	Lecture 1 Lecture 2	Relational Algebra Relational Algebra
	Lab	Functional dependencies and Normal forms
Week 5	Lecture 1 Lecture 2	SQL Scripts Advanced SQL Scripts
	Lab	Relational algebra
Week 6	Lecture 1 Lecture 2	JDBC Review
	Lab	Advanced SQL Scripting
Week 7	Lecture 1 Lecture 2	Physical Storage - record files Storage - secondary files
	Lab	Project
Week 8	Lecture 1 Lecture 2	Record Search - B-Trees Search - Hashing
	Lab	Project
Week 9	Lecture 1 Lecture 2	Concurrency - Transaction Processing (cont) Durability of Transactions and Crash Recovery
	Lab	Project
Week 10	Lecture 1 Lecture 2	Advanced SQL - schemas, views, access control Review
	Lab	Project

From you

What aspect of the module did you find most difficult?

3NF and BCNF

-

Getting head around whether to use section or projection

Currently; sorting the FDs using the Axioms

From you

What aspect of the module did you find most difficult?
3NF and BCNF

-
Getting head around wther to use section or projection
Currently; sorting the FDs using the Axioms

What can be improved to help with the module?

None

I don't know if you have any control over this but there is only one physical copy of the required reading in the library. Would be good to have more for those that prefer physical book over reading off a screen.

more examples to walk though or exercises to do for relational algebra

I think that it would be useful to release the worked solutions with the lab worksheets so that we can go through and check our answers in the lab while it is fresh and so that should we get stuck, we can check we are on the right track

From you

What aspect of the module did you find most difficult?
3NF and BCNF

-
Getting head around wther to use section or projection
Currently; sorting the FDs using the Axioms

What can be improved to help with the module?

None
I don't know if you have any control over this but there is only one physical copy of the required reading in the library. Would be good to have more for those that prefer physical book over reading off a screen.

more examples to walk though or exercises to do for relational algebra
I think that it would be useful to release the worked solutions with the lab worksheets so that we can go through and check our answers in the lab while it is fresh and so that should we get stuck, we can check we are on the right track

How was the support offered by module convenor?

Extra

Extra

Enough

None

From you

What aspect of the module did you find most difficult?
3NF and BCNF

-

Getting head around wther to use section or projection

Currently; sorting the FDs using the Axioms

What can be improved to help with the module?
None

I don't know if you have any control over this but there is only one physical copy of the required reading in the library. Would be good to have more for those that prefer physical book over reading off a screen.

more examples to walk though or exercises to do for relational algebra

I think that it would be useful to release the worked solutions with the lab worksheets so that we can go through and check our answers in the lab while it is fresh and so that should we get stuck, we can check we are on the right track

How was the support offered by module convenor?
Extra

Extra

Enough

None

What more support can be provided?
N/A

-

none Uraz is great love that he remembers names :)

n/a

From you

What aspect of the module did you find most difficult?
3NF and BCNF

-
Getting head around wther to use section or projection
Currently; sorting the FDs using the Axioms

How was the support offered by module convenor?
Extra
Extra
Enough
None

What can be improved to help with the module?
None

I don't know if you have any control over this but there is only one physical copy of the required reading in the library. Would be good to have more for those that prefer physical book over reading off a screen.

more examples to walk though or exercises to do for relational algebra

I think that it would be useful to release the worked solutions with the lab worksheets so that we can go through and check our answers in the lab while it is fresh and so that should we get stuck, we can check we are on the right track

What more support can be provided?
N/A

-
none Uraz is great love that he remembers names :)
n/a

Any other comments

N/A

-
Do we need to use semijoins and more as when reading the chapters said to for this week it goes into a lot more such as aggregate and grouping do we need to know this for the test in the summer and do we need to learn the definitions for each one aswell as could we be asked a question saying what does this mean

n/a

Resources



<https://mylab.lancaster.ac.uk>

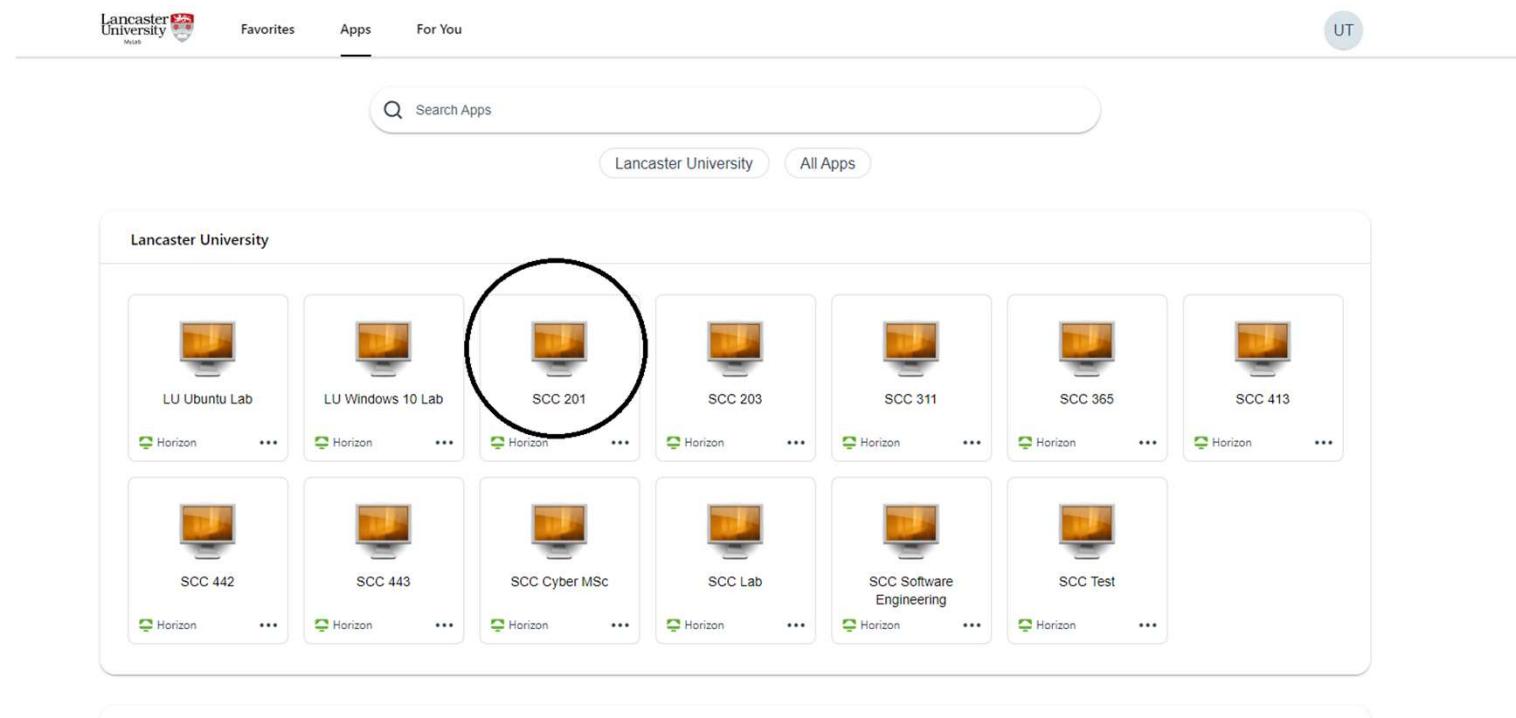
MySQL by Paul DuBois

Lancaster University  Favorites Apps For You UT

Search Apps Lancaster University All Apps

Lancaster University

LU Ubuntu Lab	LU Windows 10 Lab	SCC 201	SCC 203	SCC 311	SCC 365	SCC 413
Horizon ...	Horizon ...	Horizon ...	Horizon ...	Horizon ...	Horizon ...	Horizon ...
SCC 442	SCC 443	SCC Cyber MSc	SCC Lab	SCC Software Engineering	SCC Test	
Horizon ...	Horizon ...	Horizon ...	Horizon ...	Horizon ...	Horizon ...	





```
turker@vdi-scc201-018: ~
$ mysql
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 11
Server version: 8.0.35-0ubuntu0.22.04.1 (Ubuntu)

Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
mysql> ■
```

SCC 201 Databases

Learning outcomes

This week, you will learn the basics of SQL scripting, including notations and formulations.

You will gain hands-on experience in converting ERDs and RLs into SQL scripts.

You will learn to generate nested queries and aggregate operations using SQL.

These slides and SCC.201 are based on MySQL v8; using any other versions of SQL may lead you to lose marks in coursework and exams.

The SQL Query Language

- Developed by IBM (system R)
- Need for a standard since it is
- Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision, current)
 - SQL-99 (major extensions)

```
void myDisplay()
{
    game.display();

    // Define the reshape function
void myReshape(int width, int height)
{
    game.reshape(width, height);

    // Define the mouse click events
void myMouseClick(int button, int state, int x, int y)
{
    game.mouseClick(button, state, x, y);
```

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) > 1
```

Creating Relations in SQL

```
CREATE TABLE Students  
(sid TEXT,  
 name TEXT,  
 login TEXT,  
 age INTEGER,  
 gpa REAL);
```

Creates the Students' relation. Observe that each field's type (domain) is specified and enforced by the DBMS whenever tuples are added or modified.

CREATE Relation's Name(*Relational Schema Text in MySQL format*);

RS:

Students(sid TEXT, name TEXT, login TEXT, age TEXT, gpa REAL);

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

Built-in data types varies interpreter to another.

TEXT, VARCHAR(Length), REAL, INTEGER, and BLOB are the most common data types.

Some interpreters accepts INT and INTEGER, some accepts DATA and BOOLEAN, but others don't.

MySQL DATATYPES

DATE TYPE	SPEC	DATA TYPE	SPEC
CHAR	String (0 - 255)	INT	Integer (-2147483648 to 2147483647)
VARCHAR	String (0 - 255)	BIGINT	Integer (-9223372036854775808 to 9223372036854775807)
TINYTEXT	String (0 - 255)	FLOAT	Decimal (precise to 23 digits)
TEXT	String (0 - 65535)	DOUBLE	Decimal (24 to 53 digits)
BLOB	String (0 - 65535)	DECIMAL	"DOUBLE" stored as string
MEDIUMTEXT	String (0 - 16777215)	DATE	YYYY-MM-DD
MEDIUMBLOB	String (0 - 16777215)	DATETIME	YYYY-MM-DD HH:MM:SS
LONGTEXT	String (0 - 4294967295)	TIMESTAMP	YYYYMMDDHHMMSS
LONGBLOB	String (0 - 4294967295)	TIME	HH:MM:SS
TINYINT	Integer (-128 to 127)	ENUM	One of preset options
SMALLINT	Integer (-32768 to 32767)	SET	Selection of preset options
MEDIUMINT	Integer (-8388608 to 8388607)	BOOLEAN	TINYINT(1)

Creating Relations in SQL

As another example, the `Enrolled` table holds information about students' courses and grades.

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

Students

```
CREATE TABLE Enrolled  
  (sid TEXT,  
   cid TEXT,  
   grade TEXT);
```

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Enrolled

Recall Keys

id	name	surname	Phone number	gpa
53666	Cherry	Hardguy	123456	3.4
53688	Cherry	Goodgirl	234567	3.4

Recall Keys

id	name	surname	Phone number	gpa
53666	Cherry	Hardguy	123456	3.4
53688	Cherry	Goodgirl	234567	3.4

- Primary Key: A set of attributes that can uniquely identify tuples in a given relation.

- PK: {name,surname}



Recall Keys

id	name	surname	Phone number	gpa
53666	Cherry	Hardguy	123456	3.4
53688	Cherry	Goodgirl	234567	3.4

- Primary Key: A set of attributes that can uniquely identify tuples in a given relation.
 - PK: {name,surname}
- Candidate Key: A set of attributes that can be selected as a Primary Key.
 - CK:{id}{phone number}{name,surname}
-
-

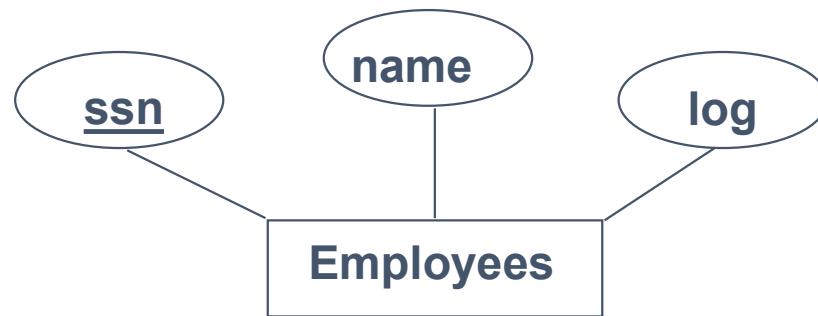
Recall Keys

id	name	surname	Phone number	gpa
53666	Cherry	Hardguy	123456	3.4
53688	Cherry	Goodgirl	234567	3.4

- Primary Key: A set of attributes that can uniquely identify tuples in a given relation.
 - PK: {name,surname}
- Candidate Key: A set of attributes that can be selected as a Primary Key.
 - CK:{id}{phone number}{name,surname}}
- Super Key: A set of attributes having the Primary Key.
 - SK:{name,surname,gpa}

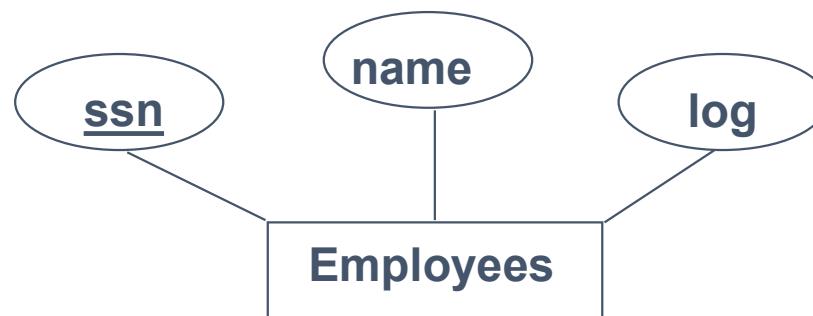
Logical DB Design: ER to Relational

Entity sets to tables.



Logical DB Design: ER to Relational

Entity sets to tables.

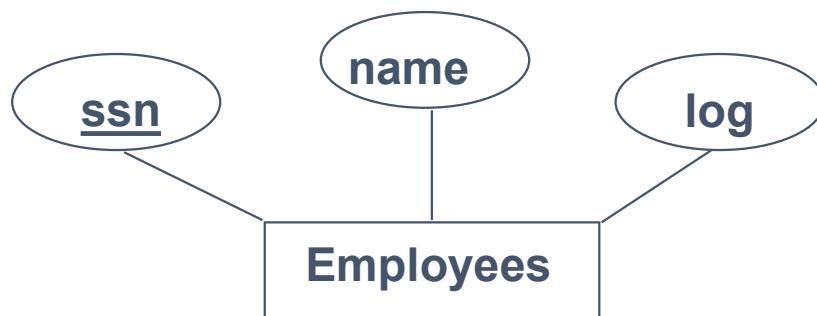


```
CREATE TABLE Employees  
  (ssn VARCHAR(11),  
   name VARCHAR(20),  
   log INTEGER,
```



Logical DB Design: ER to Relational

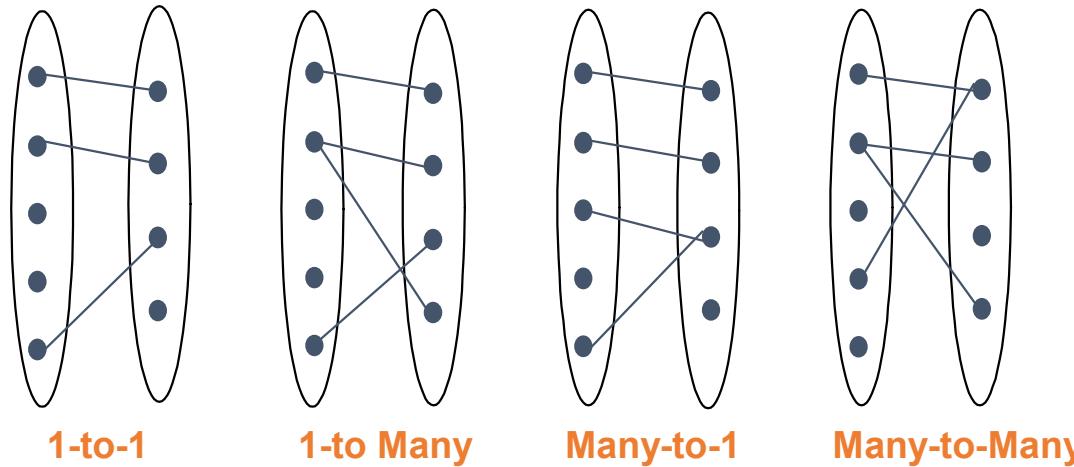
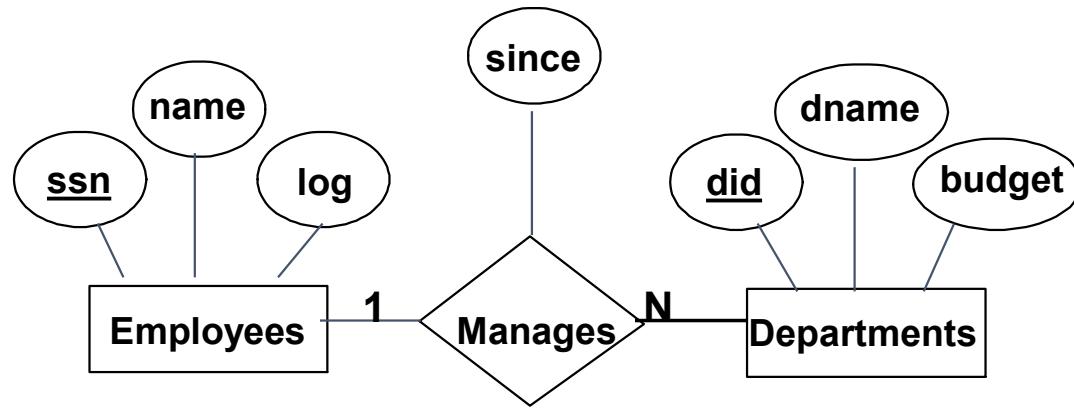
Entity sets to tables.



```
CREATE TABLE Employees
(ssn VARCHAR(11),
name VARCHAR(20),
log INTEGER,
PRIMARY KEY (ssn));
```

Review: participation and multiplicity Constraints

- Each dept has at most one manager, according to the *key constraint* on Manages.

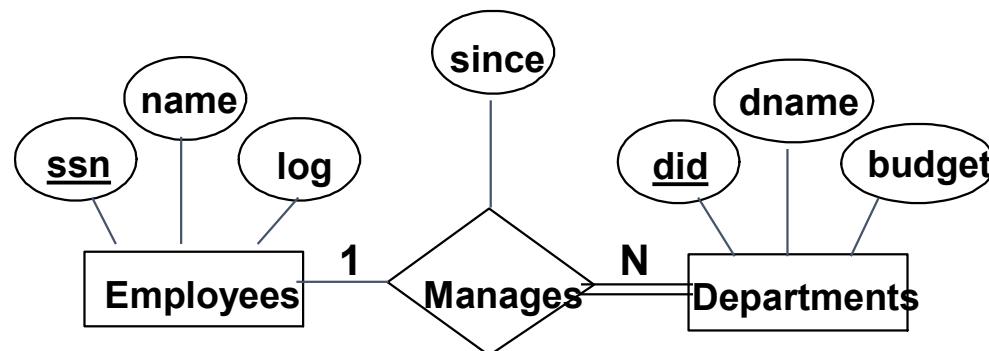


Translation to relational model?

Is this correct?



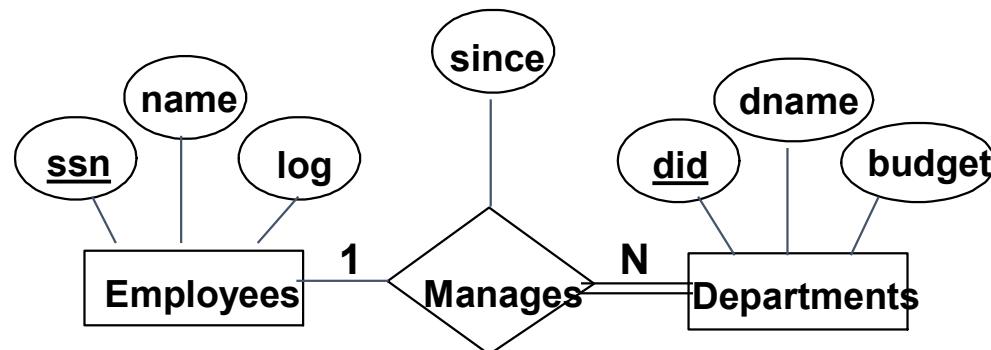
- “Employee can manage one department, a department must be managed by many employees”



Is this correct?



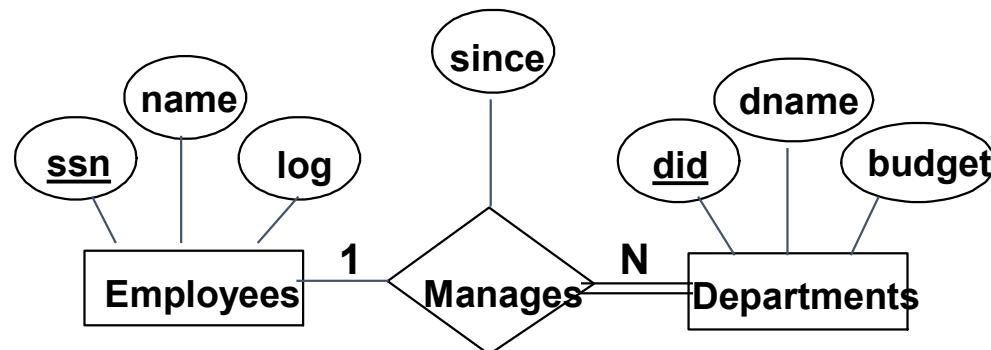
- “Employee must manage one department, a department may be managed by many employees”



Is this correct?

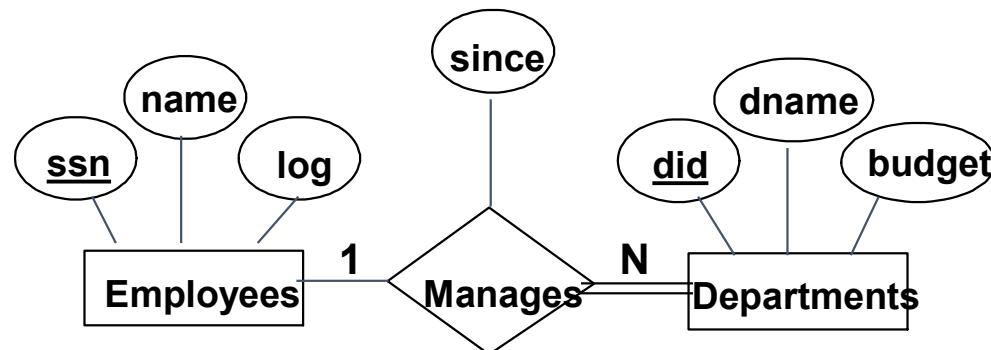


- “Employee can manage many departments, a department may be managed by many employees”



Is this correct?

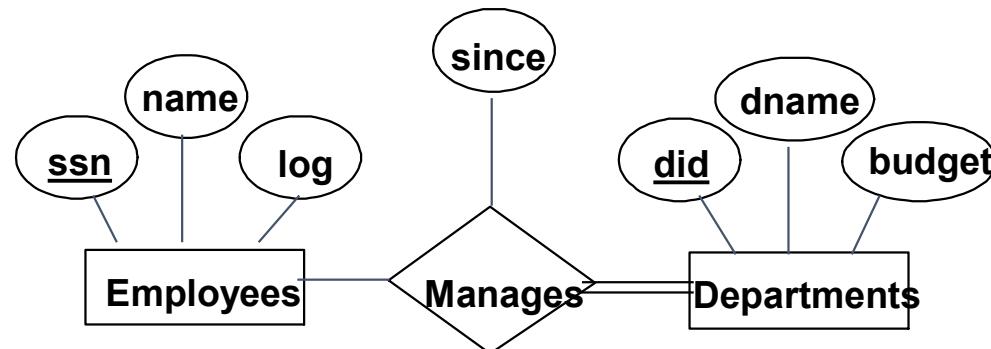
- “Employee can manage many departments, a department must be managed by one employee”



Is this correct?



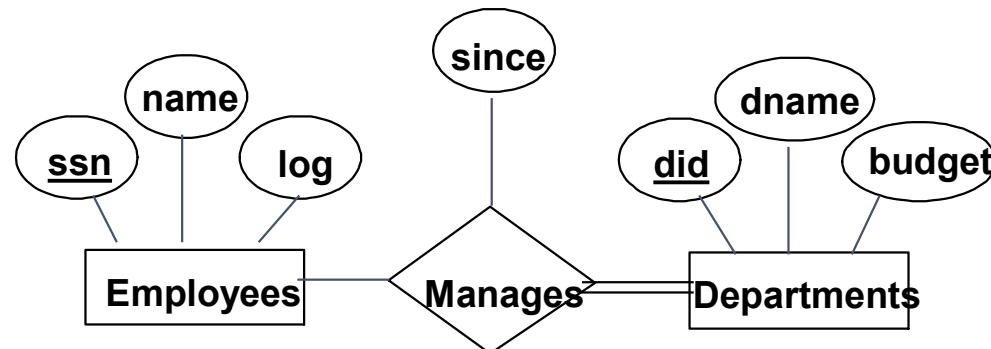
- “Employee must manage department(s) but department may be managed by employee(s).”



Is this correct?

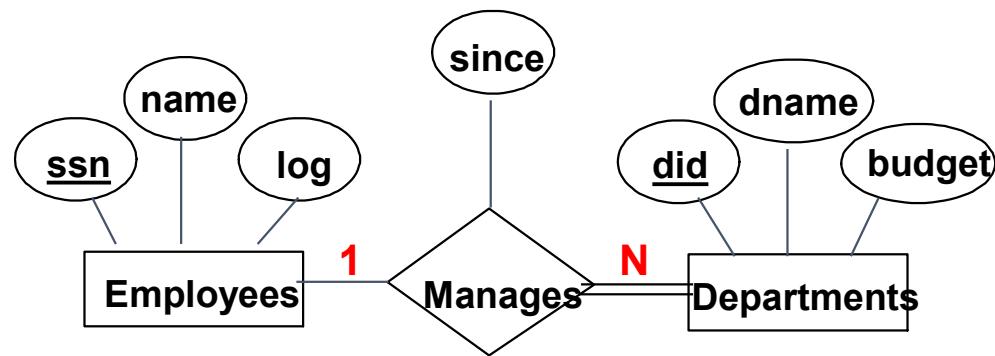


- “Employee can manage department(s) but department must be managed by employee(s).”



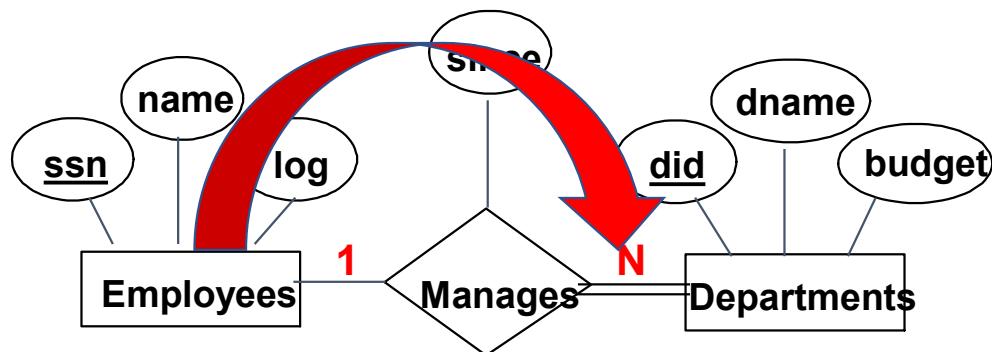
RECALL

- Multiplicity constraints are given in the **OPPOSITE** ends of the relation



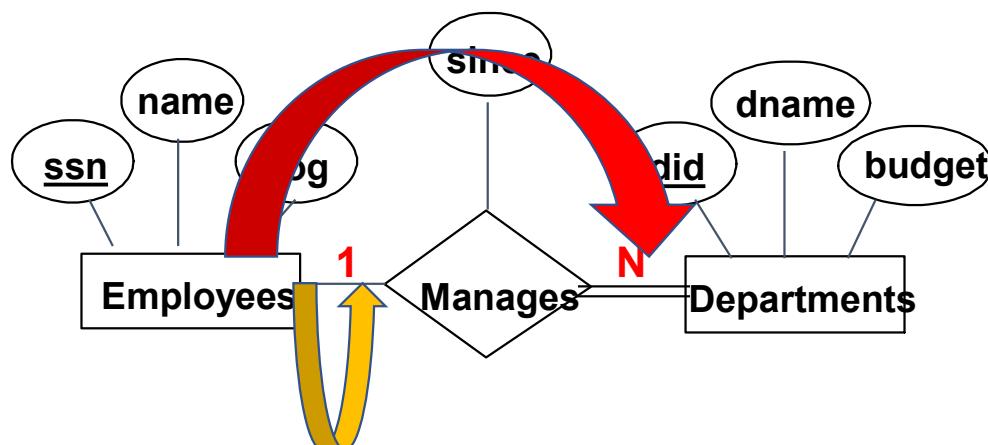
RECALL

- Multiplicity constraints are given in the **OPPOSITE** ends of the relation
- Employee manages **MANY** departments.



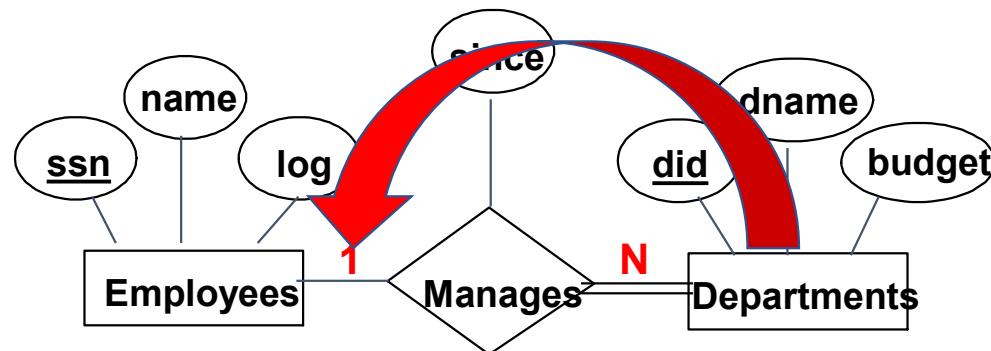
RECALL

- Participation constraints are given **BY** the entity set.
- Employee **MAY** manage **MANY** departments.



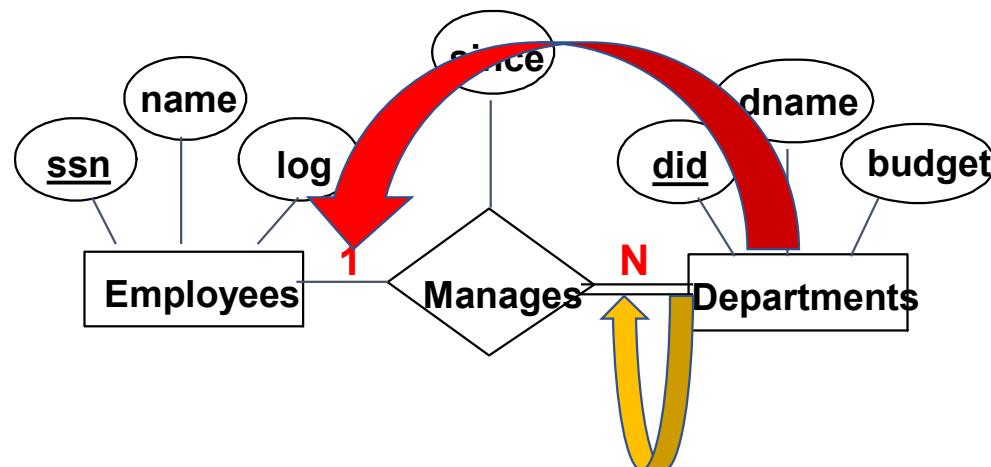
RECALL

- Multiplicity constraints are given in the **OPPOSITE** ends of the relation, Participation constraints are given **BY** the entity set.
- Employee may manage many departments.
- Department managed by **ONE** employee.



RECALL

- Multiplicity constraints are given in the **OPPOSITE** ends of the relation, Participation constraints are given **BY** the entity set.
- Employee may manage many departments.
- A department **must** be managed by **one** employee.



Foreign Keys in SQL

FOREIGN KEY (*Attribute Name*) REFERENCES *Relation Name* (*Attribute Name*)

```
CREATE TABLE Enrolled  
  (sid VARCHAR(20), cid VARCHAR(20), grade VARCHAR(2),  
   PRIMARY KEY (sid,cid),  
   FOREIGN KEY (sid) REFERENCES Students (sid));
```

Enrolled

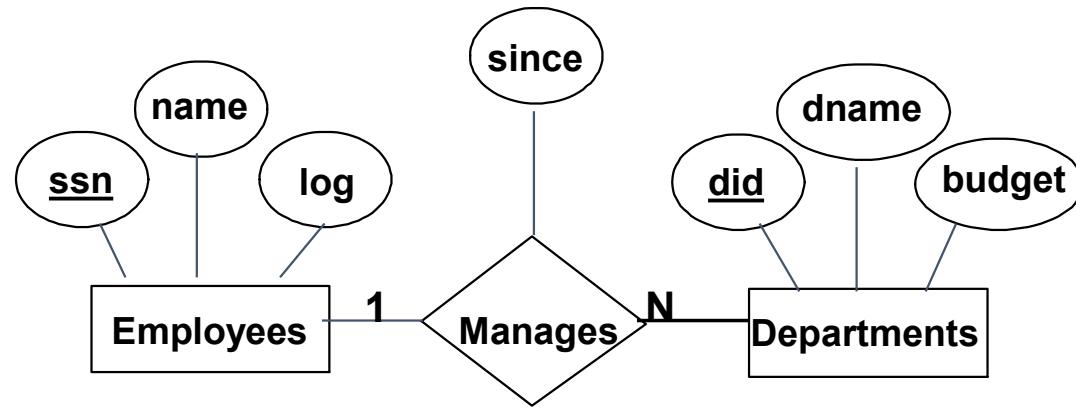
<u>sid</u>	<u>cid</u>	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

Students

<u>sid</u>	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Review: Key Constraints

- Each dept has at most one manager, according to the *key constraint* on Manages.



```

CREATE TABLE Employees
(ssn VARCHAR(11),
name VARCHAR(20),
log INTEGER,
PRIMARY KEY (ssn));
  
```

```

CREATE TABLE Departments
(did INTEGER,
dname VARCHAR(20),
budget REAL,
PRIMARY KEY (did));
  
```

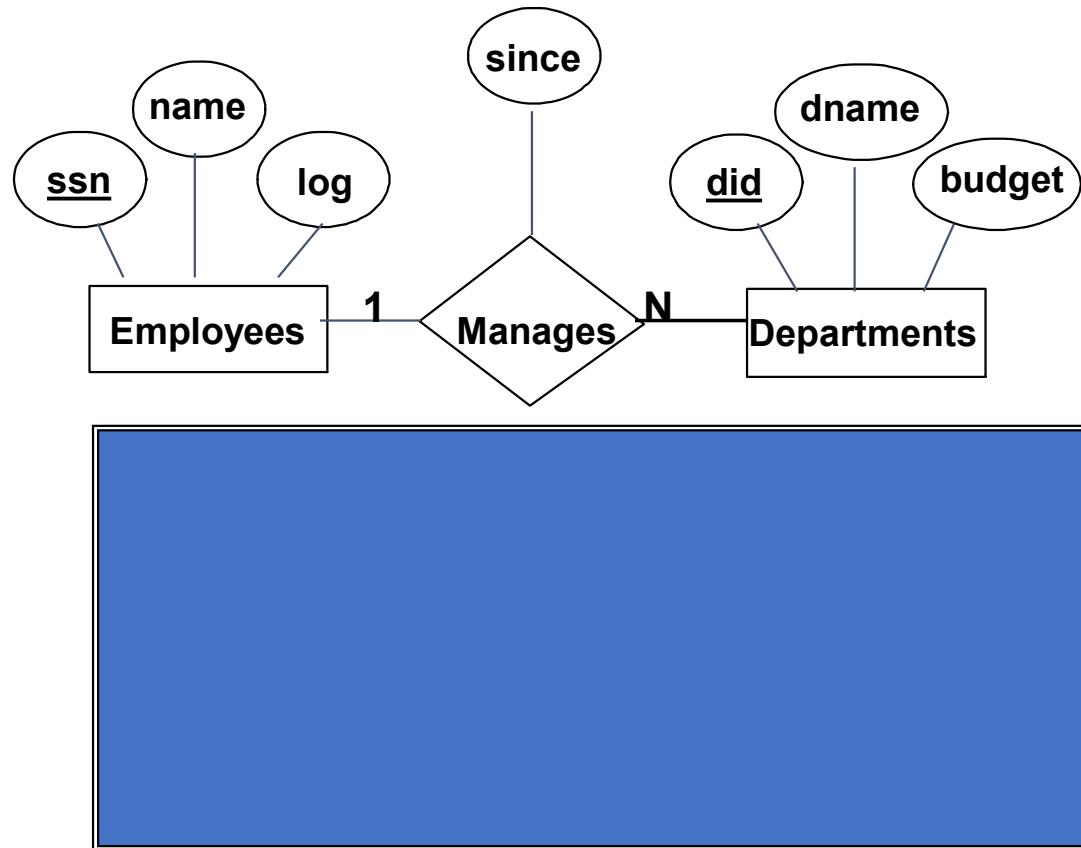


Translation to relational model?

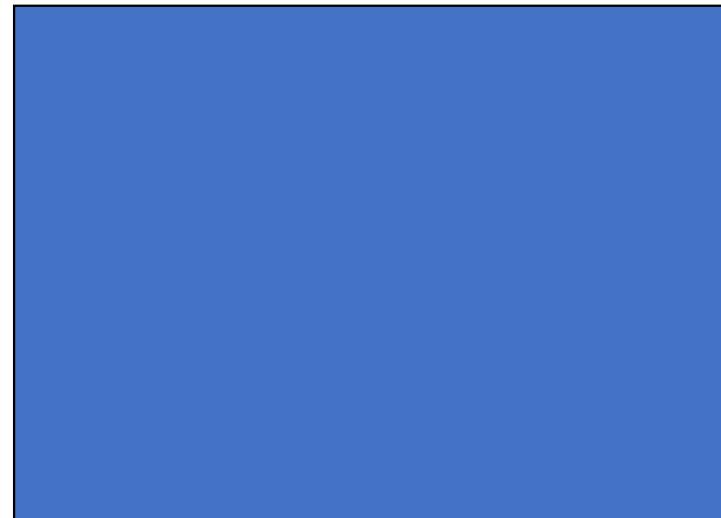
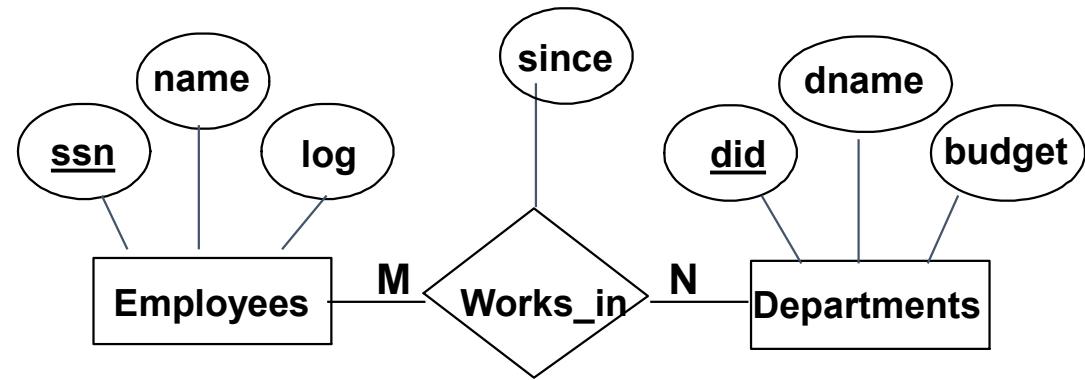
Review: Key Constraints

- Each dept has at most one manager, according to the *key constraint* on Manages.
1. Since each department has a unique manager, we could instead combine Manages and Departments.

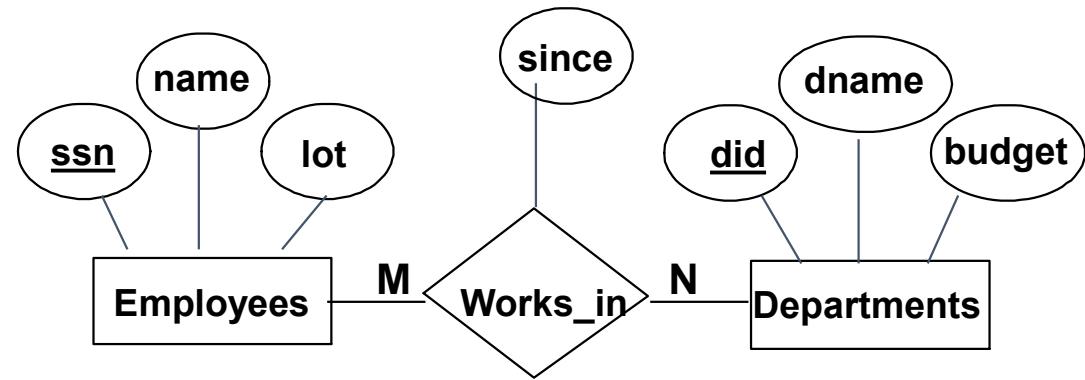
```
CREATE TABLE Employees
(ssn VARCHAR(11),
name VARCHAR(20),
log INTEGER,
PRIMARY KEY (ssn));
```



Relationship Sets to Tables

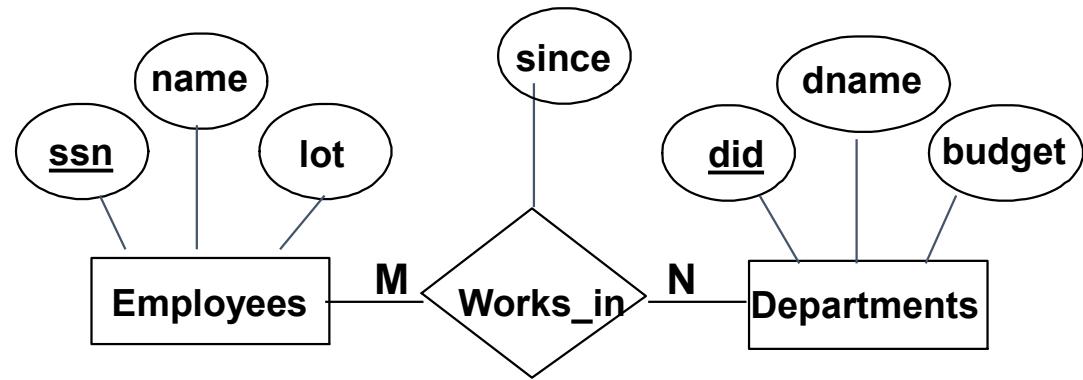


Relationship Sets to Tables



CREATE TABLE Works_In(

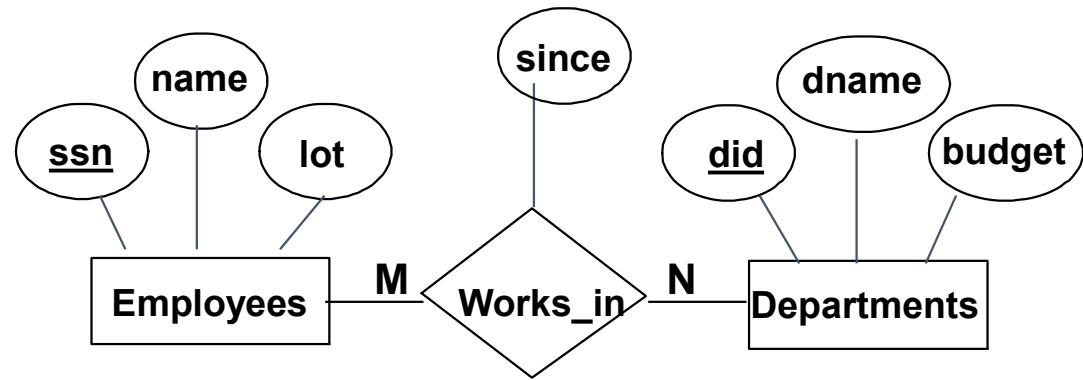
Relationship Sets to Tables



```
CREATE TABLE Works_In(  
    ssn VARCHAR(1),  
    did INTEGER,  
    since TEXT,
```



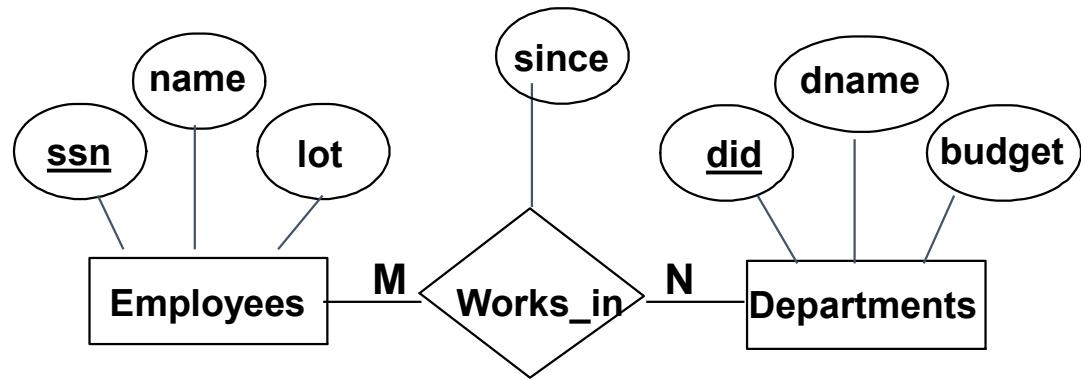
Relationship Sets to Tables



```
CREATE TABLE Works_In(  
    ssn VARCHAR(1),  
    did INTEGER,  
    since TEXT,  
    PRIMARY KEY (ssn, did),
```



Relationship Sets to Tables



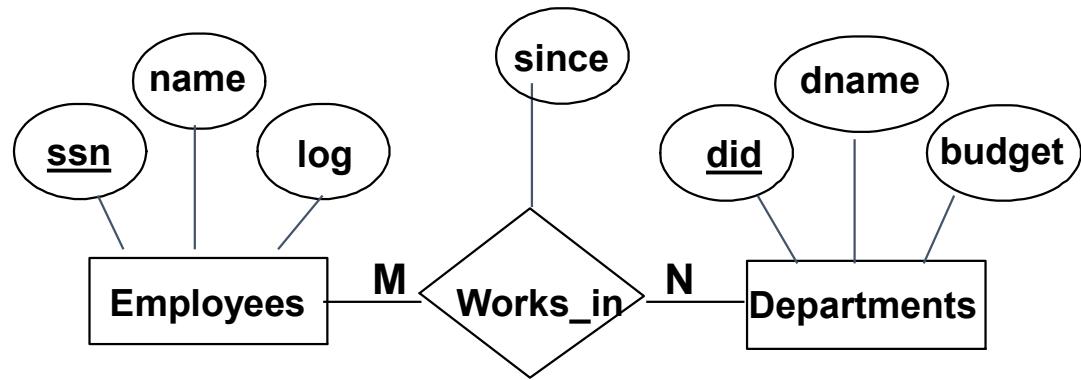
```
CREATE TABLE Works_In(  
    ssn VARCHAR(1),  
    did INTEGER,  
    since TEXT,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees(ssn),
```

Relationship Sets to Tables

- In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys).
 - All descriptive attributes.

```
CREATE TABLE Employees
(ssn VARCHAR(11),
name VARCHAR(20),
log INTEGER,
PRIMARY KEY (ssn));
```

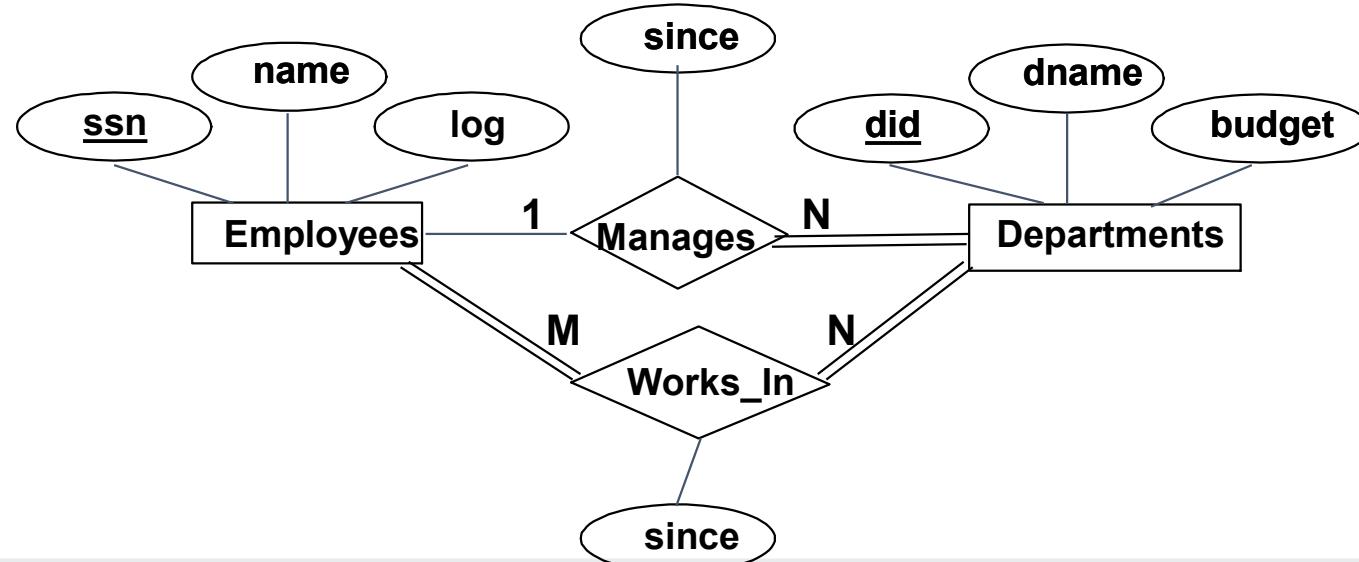
```
CREATE TABLE Departments
(did INTEGER,
dname VARCHAR(20),
budget REAL,
PRIMARY KEY (did));
```



```
CREATE TABLE Works_In(
ssn VARCHAR(11),
did INTEGER,
since TEXT,
PRIMARY KEY (ssn, did),
FOREIGN KEY (ssn)
    REFERENCES Employees(ssn),
FOREIGN KEY Departments (did)
    REFERENCES Departments (did) );
```

Review: Participation Constraints

- Does every department have a manager?
 - If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



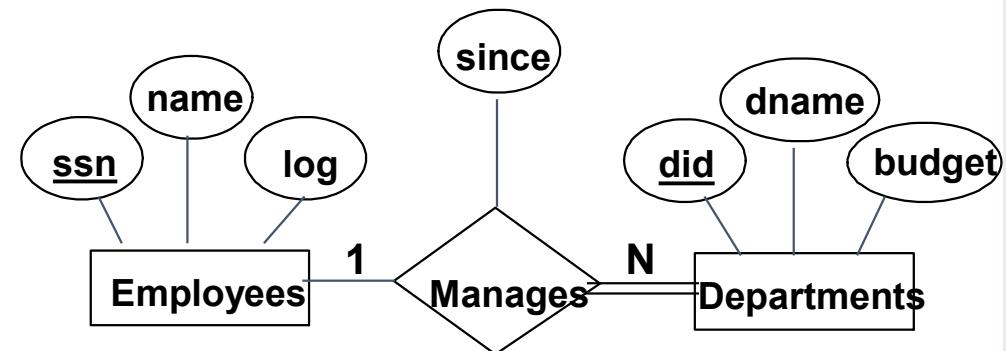
Participation Constraints in SQL

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname VARCHAR(20),  
    budget REAL,
```

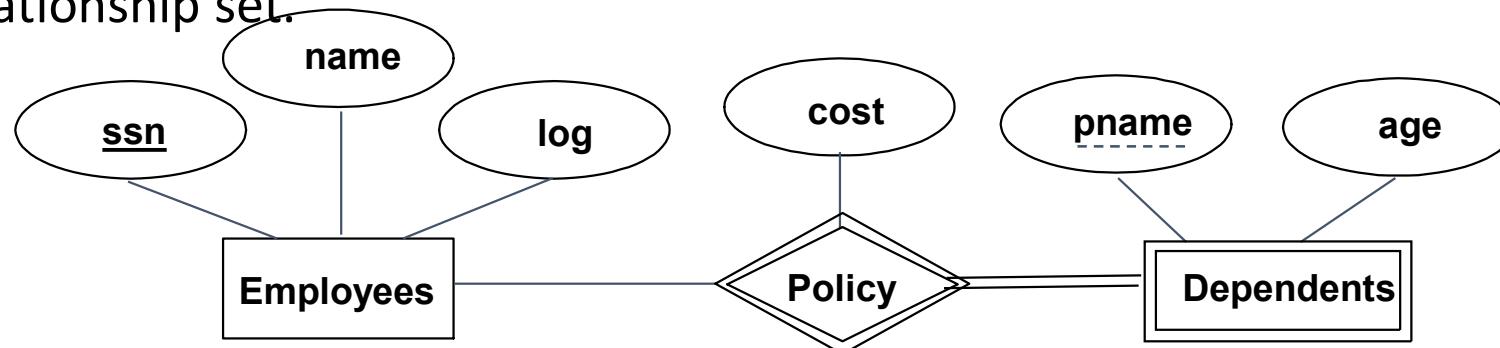
```
    since DATE,
```

```
    FOREIGN KEY (ssn) REFERENCES Employees(ssn)  
    ON DELETE CASCADE)
```



Review: Weak Entities

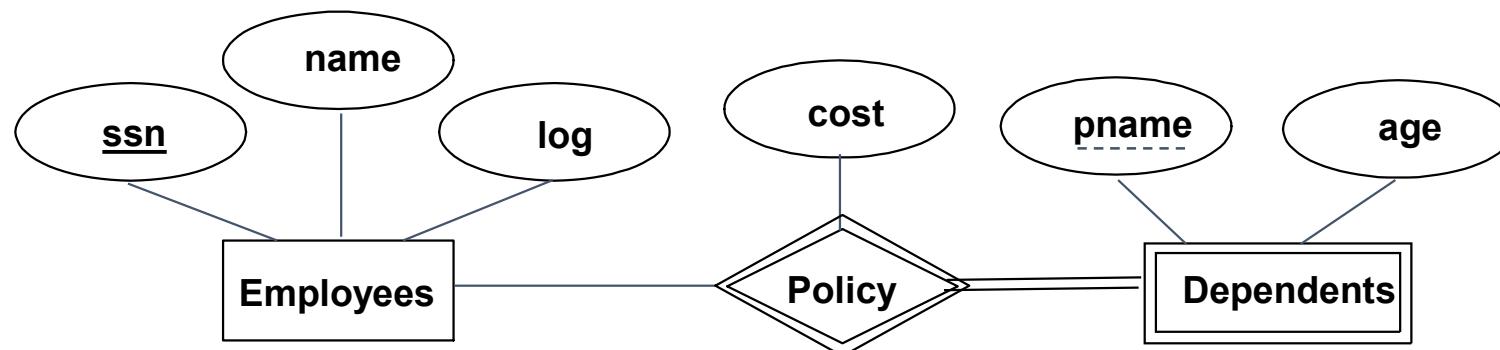
- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



Weak Entities

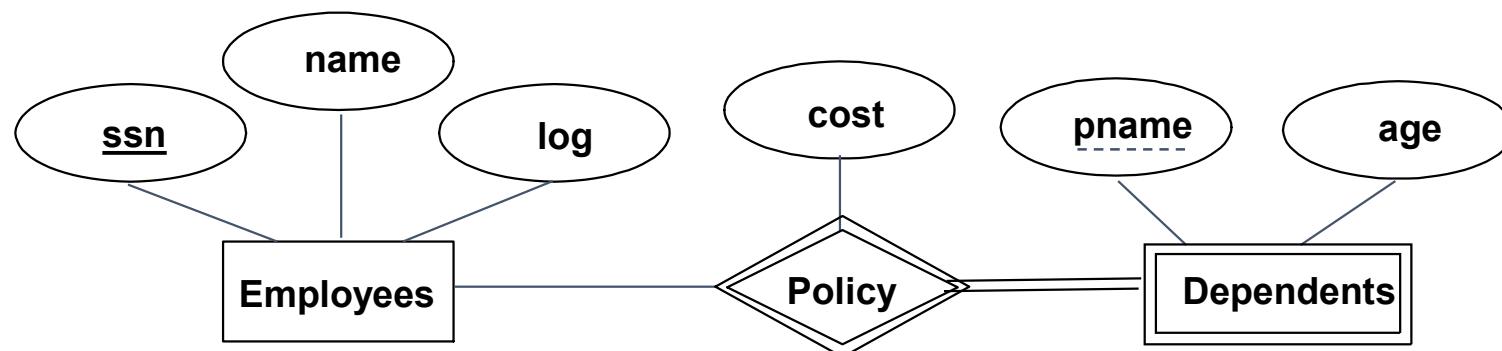
Weak entity set and identifying relationship set are translated into a single table.

- When the owner entity is deleted, all owned weak entities must also be deleted.



Weak Entities

```
CREATE TABLE Dep_Policy (  
    pname VARCHAR(20),  
    age INTEGER,  
    cost REAL,
```



Referential Integrity in SQL/92

- SQL/92 supports all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

DRAW ER DIAGRAM FOR THIS

```
CREATE TABLE Enrolled
(sid VARCHAR(20),
cid VARCHAR(20),
grade VARCHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students (sid)
ON DELETE CASCADE
ON UPDATE SET DEFAULT );
```

EXAMPLE

“For a given student and course, there is a single grade in Enrolled.”

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

```
CREATE TABLE Enrolled  
  (sid VARCHAR(20)  
   cid VARCHAR(20),  
   grade VARHAR(2),  
   PRIMARY KEY (sid,cid) );
```

UNIQUE keyword

Create Enrolled table

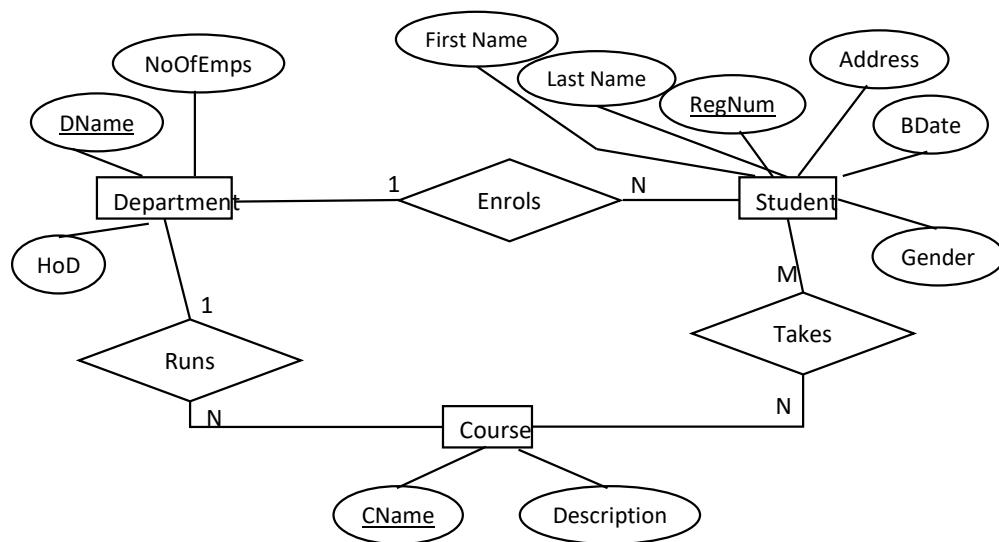
“Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”

sid	cid	grade
53834	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

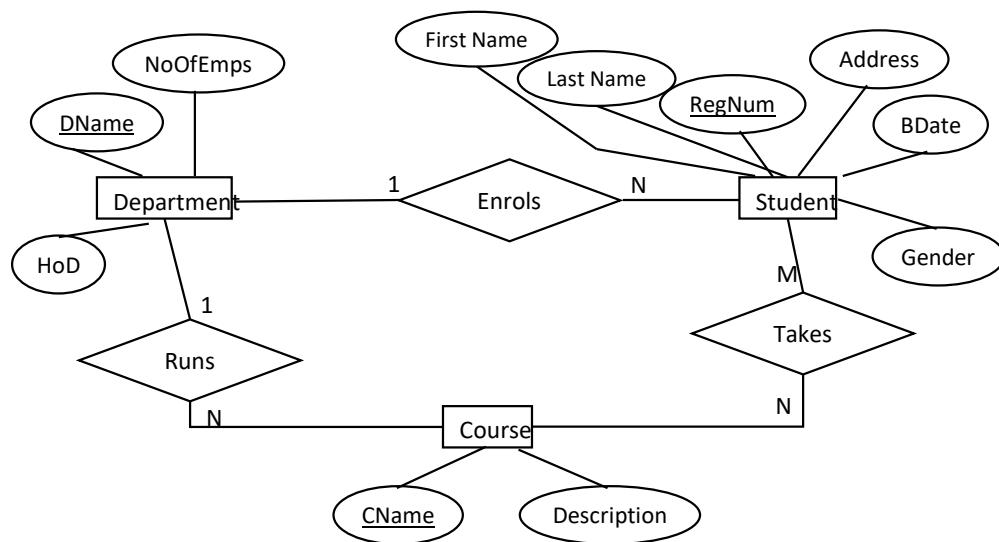
```
CREATE TABLE Enrolled  
(sid VARCHAR(20)  
cid  VARCHAR(20),  
grade VARCHAR(2),  
PRIMARY KEY (sid),  
UNIQUE (cid, grade) ;
```

Example

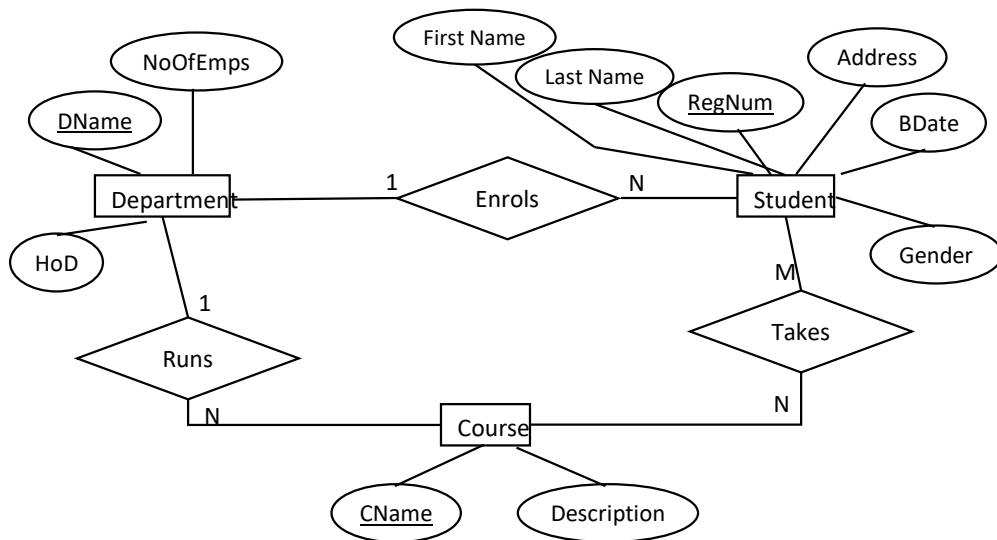


Example

- **CREATE TABLE** Department(DName TEXT, HoD TEXT, NoOfEmps INTEGER, PRIMARY KEY(DName));
-
-
-

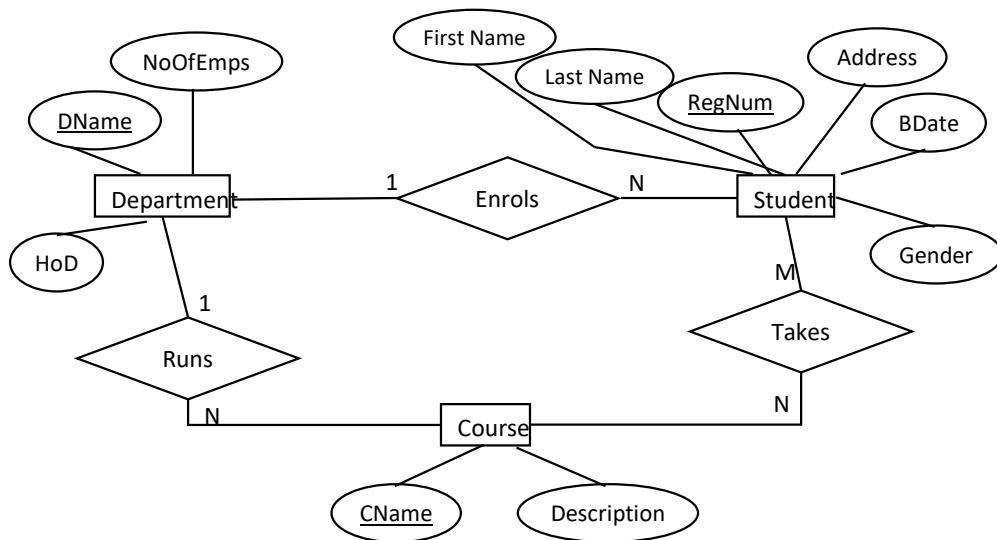


Example



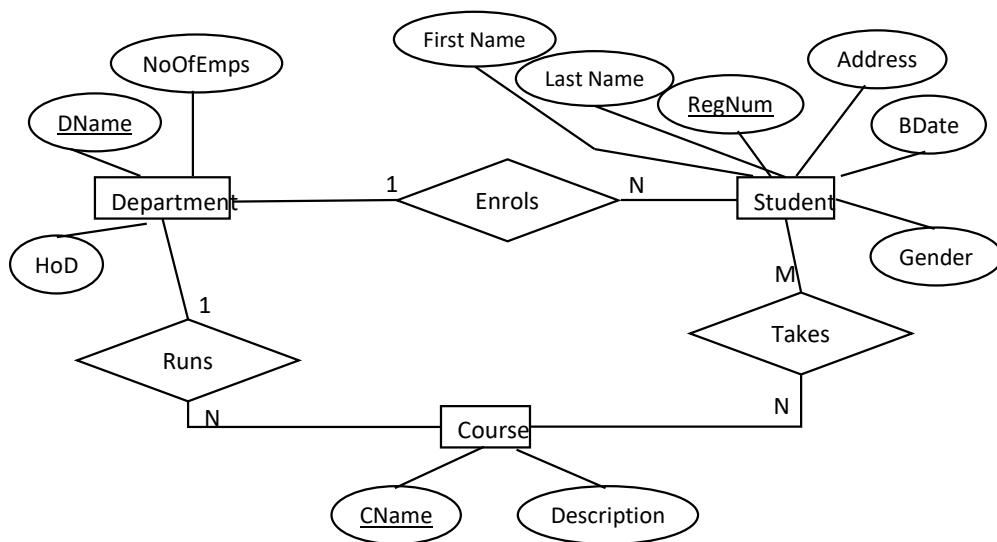
- CREATE TABLE Department(DName TEXT, HoD TEXT, NoOfEmps INTEGER, PRIMARY KEY(DName));
- **CREATE TABLE** StudentsEnrol(firstName TEXT, lastName TEXT, RegNumber INTEGER, Address TEXT, BDate TEXT, Gender TEXT, DepName TEXT, PRIMARY KEY (RegNumber), FOREIGN KEY(DepName) REFERENCES Department(DName) ON DELETE SET NULL);
-
-

Example



- CREATE TABLE Department(DName TEXT, HoD TEXT, NoOfEmps INTEGER, PRIMARY KEY(DName));
- CREATE TABLE StudentsEnrol(firstName TEXT, lastName TEXT, RegNumber INTEGER, Address TEXT, BDate TEXT, Gender TEXT, DepName TEXT, PRIMARY KEY (RegNumber), FOREIGN KEY(DepName) REFERENCES Department(DName) ON DELETE SET NULL);
- **CREATE TABLE** CourseRuns(CName TEXT, Description TEXT, DName TEXT, PRIMARY KEY (CName), FOREIGN KEY(DName) REFERENCES Department(DName) ON DELETE SET NULL);
-

Example



- CREATE TABLE Department(DName TEXT, HoD TEXT, NoOfEmp INTEGER, PRIMARY KEY(DName));
- CREATE TABLE StudentsEnrol(firstName TEXT, lastName TEXT, RegNumber INTEGER, Address TEXT, BDate TEXT, Gender TEXT, DepName TEXT, PRIMARY KEY (RegNumber), FOREIGN KEY(DepName) REFERENCES Department(DName) ON DELETE SET NULL);
- CREATE TABLE CourseRuns(CName TEXT, Description TEXT, DName TEXT, PRIMARY KEY (CName), FOREIGN KEY(Dname) REFERENCES Department(DName) ON DELETE SET NULL);
- **CREATE TABLE StTakesCourse(CName TEXT, RegNumber INTEGER, PRIMARY KEY(CName,RegNumber), FOREIGN KEY (CName) REFERENCES CourseRuns(CName), FOREIGN KEY (RegNumber) REFERENCES StudentsEnrol(RegNumber));**



Let's begin ☺

Last lecture

- We remembered



Last lecture

- We remembered
 - Keys



Last lecture

- We remembered
 - Keys
 - Relational Schema
-



Last lecture

- We remembered
 - Keys
 - Relational Schema
 - Participating and Multiplicity Constraints



Last lecture

- We remembered
 - Keys
 - Relational Schema
 - Participating and Multiplicity Constraints
 - We were introduced to:
 -
 -
 -
 -
 -

Last lecture

- We remembered
 - Keys
 - Relational Schema
 - Participating and Multiplicity Constraints
- We were introduced to:
 - Create Table
 - NOT NULL
 - Primary Key
 - UNIQUE
 -
 -
 -
 -
 -
 -

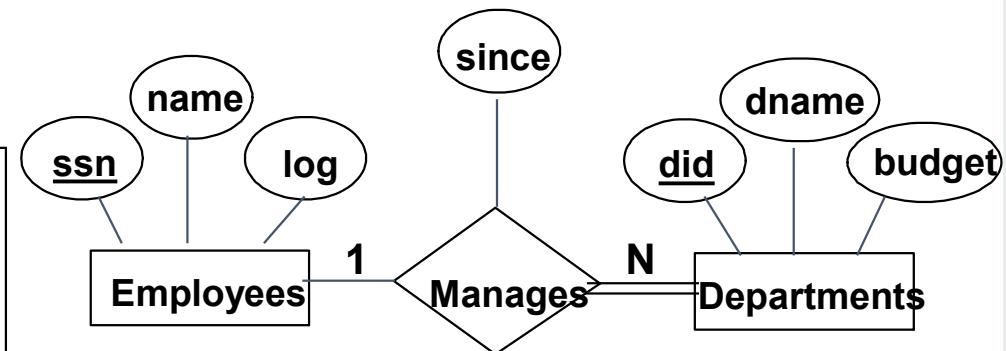
Last lecture

- We remembered
 - Keys
 - Relational Schema
 - Participating and Multiplicity Constraints
- We were introduced to:
 - Create Table
 - NOT NULL
 - Primary Key
 - UNIQUE
 - Foreign Key
 - ON DELETE
 - CASCADE
 - SET NULL
 - SET DEFAULT
 - ON UPDATE
 - SET DEFAULT
 - SET NULL

Integrity Constraints (Key, Participation) in SQL

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname VARCHAR(20),  
    budget REAL,  
    ssn VARCHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees(ssn)  
        ON DELETE CASCADE)
```



Some more introduction to SQL

- Inserting and querying

Adding and Deleting Tuples

One can insert a single tuple using

INSERT INTO <Relation Name> (<Attributes>) **Values** (<all_the_values>);

Directive

INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2);

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4

DELETE and SELECT

The **DELETE** statement is used to delete existing records in a table.

DELETE
FROM Relation-set
WHERE <Conditions>;

Adding and Deleting Tuples

Can delete all tuples satisfying some condition (e.g., name = Shero):

```
DELETE  
FROM Students S  
WHERE S.name = 'Jones';
```

Powerful variants of these commands are available; more later!

sid	name	login	age	gpa
53688	Smith	smith@ee	18	3.2

DELETE and SELECT

The **SELECT** statement is used to select data from a database. The data returned is stored in a result table, called the result set.

SELECT Attributes
FROM Relation-set
WHERE <Conditions>;

If Attributes = * the query evaluates to $\sigma_{conditions}$

If Attributes = attr1, attr2... then the query evaluates to
 $\pi_{attr1, attr2\dots}(\sigma_{conditions})$

The SQL Query Language

If Attributes = * the query evaluates to $\sigma_{age > 18}$

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2
53650	Shero	shero@math	19	3.8

Find all students
with age 18

```
SELECT *
FROM Students S
WHERE S.age=18;
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

The SQL Query Language

$$\pi_{name, login}(\sigma_{age > \text{?}})$$

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2
53650	Shero	shero@math	19	3.8

Names and logins of
all students
with age 18

```
SELECT S.name, S.login  
FROM Students S  
WHERE S.age=18;
```

name	login
Jones	jones@cs
Shero	shero@cs

Selection using Distinct to eliminate duplicates

- Basic query structure

```
SELECT      [DISTINCT] attribute-list
FROM        relation-list
WHERE       condition
```

- DISTINCT is an optional keyword indicating that duplicates should be eliminated. (Otherwise duplicate elimination is not done)

```
mysql> select DISTINCT e.NAME FROM Enrolled2 e WHERE NAME='URaz';
+-----+
| NAME |
+-----+
| URaz |
+-----+
1 row in set (0.00 sec)
```

Selection using logical operators

```
SELECT DISTINCT sname  
FROM Sailors  
WHERE Sailors.age>42 AND Sailors.rating>10
```

SELECT	[DISTINCT]	<i>attribute-list</i>
FROM		<i>relation-list</i>
WHERE		<i>condition</i>

- Conditions (ATTR *op* CONST or ATTR1 *op* ATTR2, where *op* is one of (<, >, =, ≤, ≥, ≠) combined using AND, OR and NOT.

These are logical operations!

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 -
 -
 -
 -
 -

```
SELECT      [DISTINCT] attribute-list
FROM        relation-list
WHERE       condition
```

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 -
 -
 -
 -
 -

```
SELECT      [DISTINCT]  attribute-list
FROM        relation-list
WHERE       condition
```

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they do not satisfy the *conditions*.
 -
 -
 -
-

```
SELECT      [DISTINCT]  attribute-list
FROM        relation-list
WHERE       condition
```

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they do not satisfy the *conditions*.
 - Display attributes that are in *attribute-list*.
 -
-

```
SELECT      [DISTINCT]  attribute-list
FROM        relation-list
WHERE       condition
```

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they do not satisfy the *conditions*.
 - Display attributes that are in *attribute-list*.
 - If DISTINCT is specified, eliminate duplicate rows.
-

```
SELECT      [DISTINCT]  attribute-list
FROM        relation-list
WHERE       condition
```

What will happen if more than relations are used: Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they do not satisfy the *conditions*.
 - Display attributes that are in *attribute-list*.
 - If DISTINCT is specified, eliminate duplicate rows.
- This strategy is probably the least efficient way to compute a query! An optimiser will find more efficient strategies to compute *the same answers*.

```
SELECT      [DISTINCT]  attribute-list
FROM        relation-list
WHERE       condition
```

Example of Conceptual Evaluation

```
SELECT sname
FROM Sailors, Reserves
WHERE Sailors.sid=Reserves.sid AND Reserves.bid=103
```

SELECT	[DISTINCT]	attribute-list
FROM	relation-list	
WHERE	condition	

$$\pi_{sname}(Sailors \bowtie_{Sailors.sid=Reserves.sid \wedge age>18} Reserves)$$

Query: Find names of sailors who Reserved boat number 103

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Figure 4.15 An Instance S_3 of Sailors

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Figure 4.16 An Instance R_2 of Reserves

Range Variables

Query: Find names
of sailors who
Reserved boat
number 103

```
SELECT sname  
FROM   Sailors, Reserves  
WHERE  Sailors.sid=Reserves.sid AND Reserves.bid=103
```

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=103
```

Range variables are
necessary when joining a
table with itself !!!

Expressions and Strings

SELECT	[DISTINCT]	<i>attribute-list</i>
FROM	<i>relation-list</i>	
WHERE		<i>condition</i>

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

- Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two new attributes defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
-
-

Expressions and Strings

SELECT	[DISTINCT]	<i>attribute-list</i>
FROM	<i>relation-list</i>	
WHERE	<i>condition</i>	

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2
FROM Sailors S
WHERE S.sname LIKE 'B_%B'
```

- Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two new attributes defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- **AS** and **=** are two ways to name fields in result.
-

Expressions and Strings

SELECT	[DISTINCT]	<i>attribute-list</i>
FROM	<i>relation-list</i>	
WHERE		<i>condition</i>

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

- Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two new attributes defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- **AS** and **=** are two ways to name fields in result.
- **LIKE** is used for string matching. **'_'** stands for any one character and **'%'** stands for 0 or more arbitrary characters.

Expressions and Strings

SELECT	[DISTINCT]	<i>attribute-list</i>
FROM	<i>relation-list</i>	
WHERE		<i>condition</i>

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

- Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two new attributes defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- **AS** and **=** are two ways to name fields in result.
- **LIKE** is used for string matching. **'_'** stands for any one character and **'%'** stands for 0 or more arbitrary characters.

Note: we can create attributes! In this example they are age1 and age2

Find sid's of sailors who've reserved a red or a green boat

B		R	
BID	Colr	SID	BID
b1	red	s1	b1
b2	grn	s1	b2
		s2	b1

- UNION: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
AND (B.color='red' OR B.color='green')
```

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
AND B.color='red'  
UNION  
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
AND B.color='green'
```

B.BID	B.Colr	R.SID	R.BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

Find sid's of sailors who've reserved a red but not a green boat

- EXCEPT : Used to compute the set difference of two *union-compatible* sets of tuples
- What do we get if we replace UNION with EXCEPT in the previous SQL query?

B

<u>BID</u>	Colr
b1	red
b2	grn

R

SID	<u>BID</u>
s1	b1
s1	b2
s2	b1

```
SELECT R.sid          A
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
EXCEPT
SELECT R.sid          B
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat



- AND : Used to compute the set intersection of two *union-compatible* sets of tuples

BID	Colr
b1	red
b2	grn

SID	BID
s1	b1
s1	b2
s2	b1

B.BID	B.Colr	R.SID	R.BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
AND (B.color='red' AND B.color='green')

Find sid's of sailors who've reserved a red and a green boat



B1 X R1

BID	Colr	SID	BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

Let us use *Tmp1* to for sid's of sailors who reserved red boats

B2 X R2

BID	Colr	SID	BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

Let us use *Tmp2* to for sid's of sailors who reserved green boats

Find sid's of sailors who've reserved a red and a green boat



B1 X R1

BID	Colr	SID	BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

Let us use *Tmp1* to for sid's of sailors who reserved red boats

Take cross product of these two relations.

B2 X R2

BID	Colr	SID	BID
b1	red	s1	b1
b1	red	s1	b2
b1	red	s2	b1
b2	grn	s1	b1
b2	grn	s1	b2
b2	grn	s2	b1

Let us use *Tmp2* to for sid's of sailors who reserved green boats

B2 X R2 X B1 X R1

BID	Colr	SID	BID	BID	Colr	SID	BID
b1	red	s1	b1	b1	red	s1	b1
b1	red	s1	b1	b1	red	s1	b2
b1	red	s1	b1	b1	red	s2	b1
b1	red	s1	b1	b2	grn	s1	b1
b1	red	s1	b1	b2	grn	s1	b2
b1	red	s1	b1	b2	grn	s2	b1
b1	red	s1	b2	b1	red	s1	b1
b1	red	s1	b2	b1	red	s1	b2
b1	red	s1	b2	b1	red	s2	b1
b1	red	s1	b2	b2	grn	s1	b1
b1	red	s1	b2	b2	grn	s1	b2
b1	red	s1	b2	b2	grn	s2	b1
b1	red	s2	b1	b1	red	s1	b1
b1	red	s2	b1	b1	red	s1	b2
b1	red	s2	b1	b1	red	s2	b1
b1	red	s2	b1	b2	grn	s1	b1
b1	red	s2	b1	b2	grn	s1	b2
b1	red	s2	b1	b2	grn	s2	b1

B2 X R2 X B1 X R1 (Continue)

BID	Colr	SID	BID	BID	Colr	SID	BID
b1	grn	s1	b1	b1	red	s1	b1
b1	grn	s1	b1	b1	red	s1	b2
b1	grn	S1	b1	b1	red	s2	b1
b1	grn	s1	b1	b2	grn	s1	b1
b1	grn	s1	b1	b2	grn	s1	b2
b1	grn	s1	b1	b2	grn	s2	b1
b1	grn	s1	b2	b1	red	s1	b1
b1	grn	s1	b2	b1	red	s1	b2
b1	grn	s1	b2	b1	red	s2	b1
b1	grn	s1	b2	b2	grn	s1	b1
b1	grn	s1	b2	b2	grn	s1	b2
b1	grn	s1	b2	b2	grn	s2	b1
b1	grn	s2	b1	b1	red	s1	b1
b1	grn	s2	b1	b1	red	s1	b2
b1	grn	s2	b1	b2	grn	s1	b1
b1	grn	s2	b1	b2	grn	s1	b2
b1	grn	s2	b1	b2	grn	s2	b1



```

SELECT R1.sid
FROM Boats B1, Reserves R1, Boats B2,
Reserves R2
WHERE R1.sid = R2.sid AND R1.bid=B1.bid
AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='green')

```

Find sid's of sailors who've reserved
a red and a green boat

B2 X R2 X B1 X R1 (Continue)

BID	Colr	SID	BID	BID	Colr	SID	BID
<u>b1</u>	<u>grn</u>	<u>s1</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b1</u>
<u>b1</u>	<u>grn</u>	<u>s1</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b2</u>
<u>b1</u>	<u>grn</u>	<u>S1</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s2</u>	<u>b1</u>
b1	grn	s1	b1	b2	grn	s1	b1
b1	grn	s1	b1	b2	grn	s1	b2
b1	grn	s1	b1	b2	grn	s2	b1
<u>b1</u>	<u>grn</u>	<u>s1</u>	<u>b2</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b1</u>
<u>b1</u>	<u>grn</u>	<u>s1</u>	<u>b2</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b2</u>
<u>b1</u>	<u>grn</u>	<u>s1</u>	<u>b2</u>	<u>b1</u>	<u>red</u>	<u>s2</u>	<u>b1</u>
b1	grn	s1	b2	b2	grn	s1	b1
b1	grn	s1	b2	b2	grn	s1	b2
b1	grn	s1	b2	b2	grn	s2	b1
<u>b1</u>	<u>grn</u>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b1</u>
<u>b1</u>	<u>grn</u>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s1</u>	<u>b2</u>
<u>b1</u>	<u>grn</u>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<u>red</u>	<u>s2</u>	<u>b1</u>
b1	grn	s2	b1	b2	grn	s1	b1
b1	grn	s2	b1	b2	grn	s1	b2
b1	grn	s2	b1	b2	grn	s2	b1

```
SELECT R1.sid
FROM Boats B1, Reserves R1, Boats B2,
Reserves R2
WHERE R1.sid = R2.sid AND R1.bid=B1.bid
AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='green')
```

B2 X R2 X B1 X R1 (Continue)

SELECT R1.sid
 FROM Boats B1, Reserves R1, Boats B2,
 Reserves R2
 WHERE R1.sid = R2.sid AND R1.bid=B1.bid
 AND R2.bid=B2.bid
 AND (B1.color='red' **AND** B2.color='green')

BID	Colr	SID	BID	BID	Colr	SID	BID
<u>b1</u>	<i>grn</i>	<u>s1</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s1</u>	<u>b1</u>
<u>b1</u>	<i>grn</i>	<u>s1</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s1</u>	<u>b2</u>
<u>b1</u>	<i>grn</i>	<u>S1</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s2</u>	<u>b1</u>
b1	<i>grn</i>	s1	b2	b1	red	s1	b1
b1	<i>grn</i>	s1	b2	b1	red	s1	b2
b1	<i>grn</i>	s1	b2	b1	red	s2	b1
<u>b1</u>	<i>grn</i>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s1</u>	<u>b1</u>
<u>b1</u>	<i>grn</i>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s1</u>	<u>b2</u>
<u>b1</u>	<i>grn</i>	<u>s2</u>	<u>b1</u>	<u>b1</u>	<i>red</i>	<u>s2</u>	<u>b1</u>

B2X R2 X B1 X R1 (Continue)

BID	Colr	SID	BID	BID	Colr	SID	BID
<u>b1</u>	<i>grn</i>	<i>s1</i>	<u>b1</u>	<u>b1</u>	<i>red</i>	<i>s1</i>	<u>b1</u>
b1	grn	s1	b1	b1	red	s1	b2
<u>b1</u>	<i>grn</i>	<i>S1</i>	<u>b1</u>	<u>b1</u>	<i>red</i>	<i>s2</i>	<u>b1</u>
<u>b1</u>	<i>grn</i>	<i>s2</i>	<u>b1</u>	<u>b1</u>	<i>red</i>	<i>s1</i>	<u>b1</u>
b1	grn	s2	b1	b1	red	s1	b2
<u>b1</u>	<i>grn</i>	<i>s2</i>	<u>b1</u>	<u>b1</u>	<i>red</i>	<i>s2</i>	<u>b1</u>



B2 X R2 X B1 X R1 (Continue)



```
SELECT R1.sid
FROM Boats B1, Reserves R1, Boats B2,
Reserves R2
WHERE R1.sid = R2.sid AND R1.bid=B1.bid
AND R2.bid=B2.bid
AND (B1.color='red' AND B2.color='green')
```

BID	Colr	SID	BID	BID	Colr	SID	BID
b1	<i>grn</i>	s1	b1	b1	<i>red</i>	s1	b1
b1	<i>grn</i>	S1	b1	b1	red	s2	b1
b1	<i>grn</i>	s2	b1	b1	red	s1	b1
b1	<i>grn</i>	s2	b1	b1	red	s2	b1

B2X R2 X B1 X R1 (Continue)

SO the answer to the query is s1



Find names of sailors who've reserved a red and a green boat.

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Figure 4.15 An Instance *S3* of Sailors

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Figure 4.16 An Instance *R2* of Reserves

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Figure 4.17 An Instance *B1* of Boats

$$\rho(\text{Tempred}, \pi_{\textit{sid}}((\sigma_{\textit{color}=\text{'red'}} \textit{Boats}) \bowtie \textit{Reserves}))$$

$$\rho(\text{Tempgreen}, \pi_{\textit{sid}}((\sigma_{\textit{color}=\text{'green'}} \textit{Boats}) \bowtie \textit{Reserves}))$$

$$\pi_{\textit{sname}}((\text{Tempred} \cap \text{Tempgreen}) \bowtie \textit{Sailors})$$

Find sid's of sailors who've reserved a red and a green boat



- **INTERSECT:** Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- Included in the SQL/92 standard, but some systems don't support it.
- Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

```
SELECT R.sid  
FROM Boats B1, Reserves R1,  
      Boats B2, Reserves R2  
WHERE R1.sid = R2.sid AND  
      R1.bid=B1.bid AND R2.bid=B2.bid  
      AND (B1.color='red' AND B2.color='green')
```

```
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
      AND B.color='red'  
INTERSECT  
SELECT R.sid  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid  
      AND B.color='green'
```

Nested Queries

Find names of sailors who've reserved boat #103:

```
SELECT S.sname
  FROM Sailors S
 WHERE S.sid IN (SELECT R.sid (ONLY ONE Column)
                  FROM Reserves R
                 WHERE R.bid=103)
```

- **IN** returns true if the first value is contained in the set returned (single column!)
-)WHERE clause can itself contain an SQL query! (As well as FROM and HAVING clauses which we will see later on.)

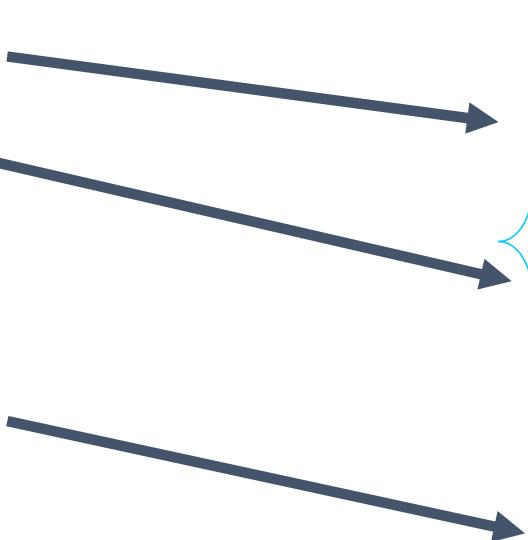
Nested Queries

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid IN (SELECT R.sid  
FROM Reserves R  
WHERE R.bid=103)
```

Find names of sailors who've reserved boat #103:



<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5



<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Nested Queries

```
SELECT S.sname
FROM Sailors S
WHERE S.sid NOT IN (SELECT R.sid
                     FROM Reserves R
                     WHERE R.bid=103)
```

*Find names of sailors
Who reserved a boat that is not
boat #103:
Answer must be 64*



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

*Find names of sailors
who did NOT reserve boat #103:*

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Nested Queries with Correlation

- EXISTS returns true **TRUE** if the set, is nonempty.
EXISTS operator is another set comparison operator, like IN.

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
               FROM Reserves R
              WHERE R.bid=103 AND S.sid=R.sid)
```

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

*Find names of sailors
who've reserved boat #103:*

- To understand semantics of correlated queries, think of a nested loops evaluation:
For each Sailors tuple, check the qualification by computing the subquery.

```
For (i=1...10) do {
    For (j=1...5) do {
        y=i+1;
    }
}
```

Nested Queries with Correlation



Find names of sailors who reserved a boat at most once

- UNIQUE construct can be used.
- UNIQUE checks for duplicate tuples. Returns TRUE if the corresponding set **does not** contain duplicates.

```
SELECT S.sname  
FROM Sailors S  
WHERE UNIQUE (SELECT R.bid  
              FROM Reserves R  
              WHERE S.sid=R.sid)
```

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *condition* ANY, *condition* ALL $>, <, =, \geq, \leq, \neq, \not{<}'$
- Find sailors whose rating is greater than that of some sailor called Horatio:*

ANY (returns true if there exist tuples (returned from nested WHERE clause) obey the condition!)

ALL (returns true if all tuples (returned from nested WHERE clause) obey the condition!)

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.sname='Horatio')
```

sid	sname	rating	age			
22	Dustin	7	45.0	22	101	10/10/98
29	Brutus	1	33.0	22	102	10/10/98
31	Lubber	8	55.5	22	103	10/8/98
32	Andy	8	25.5	22	104	10/7/98
58	Rusty	10	35.0	31	102	11/10/98
64	Horatio	7	35.0	31	103	11/6/98
71	Zorba	10	16.0	31	104	11/12/98
74	Horatio	9	35.0	64	101	9/5/98
85	Art	3	25.5	64	102	9/8/98
95	Bob	3	63.5	74	103	9/8/98

More on Set-Comparison Operators

- We've already seen IN, EXISTS and UNIQUE. Can also use NOT IN, NOT EXISTS and NOT UNIQUE.
- Also available: *condition* ANY, *condition* ALL $>, <, =, \geq, \leq, \neq, \not{<}'$
- Find sailors whose rating is greater than that of ALL sailor's called Horatio:***

ANY (returns true if there exist tuples (returned from nested WHERE clause) obey the condition!)

ALL (returns true if all tuples (returned from nested WHERE clause) obey the condition!)

```
SELECT *
FROM Sailors S
WHERE S.rating > ALL (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.sname='Horatio')
```

sid	sname	rating	age			
22	Dustin	7	45.0	22	101	10/10/98
29	Brutus	1	33.0	22	102	10/10/98
31	Lubber	8	55.5	22	103	10/8/98
32	Andy	8	25.5	22	104	10/7/98
58	Rusty	10	35.0	31	102	11/10/98
64	Horatio	7	35.0	31	103	11/6/98
71	Zorba	10	16.0	31	104	11/12/98
74	Horatio	9	35.0	64	101	9/5/98
85	Art	3	25.5	64	102	9/8/98
95	Bob	3	63.5	74	103	9/8/98

Division in SQL

Find sailors who've reserved all boats.

`EXISTS` (returns true if the set is nonempty)

`NOT EXISTS` (returns true if the set is empty)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
  ((SELECT B.bid
    FROM Boats B)
   EXCEPT
   (SELECT R.bid
    FROM Reserves R
     WHERE R.sid = S.sid))
```

A - B

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Figure 4.15 An Instance S3 of Sailors

sid	bid	day
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Figure 4.16 An Instance R2 of Reserves

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Figure 4.17 An Instance B1 of Boats

Set A= The set of all boats.

Set B= The set of all boats reserved by the current sailor S.sid

A-B is the boats that S.sid did not reserve. So if this set is empty, then return this s.sname!