

SCC.211 Operating Systems

Session 8

Dr Andrew Scott
a.scott@lancaster.ac.uk

Overview

- Topic 7: Memory Protection (summary/ recap)
- Topic 8: Paging (introduction)

7. Memory Protection

Summary/ Recap


Overview

- Basic approach
- The Memory Management Unit
- Segmentation
- Paging

Memory Protection

- How to split memory between processes
 - Problems and techniques apply more generally
 - e.g. disks and filesystems
- How to protect a process's memory from other processes
- How to limit access to memory-mapped devices
 - They're just seen within/ as part of 'normal' memory
 - In theory, any code can access these locations
- Remember privileged I/O instructions 'protect' isolated I/O devices
 - Normal (user level) code cannot use/ execute these instructions
 - We don't generally have to worry about user code accessing these

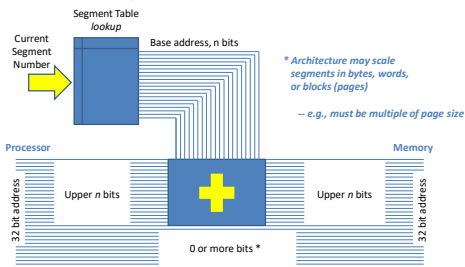
Some Simple Process Questions




Paging vs. Segmentation

- Segmentation
 - Each segment has own address space: offset from base
 - Works by adding segment base to given address/ offset
- Paging
 - All processes see same address range
 - Translation allows any page to map to any physical frame
 - Works by translating/ switching upper bits of given address, frame for page

Segmentation operation



Paging operation



Paging vs. Segmentation: *Instructions*

Segmentation

- Two options for segment table(s):
 - Code & Data share same table
 - Just goes off segment number
 - Code & Data each have own tables
 - Code → code segment table
 - Data → data segment table

LOAD 4, 0x5624 ; 4 is a data seg.

JUMP 3, 0x0088 ; 3 a code seg.

May also have one or more segment registers

-- giving default, if instructions omit explicit segment number

Paging

- Simple linear address
 - Just as we're used to:

LOAD 0x00fec90b

JUMP 0x04ad8724

Fixed Sized Memory Allocation: *Paging*

The diagram illustrates memory layout and paging. On the left, a vertical stack of memory segments is shown from Address 0 at the bottom to Address Limit at the top. The segments are: Code (purple), Data (green), Heap (red, with an upward arrow), and Stack (blue, with a downward arrow). On the right, a vertical stack of physical memory pages is shown. A bracket labeled 'Add pages as needed' spans several pages. Arrows indicate the mapping: Code maps to Read Only pages; Data maps to Read Only pages; Heap and Stack map to Read Only/Execute pages.

Segmentation

- Now rarely (actively) use segmentation
 - Offers flexible protection (possibly even to object level), but...
 - More difficult to share small amounts of memory (code or data)
 - ...soon end up with many segments to manage
 - Unlike paging, segmentation visible to code/ programmer

© Andrew Scott 2020-2022

4

Segmentation: *Common Configuration*

- In mixed/ hybrid systems, typically configure system to rely 'solely' on paging
 - In hybrid system, segmentation still provides User/ System separation, and Code/ Data protection
 - ...user segments configured for least privilege (ring 3), and so 'data' can't be run as code (-x)
 - In this case, all segments configured to span full memory range, as in common x86 set up:

Max Memory	Configured Segment Descriptors				
	Seg. 0	Seg. 1	Seg. 2	Seg. 3	Seg. 4
0	NULL/ empty	Kernel Code (Ring 0 +x)	Kernel Data (Ring 0 -x)	User Code (Ring 3 +x)	User Data (Ring 3 -x)


Moving from System to Kernel mode

- For example, an x86 kernel finishes initialisation and initiates the first user process...

```
printk ("\\n\\n===== Entering x86 user mode =====\\n\\n");
asm volatile (
    "= usermode:      call $40+3, $usermode;      "
    );
```

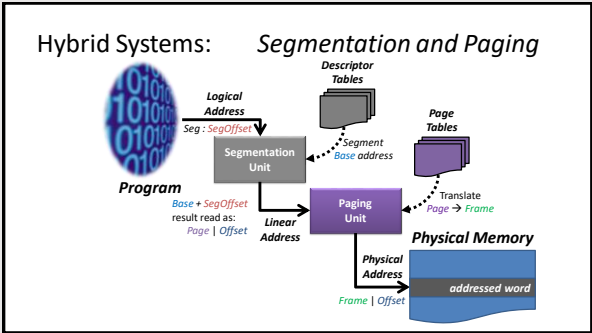
→ 16 → 32 → 48				
Seg. 0	Seg. 1	Seg. 2	Seg. 3	Seg. 4
NULL/ empty	Kernel Code (Ring 0 +x)	Kernel Data (Ring 0 -x)	User Code (Ring 3 +x)	User Data (Ring 3 -x)

Hybrid Systems: *Segmentation...*



Unless an x86 instruction explicitly specifies a segment* the segment is taken to be that set in the code or data segment register

* In protected mode this is an offset/index into the Descriptor Table, e.g. JUMP KernelCodeSeg, offset



8. Paging

Introduction

Paging Topics

- Handling large address spaces
 - Translations become increasingly expensive as memory size grows
- Demand paging
 - Growing a process by adding more memory/ pages
- Shared memory
 - Supporting everything from libraries to inter-process communication

Paging Topics II

- Translation Look-aside Buffer
 - Reducing the cost of caching translations
- Speeding up process creation
 - Using paging to reduce cost of *fork()* and *fork()...exec()*
