

# Unit 13: Instruction Processing, JIT and Stack Machines

(aka. "Fitting it all Together")

SCC 312 Compilation

John Vidler

[j.vidler@lancaster.ac.uk](mailto:j.vidler@lancaster.ac.uk)



# Split-Stage Compilation

---

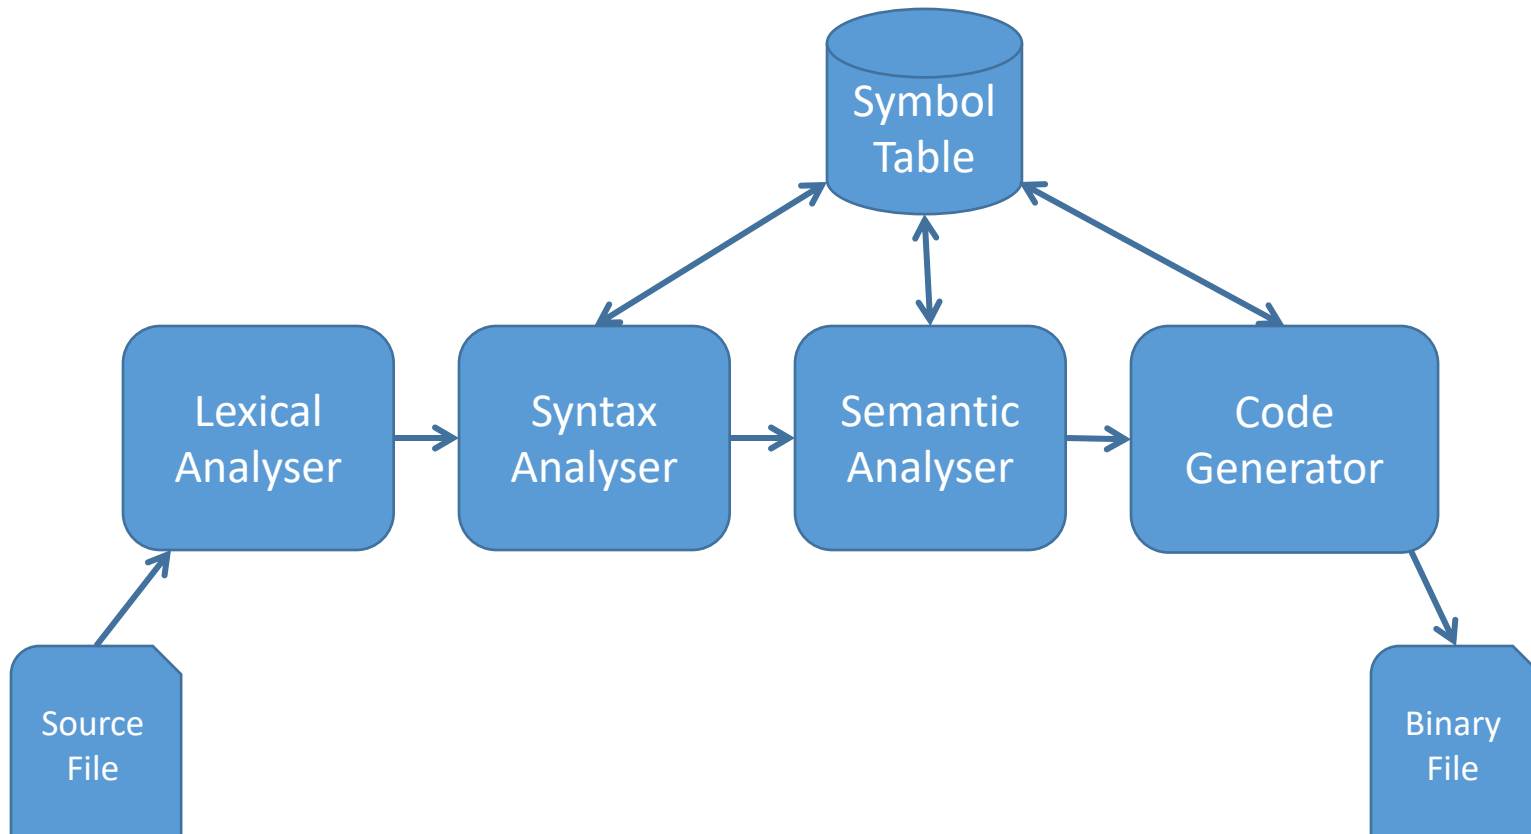
- This exists in two forms
  - Compilers with an intermediate language (IL)
  - Compilers which export code intended for an *interpreter*
- In both cases, rather than compile to a specific target, a halfway language is used
- This must be high level enough to be flexible ...
  - ... but cannot be too high level, or the final compiler steps will take too long!



# Intermediate Language Compilers

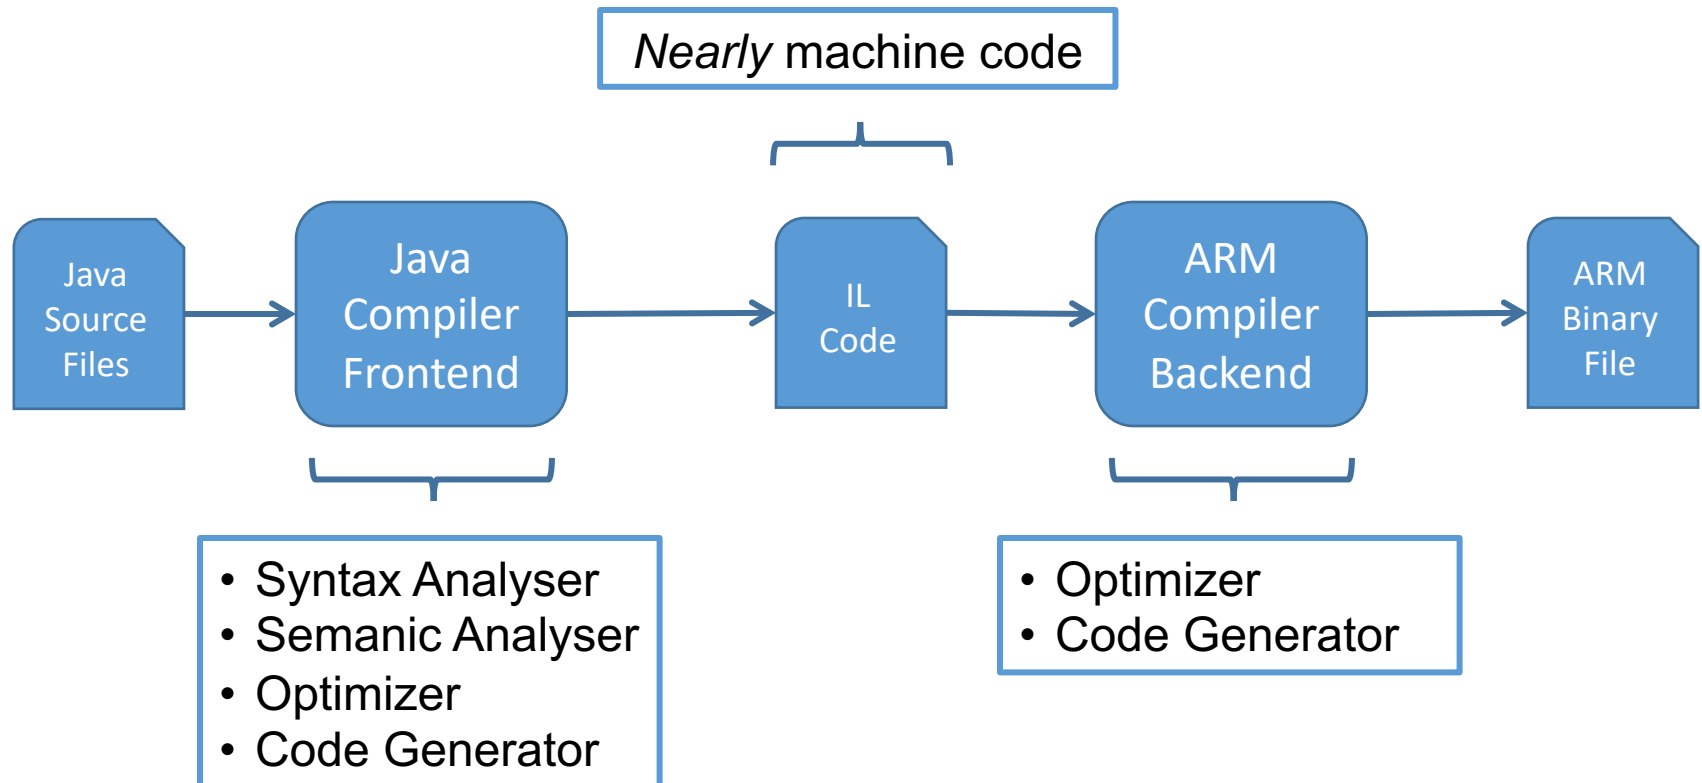
# A Structure Reminder

What we have seen so far



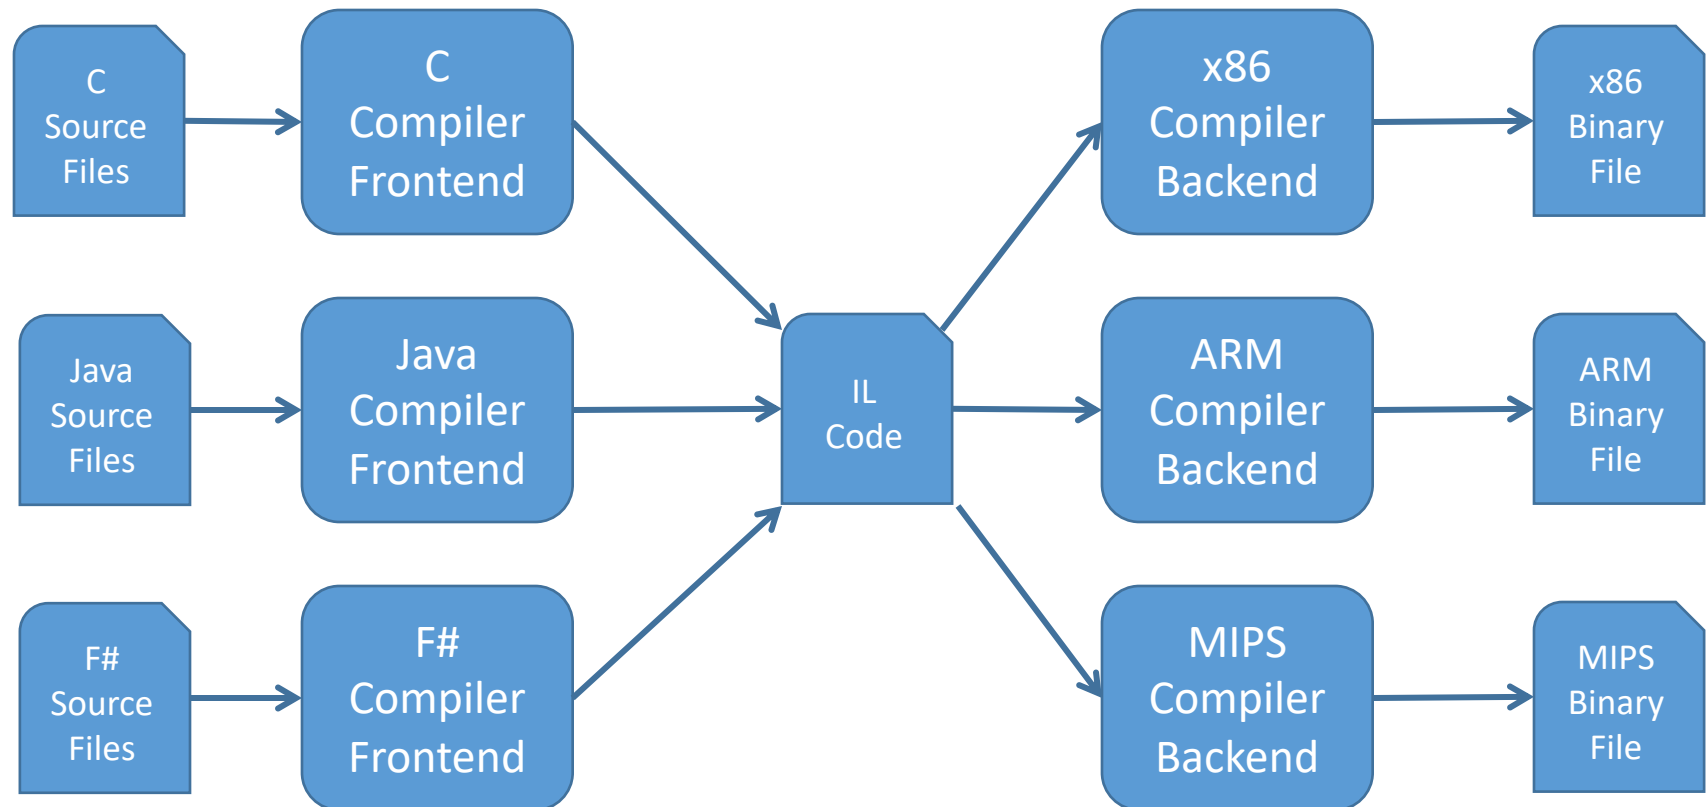
# Split Stage Compilation

(Via an Intermediate Language)



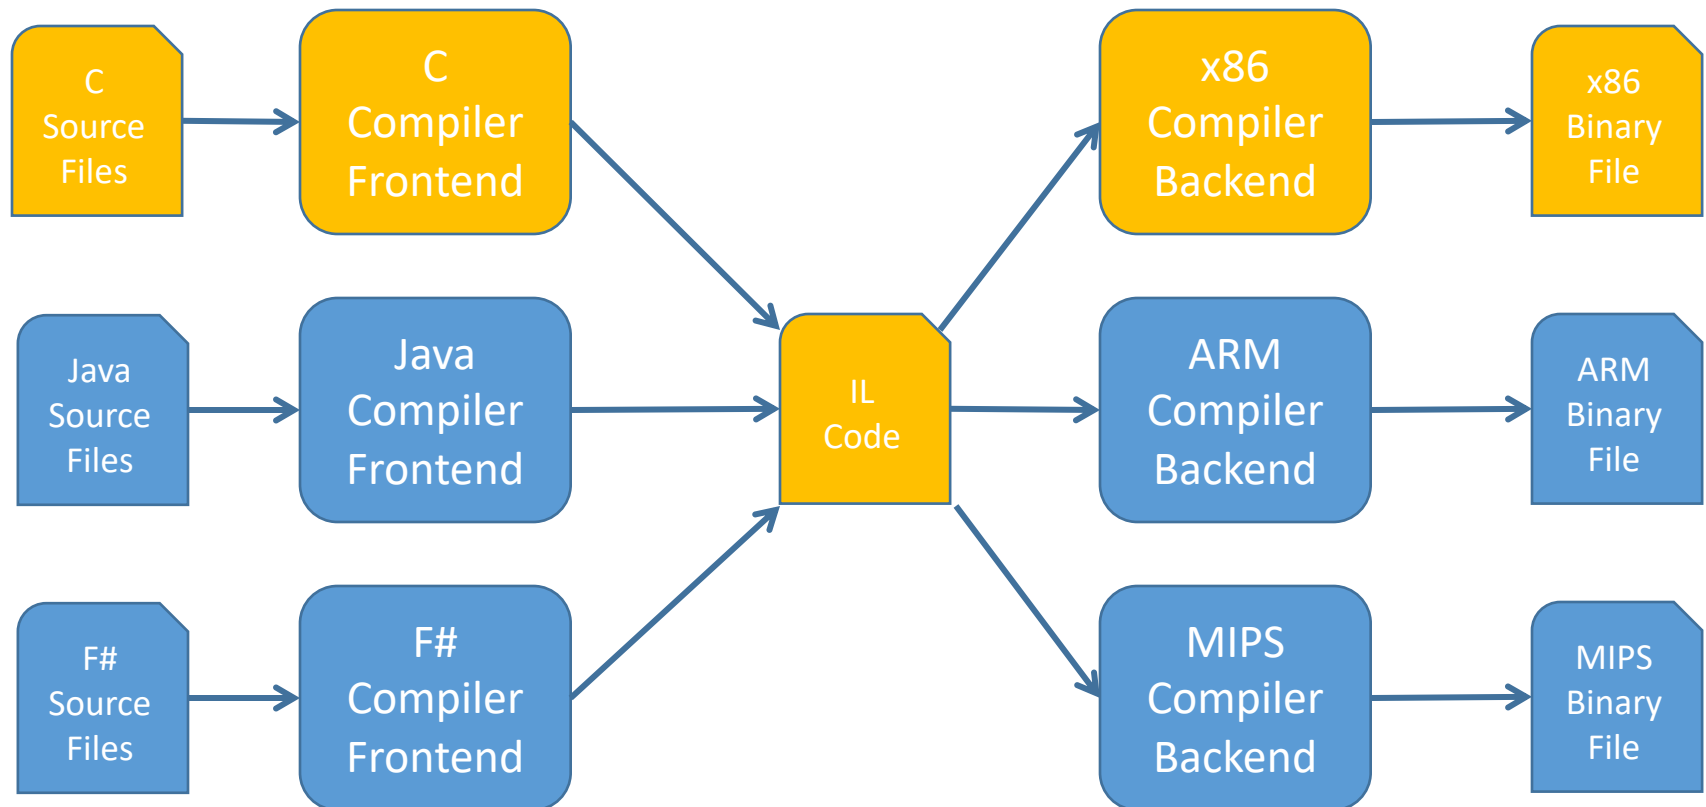
# Split Stage Compilation

(Via an Intermediate Language)



# Split Stage Compilation

(Via an Intermediate Language)

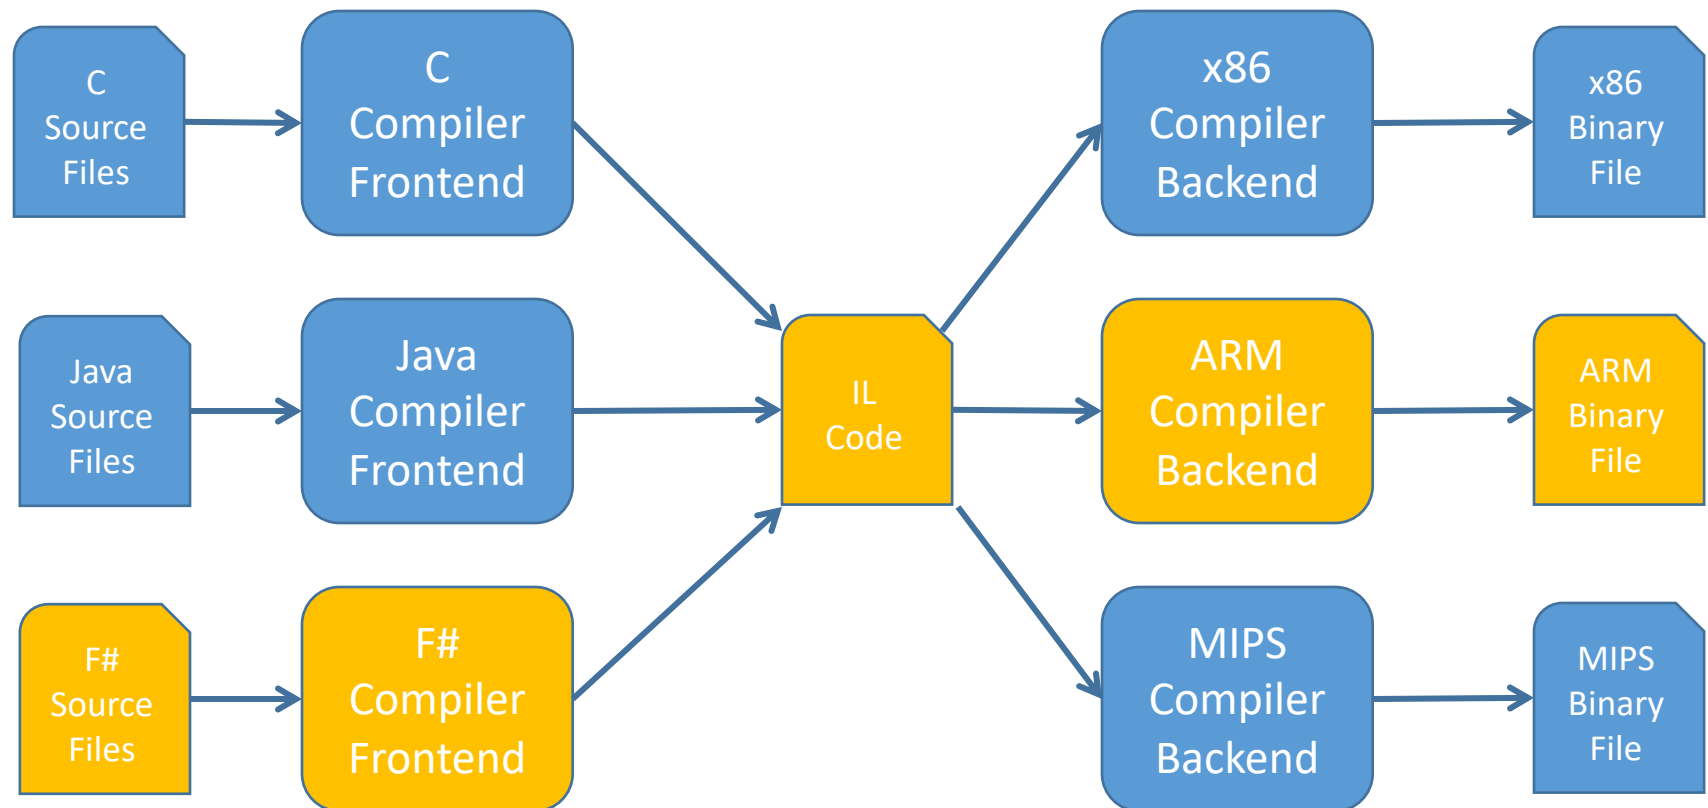


Building C code for x86 target platforms



# Split Stage Compilation

(Via an Intermediate Language)



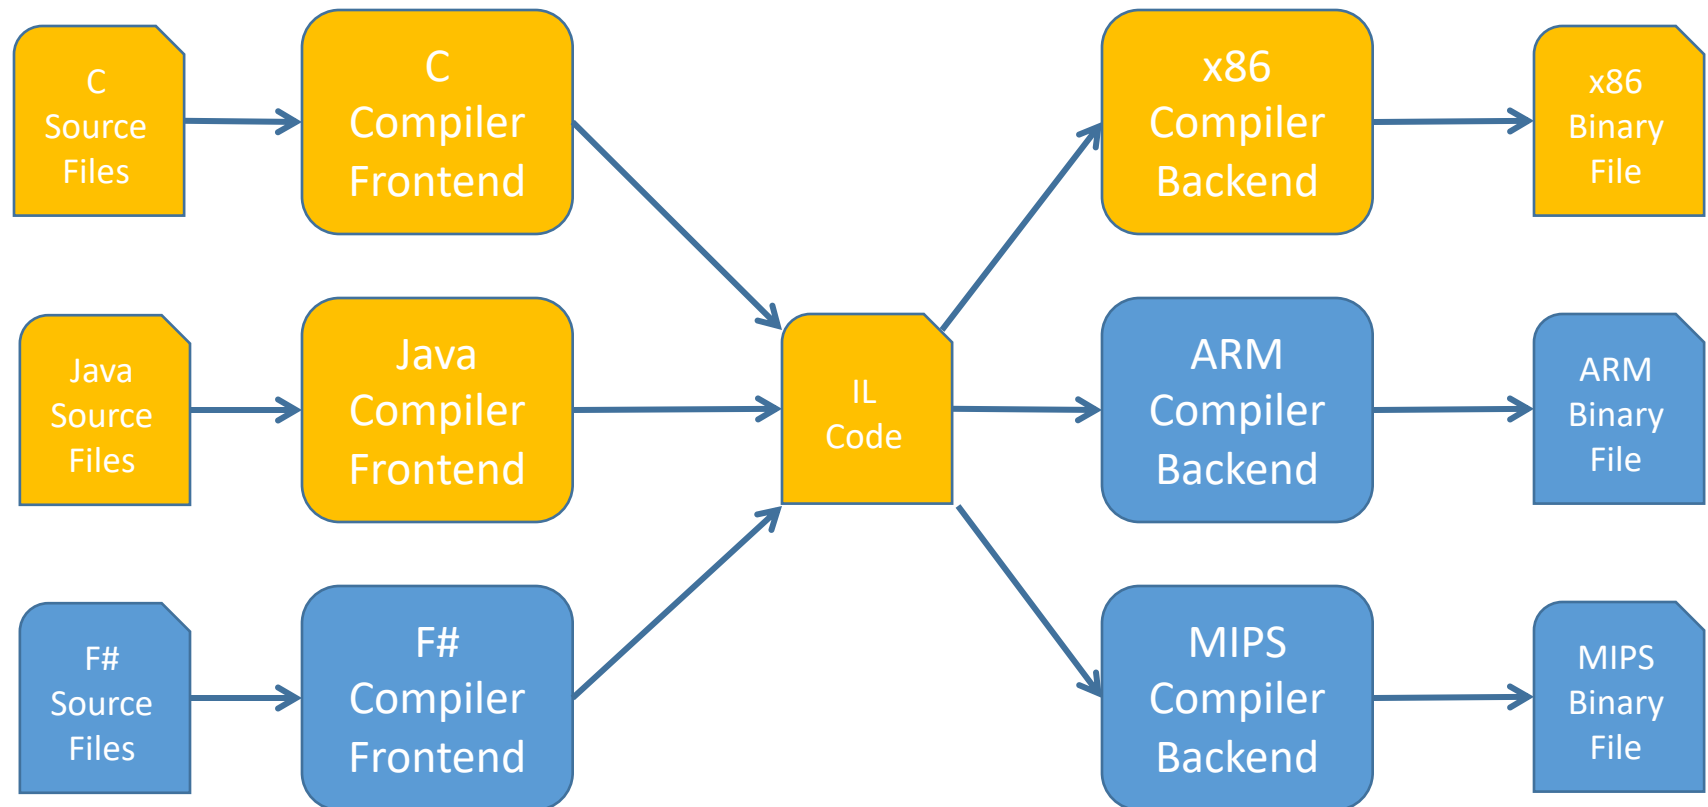
Building F# code for ARM target platforms





# Split Stage Compilation

(Via an Intermediate Language)

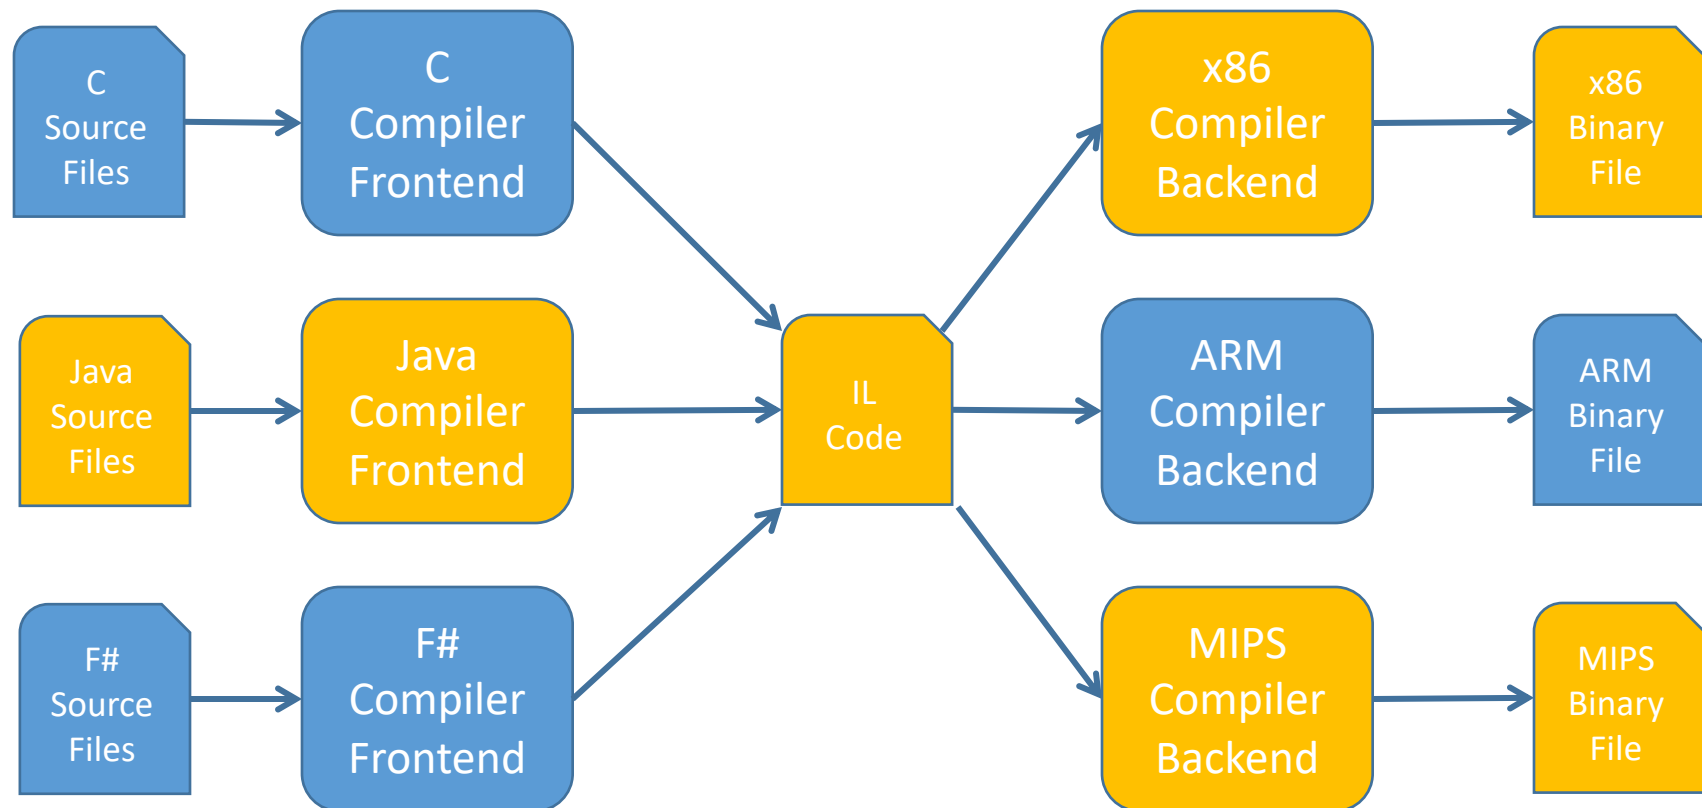


Building **mixed** C and Java code for x86 target platforms



# Split Stage Compilation

(Via an Intermediate Language)



Building Java code for x86 **and** MIPS target platforms



# Split Stage Compilation

(Via an Intermediate Language)

- This compiler design still has problems
  - Constraints in the IL language may lead to over-optimisation in the front-end
  - Constraints in the back-end may lead to overly verbose code for some targets
  - The lack of communication between the layers can lead to lost semantics between the two halves.
    - The parse tree is gone by the time we reach the back-end
    - No token metadata is available
- Linkage in *mixed* configurations may break down due to back-end workings being different
  - Call semantics may be different between MIPS and x86, for example



# Split Stage Compilation

(Via an Intermediate Language)

---

- The Intermediate Language may also constrain the types of targets that are possible (or reasonable) to create
  - Register-less machines are hard to handle with an IL that expresses things in terms of registers
  - Accumulator-based machines can result in overly verbose code, requiring more optimisation at the backend than normal
- But overall, the design provides a good degree of freedom, and prevents the internals of the compiler becoming too convoluted.
  - That said... look at the LLVM compiler code some time - it's not a fun place :|



# Just In Time Compilation

# Just in Time Compilation

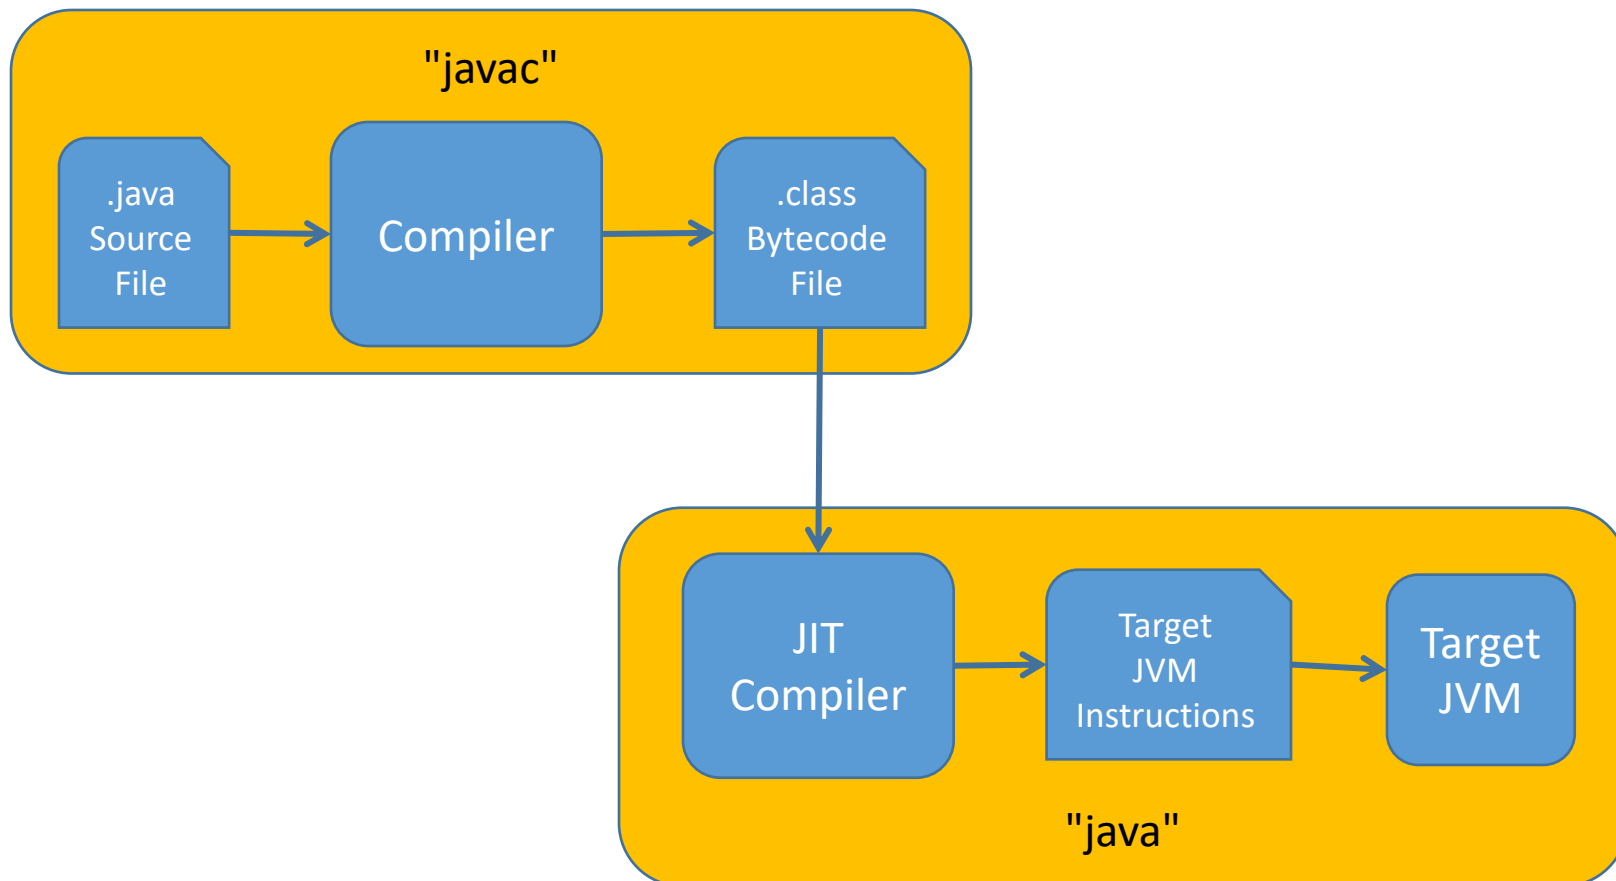
---

- Rather than Split-Stage compilers, a JIT compiler does its final stage as part of the execution environment.
  - The intermediate language becomes a merge of high-level and low-level opcodes
- The local high-level opcodes are then transcoded just before the runtime
  - This means we can do this for target-specific reasons
    - Runtime-environment patches
  - This also means we can effectively update the binary *after* the main compilation stage
    - Patches to bad behaviour, extensions, performance increases
    - Runtime flags can literally modify the code that actually runs



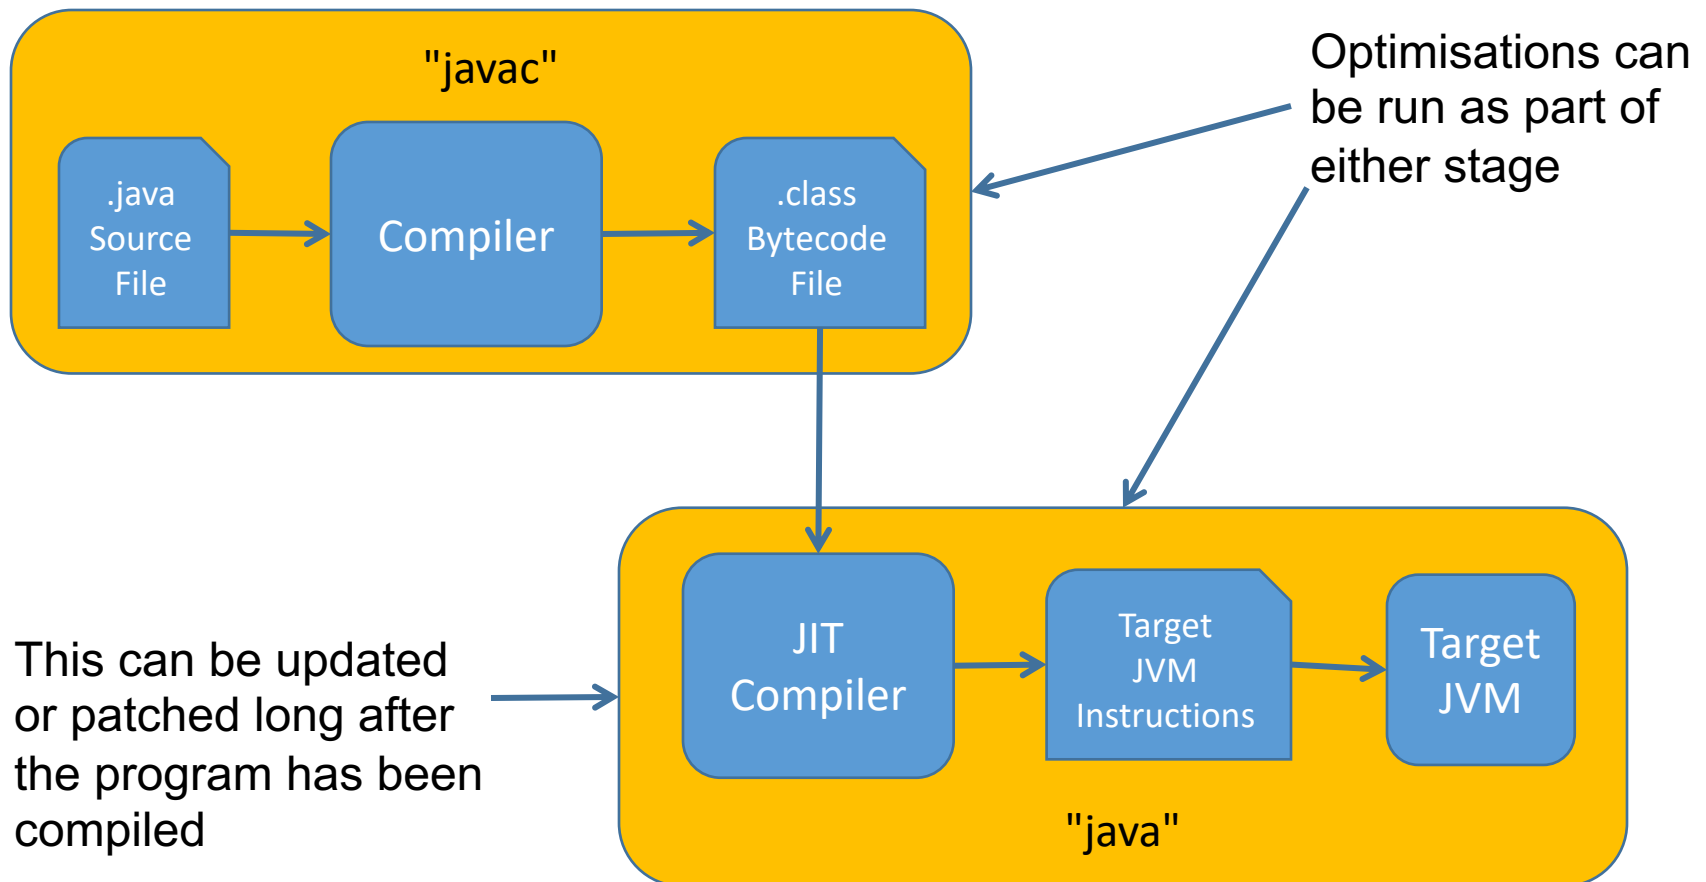
# Just In Time (JIT) Compilation + Interpreter

(Roughly how Java works)



# Just In Time (JIT) Compilation + Interpreter

(Roughly how Java works)





# JIT - Compiler:

## Java to Bytecode

---

- The compiler starts with the java source
  - We assume that this is from a method - otherwise it would fail in the syntax analyser!

```
outer:
for (int i = 2; i < 1000; i++) {
    for (int j = 2; j < i; j++) {
        if (i % j == 0)
            continue outer;
    }
    System.out.println (i);
}
```

- At this stage, all the usual steps are taken:
  - Lexical Analysis, Syntax Analysis, Semantic Analysis
  - But we finish with *Bytecode Generation*



# JIT - Compiler:

## Java to Bytecode

0: <code>iconst_2</code>	19: <code>ifne 25</code>
1: <code>istore_1</code>	22: <code>goto 38</code>
2: <code>iload_1</code>	25: <code>iinc 2, 1</code>
3: <code>sipush 1000</code>	28: <code>goto 11</code>
6: <code>if_icmpge 44</code>	31: <code>getstatic #84;</code>
9: <code>iconst_2</code>	34: <code>iload_1</code>
10: <code>istore_2</code>	35: <code>invokevirtual #85;</code>
11: <code>iload_2</code>	38: <code>iinc 1, 1</code>
12: <code>iload_1</code>	41: <code>goto 2</code>
13: <code>if_icmpge 31</code>	44: <code>return</code>
16: <code>iload_1</code>	
17: <code>iload_2</code>	
18: <code>irem</code>	

Method `java/io/PrintStream.println:(I)V`

Field `java/lang/System.out:Ljava/io/PrintStream;`



# JIT - Compiler:

## Java to Bytecode

---

- Most of the instructions correlate well with assembly-level instructions (direct translations)
  - `iload` -> `load`
  - `istore` -> `store`
  - `iconst` -> `literal integer constant value`
  - `goto` -> `jump`
- Others are completely alien to most architectures!
  - `getstatic` -> `load static table reference`
  - `invokevirtual` -> `invoke method from virtual table`



# Taking an LR(0) parser to an Interpreter

# Stack Machines as Interpreters

---

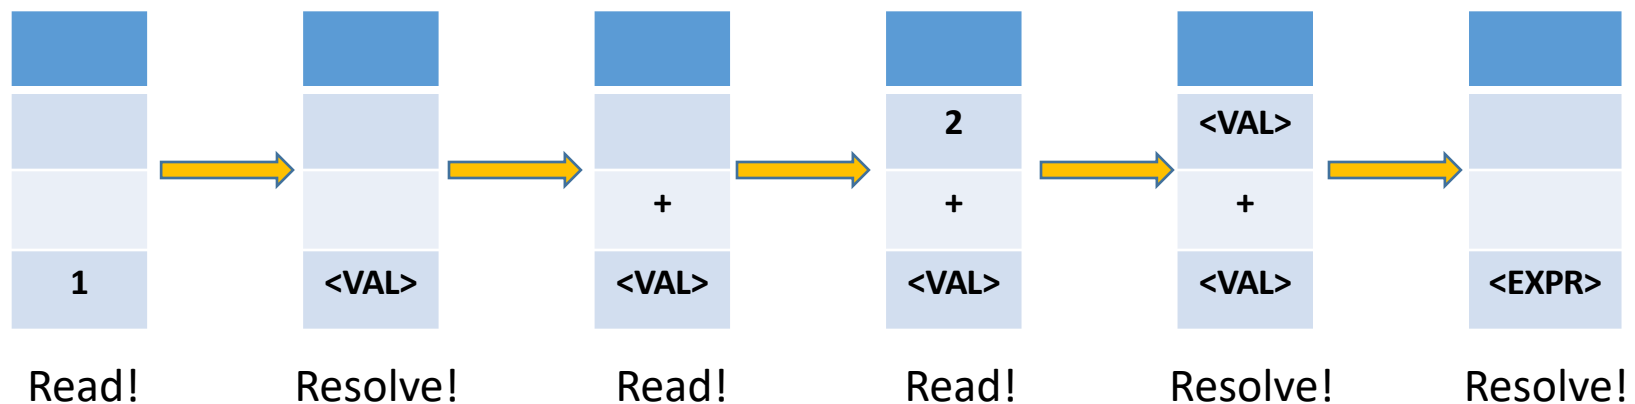
- Use a bottom-up parse strategy
  - Needs a language with an LR(0) compatible grammar
- Handle each operation as it resolves
- Sanction rules as they are processed
- Hope for no ambiguity or errors!
  - If the input is ambiguous, the interpreter is unable to proceed
    - ... most of the time

We have already seen how we can design our grammar to avoid this

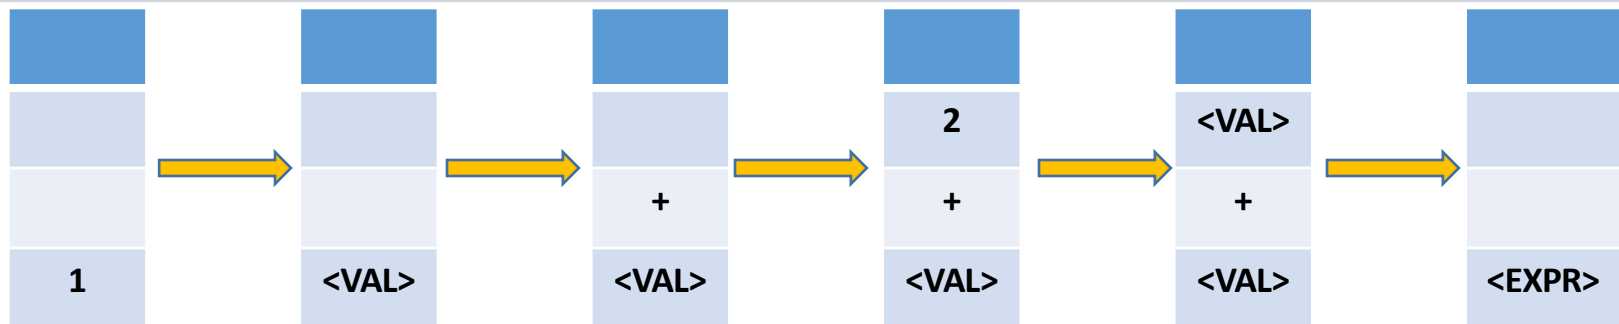


# Stack Machines as Interpreters


- Recall the way LR(0) parsers generate their tree
  - Push tokens onto a stack
  - Sanction rules by popping the components off the stack and replacing them with the non-terminal




# Stack Machines as Interpreters



- While processing the input stream, we can also generate opcodes on the fly.
  - We do need to keep some of the original metadata around on the stack to correctly handle this example - the values are lost in this representation.
  - Instead we keep the result of the rule sanctioning as well as the non-terminal



# Stack Machines as Interpreters



Note  
Our notation here is  
**<RULE : VALUE>**

Sanction

**<expr> := <val> + <val>**


AND execute:

**ADD 2, 1**





# Stack Machines as Interpreters



- We're also tracking which memory, registers, etc. we have in use
  - Register Table
  - Memory Map
- If interpreting, we do the operations **live**
  - If JIT-ing, we export new binary without the evaluation

Sanction

$\langle \text{expr} \rangle := \langle \text{val} \rangle + \langle \text{val} \rangle$

AND execute:

ADD 2, 1



# Processing Complete Expressions

$x = 1 + y$ ; LR(1) style

- LR(0) has some issues for interpreted languages
  - Without any lookahead, the grammar must have **no ambiguity**
  - Any ambiguous statements end up being processed in-order, which is probably not the order they should be handled.
- LR(1) allows us to fix this problem to a great extent
  - Is there more expression to evaluate?
  - Are there precedence modifiers in the input stream?
- Precedence Modifiers -> "(" and ")", for example.
  - We can use the presence of these to suppress the evaluation of the current top stack operation



# Processing Complete Expressions

$x = 1 + y$ ; LR(1) style

NEXT	Stack State --> Top						Symbol	Value
"x"							x	4
"="	SYM (x)						y	2
"1"	SYM (x)	ASSIGN						
"+"	SYM (x)	ASSIGN	LT (1)					
	Here LR(1) is useful to see if we have a completed expression ie. is the next token an EOL ';' ?							
"y"	SYM (x)	ASSIGN	LT (1)	OP (+)				
;	SYM (x)	ASSIGN	LT (1)	OP (+)	SYM (y)			
	Lookahead lets us know we have a complete expression. Now begin evaluating							
	SYM (x)	ASSIGN	LT (3)				ADD Op	
	SYM (x)						ASSIGN Op	



# Stack Machines as Interpreters

---

- Interpreter Problems
  - Forward declarations are an issue
    - If we haven't seen a function yet, but we're already calling it - how do we jump to it?
    - If we're careful we can build a partially compiled tree for evaluation later
      - Just grab the function blocks, save them for evaluation later
    - OR just enforce a declare-first language
      - Function headers always before calls
- But! We can combine the both JIT and Interpreter techniques for a better interpreter!
  - Use the Semantic Analyser to track scopes
  - Keep scope code we parse in memory ready for later
  - Build relations between scopes



# The END