


SCC311

# Designing Complex Distributed Systems: Google Infrastructure

Part 2

Onur Ascigil



# Key Design Philosophies

## ■ ***Simplicity***

- do one thing and do it well, avoiding feature-rich designs

## ■ ***Performance***

- ‘every millisecond counts’
- estimate performance through back-of-the-envelope calculations of primitive operations
  - e.g. memory/disk access, sending packets, locking/unlocking a mutex, etc.


## ■ ***Testing***

- ‘if it ain’t broke, you are not trying hard enough’
- complemented by a strong emphasis on *logging* and *tracing* to detect and resolve faults

## ■ Not adhering to any convention/school of thought. Whatever works best.

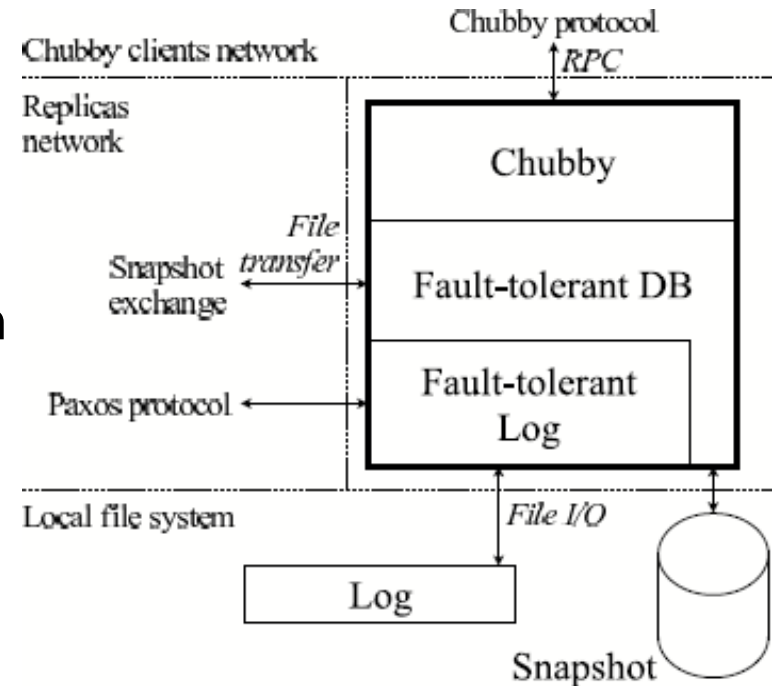
## ■ Use cheapest commodity hardware (risk analysis).

# Distributed Systems Infrastructure



# Chubby

- A filesystem to store small files and locks
- **Primary focus:** reliability and availability to a moderately large set of clients
  - Throughput and storage capacity were considered secondary
- Supports coarse grained synchronisation
  - Typical use case: locks that are acquired for long term (minutes or more)
- Example uses
  - ➔ To achieve distributed coordination
  - ➔ As a file system to store small files
    - e.g., for storing meta-data



# Chubby

- *Initially*: each file is a lock
- *Then*: associate a (small) data object with each lock
- **Atomic** read and write operations that are designed for small files
  - whole-file operations (open close, delete, ...) to discourage storing large files
  - no moving of files
  - minimal metadata

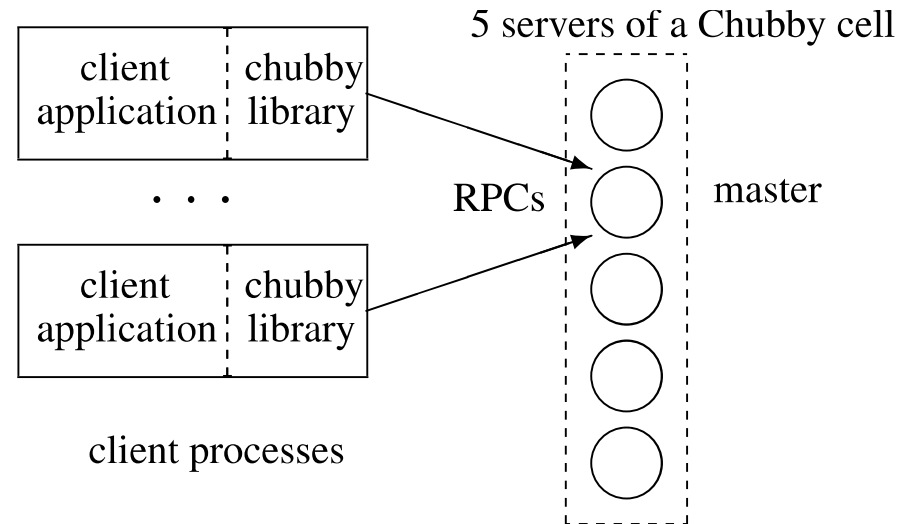
# Chubby API

<i>Role</i>	<i>Operation</i>	<i>Effect</i>
General	<i>Open</i>	Opens a given named file or directory and returns a handle
	<i>Close</i>	Closes the file associated with the handle
	<i>Delete</i>	Deletes the file or directory
File	<i>GetContentsAndStat</i>	Returns (atomically) the whole file contents and metadata associated with the file
	<i>GetStat</i>	Returns just the metadata
	<i>ReadDir</i>	Returns the contents of a directory – that is, the names and metadata of any children
	<i>SetContents</i>	Writes the whole contents of a file (atomically)
Lock	<i>SetACL</i>	Writes new access control list information
	<i>Acquire</i>	Acquires a lock on a file
	<i>TryAquire</i>	Tries to acquire a lock on a file
	<i>Release</i>	Releases a lock

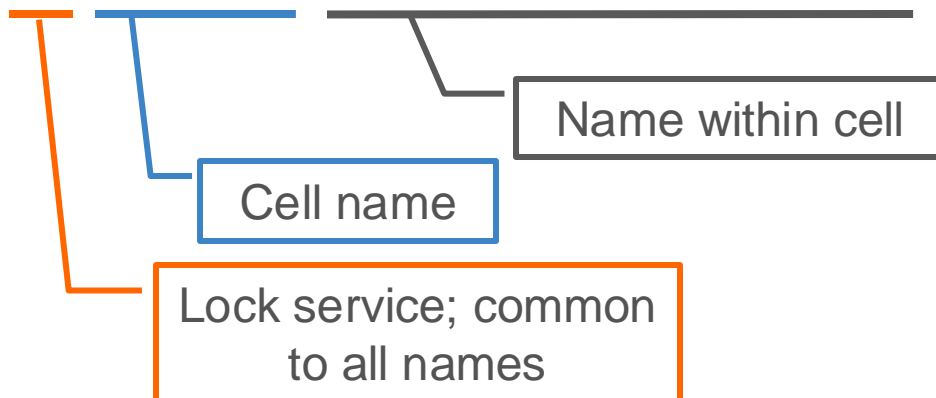
# Chubby

## ■ Simple directory structure

- root=cells
- leaves=files

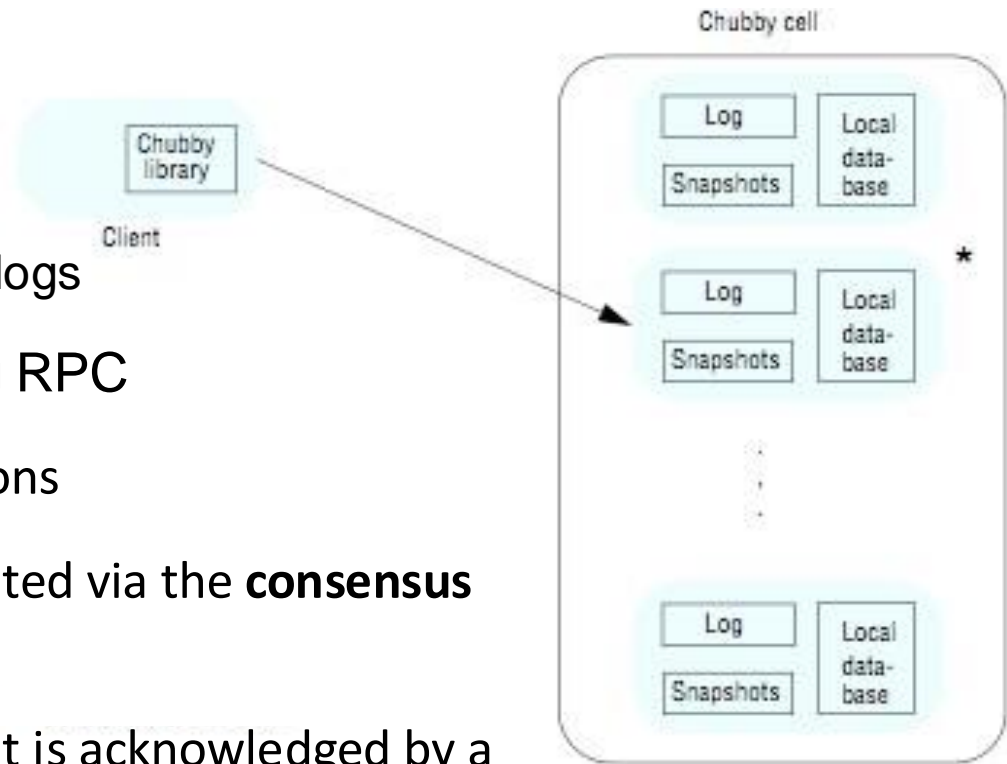


/ls/chubby\_cell/directory\_name/.../file\_name




# Replication in Chubby

- Each cell consists of replicas (default=5)
- Each replica is made up of:
  - a DB of directories and files
  - logs to ensure consistency
  - snapshots to reduce size of logs
- Clients connect to cells using RPC
- Atomic Read and Write operations
  - Write requests are propagated via the **consensus protocol** to all replicas.
  - Writes are accepted when it is acknowledged by a majority of the replicas in the cell.
  - Read requests are satisfied by the master alone





# Caching and Event Subscription



- Clients connect to cells using RPC
  - Sessions maintained using *KeepAlive* RPC calls
- Client caching of files is used to improve performance
  - Consistency maintained by invalidating all caches (piggybacked on *KeepAlives*) during which a write call is blocked
- **Subscription:** an application can register for event notification such as modification on a file (notification sending is piggybacked on *KeepAlives*)

# Use-cases

- Name Service
  - A very simple protocol for cache consistency with deterministic semantics
- Important for client services
  - e.g. BigTable stores Access Control Lists (ACLs)
- Master replica election




# Name Service using Chubby

- Applications run jobs involving **thousands of processes**.
  - Each process to communicate with every other leads to a quadratic number of DNS lookups.
- DNS supports caching with Time-to-Live (TTL)
  - A cached entries is discarded when its TTL expires
- A small TTL value is desirable for Google for prompt replacement of failed services
  - However, a small TTL can lead to a massive number of requests at the Google's DNS server.
  - Example: 3,000 clients would require about 150,000 lookups per second (estimate by Google)
- The caching semantics provided by Chubby invalidates cached entries pro-actively and avoids large number of periodic DNS resolution requests.

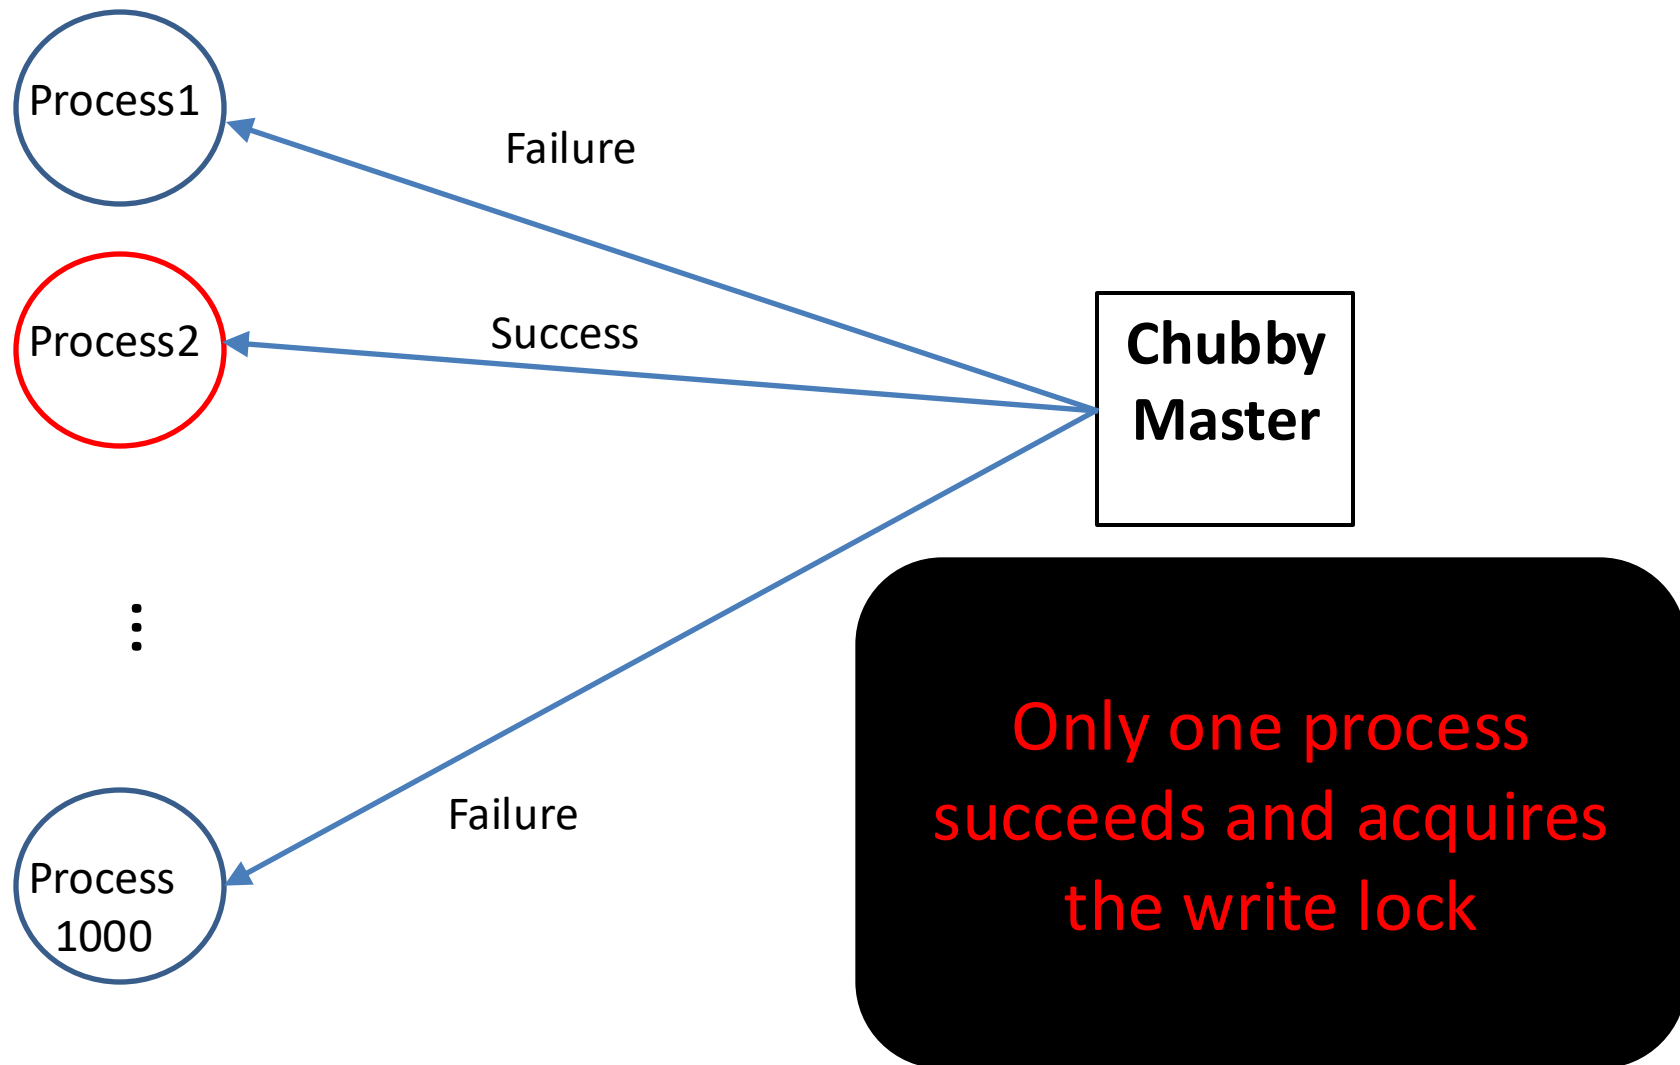
# Master Election using Chubby

- Both GFS and BigTable use Chubby to elect a master
- Processes can use the Chubby API to open a lock file and attempt to acquire a write lock.
- Only one succeeds and becomes the master.
- The others act as replicas.
- The primary then writes its identity to the lock file with `SetContents()`.
- The other replicas can call `GetContentsAndStat()` on the lock file (perhaps in response to a file modification event) to find out who the primary is.

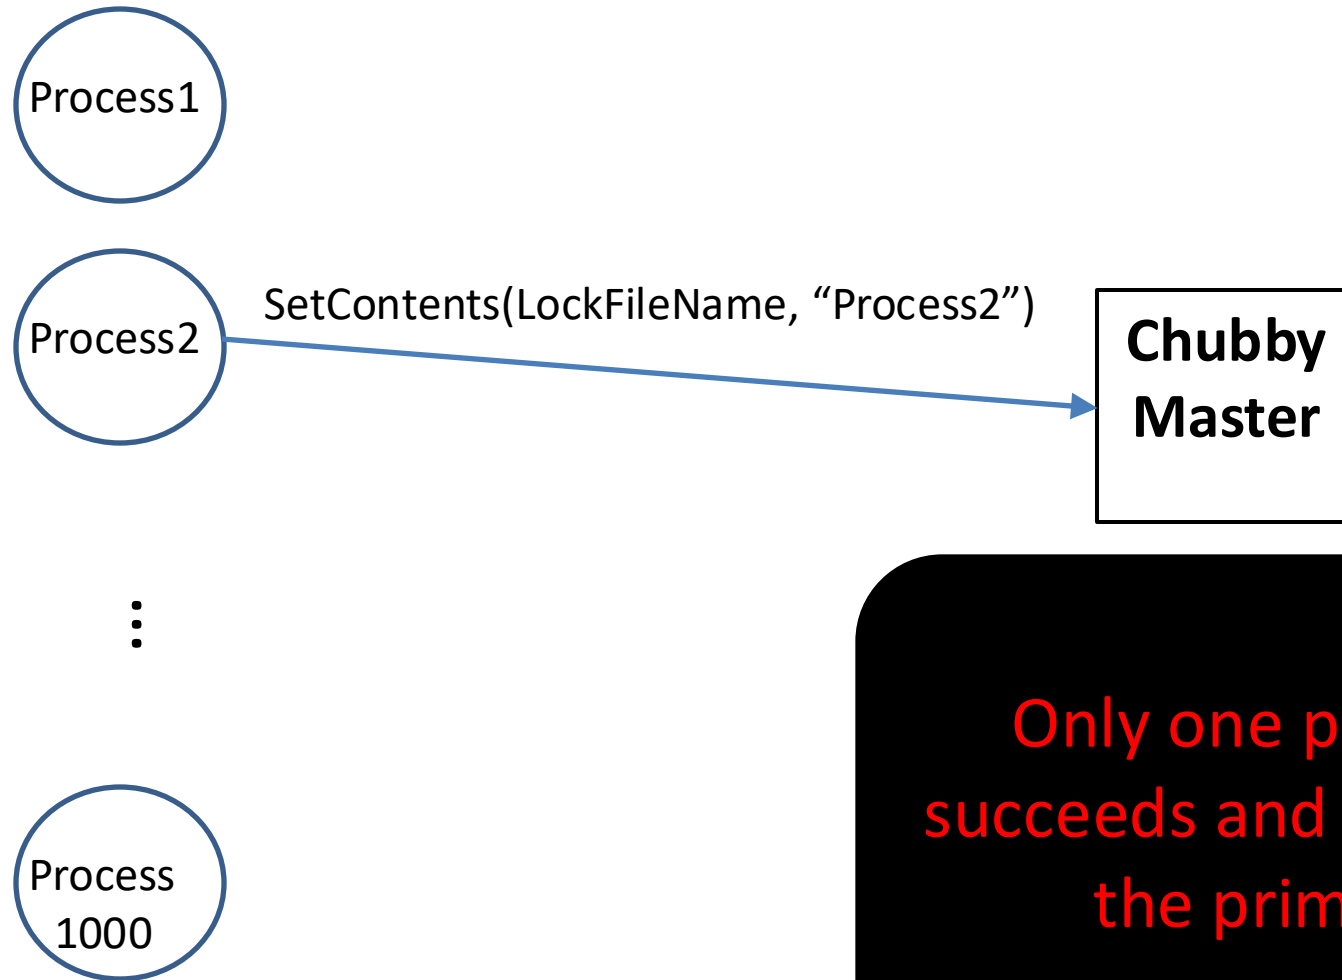
# Primary Election using Chubby



# Primary Election using Chubby

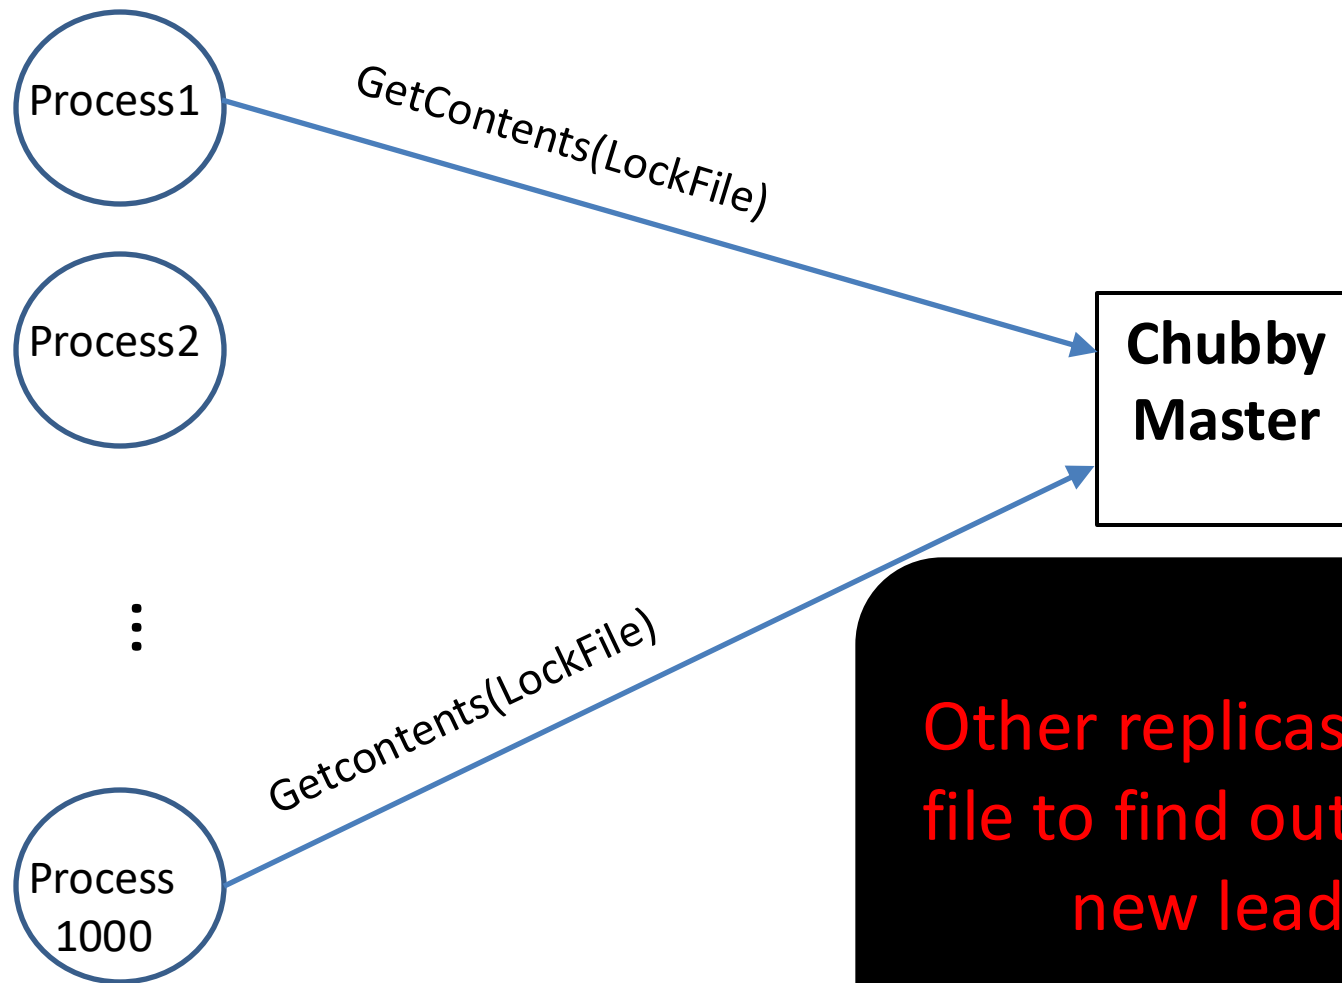


# Primary Election using Chubby



Only one process  
succeeds and becomes  
the primary

# Primary Election using Chubby



Other replicas read the file to find out who the new leader is



# Primary Election using Chubby


- **Main advantage:** by using Chubby to elect a leader, other applications and services do not have to implement their own distributed consensus library.
- Implementing a master (i.e., primary) election requires only a couple of RPC calls
- GFS and BigTable both use Chubby to elect a master.

# Chubby's use of Paxos


- Any node in a distributed system can fail, including a leader
  - *Paxos, a distributed consensus protocol for asynchronous systems, can be used for implementing a flexible election process (See: Week 5, Lecture 2)*
- **Step 1: Electing a leader**
  - *Leaders are ordered using sequence numbers*
  - *Each node maintains the highest sequence number seen so far*
  - *If bidding to be leader, a node will pick a higher unique number, and broadcasts this to all other nodes [Propose]*
  - *On receiving a Propose message, other nodes will reply with a [Promise] message if they have not seen a higher bidder*
  - *If a majority of promise messages is received (a quorum is achieved), a bidding node knows it has mandate to start acting as a leader*

# Chubby's version of Paxos


## Step 1: Electing a leader



## Step 2: Seeking a Consensus



## Step 3: Achieving Consensus



# Leader Election in Chubby

- The replicas use Paxos to elect a master
- A lease-based leadership
- The master must obtain votes from a majority of the replicas and promises that those replicas will not elect a different master for an interval of a few seconds known as the master lease
- The master lease is periodically renewed by the replicas provided the master continues to win a majority of the vote.

# Bigtable




- GFS and Chubby solve file storage and lock systems
  - What about large indexed datasets?
  - Relational DBs are too slow and too complex
- The Google way:
  - Strip away all performance-deteriorating features of RDBMS
  - Offer a minimal interface to keep things simple
- Bigtable!
  - A **very large-scale distributed table** built on top of GFS and Chubby
  - Supports **semi-structured / structured data**
    - without full relational operators, e.g. join
  - Able to scale to billions of rows and thousands of columns
  - Used by over 60 Google products including Search, Maps, Analytics

# Bigtable



- Table is indexed by **row, column and timestamp**
- Table split into fixed size **tablets** optimised for GFS
  - Subsequences of rows map onto tablets, which are the unit of distribution and placement
- Same **client, leader, worker pattern** as with GFS (workers manage tablets) but also with select (client) caching for tablet location
- 'Cheap' optimisations:
  - Rows in a table are lexicographically ordered
  - Columns are accessed using "family:qualifier"
  - Considerably simpler than RDBMS: No support for complex relational operations such as joins, unions, intersections, etc.
- Often used with MapReduce for computations on large-scale datasets
  - e.g. Web search:
    - row\_key=URL, columns=webpage\_attributes, timestamp=when\_captured

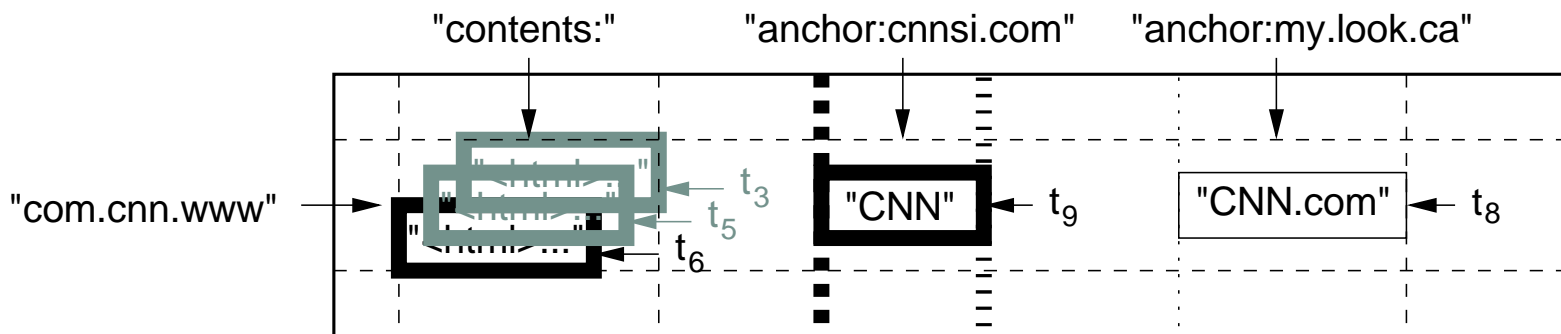
# Bigtable (continued)



**Bigtable maintains a lexicographic ordering of a given table by row key**

# Example use case: Storing Web pages

- It is common within Google to process information about web pages
- The row name is a reversed URL
- The **contents** column family contains the page contents, and the **anchor** column family contains the text of any anchors that reference the page






# BigTable API

Bigtable supports an API that provides a wide range of operations, including:

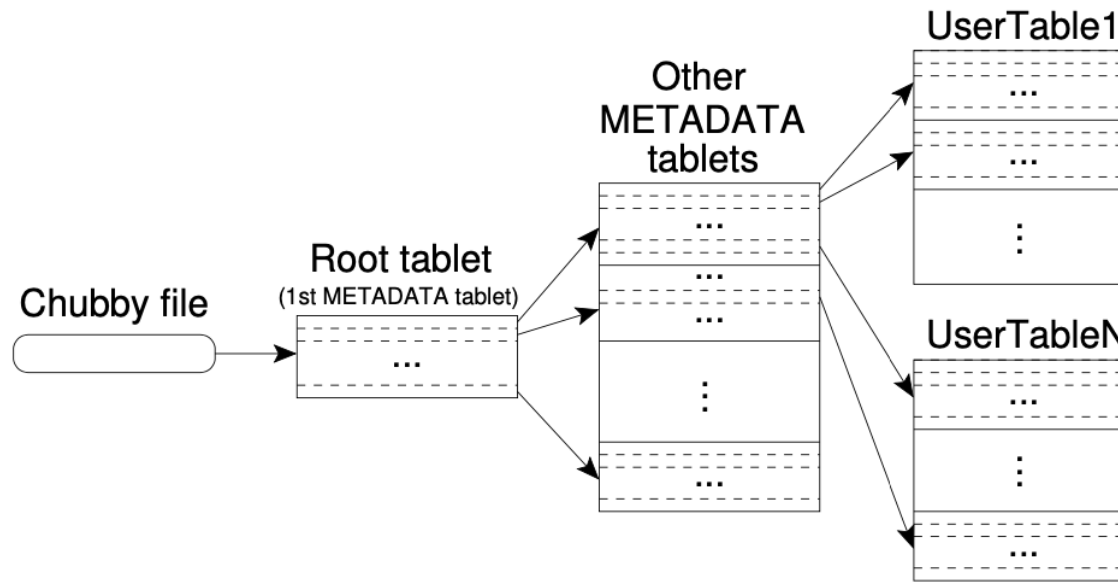
- The creation and deletion of tables
- The creation and deletion of column families within tables;
- Accessing data from given rows;
- Writing or deleting cell values;
- Carrying out **atomic row mutations** including data accesses and associated write and delete operations (more global, cross-row transactions are not supported);
- Iterating over different column families, including the use of regular expressions to identify column ranges;
- Associating metadata such as access control information with tables and column families.

# BigTable Architecture



- A master server and a potentially large number of tablet servers.
- A master mainly does two things:
  - Monitoring the status of tablet servers
  - Ensuring effective load balancing
- Clients communicate directly with the tablet servers (each holding a tablet, i.e., a range of row keys) for reads and writes.

# Tablet Index



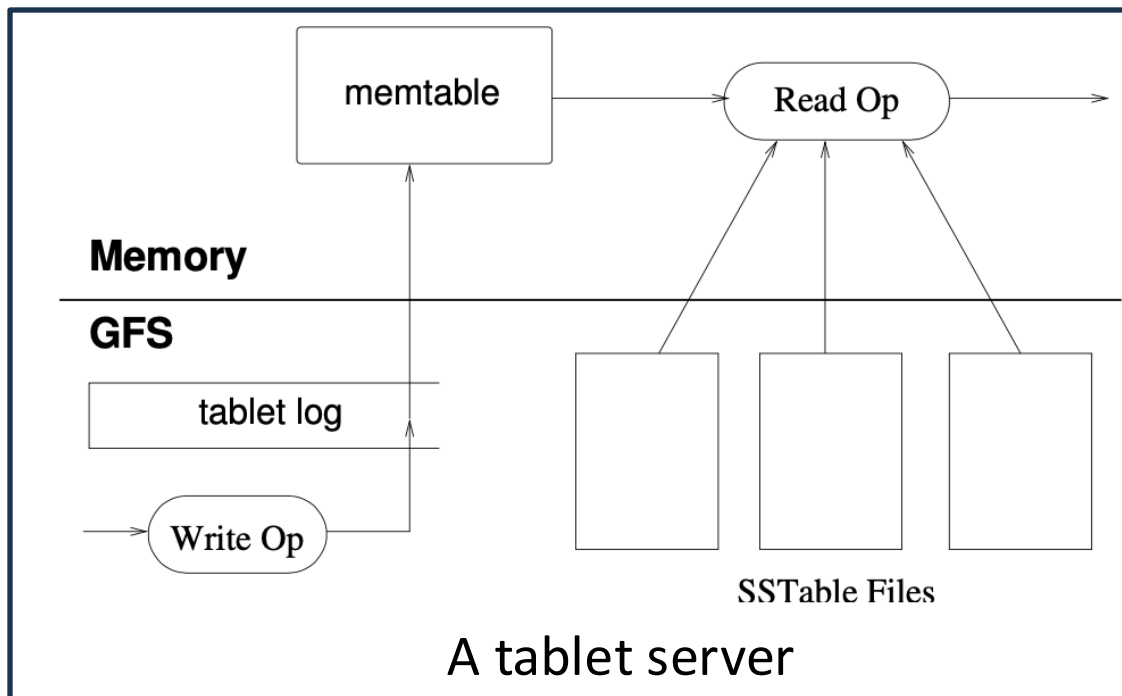
- Root tablet contains pointers to METADATA tablets, and each METADATA tablet, in turn, stores the location of user tablets by encoding their start and end row keys.
- Clients first consult a Root metadata table to find out which tablet server stores the metadata for the row key range they need.
- Once the client knows which tablet server holds the desired tablet, it sends requests directly to that server (a similar design to GFS).

# SSTable(1)

- The tablets are stored as one or more **SSTables** (Sorted Strings Table) in GFS: An ordered, **immutable** map from keys to values.

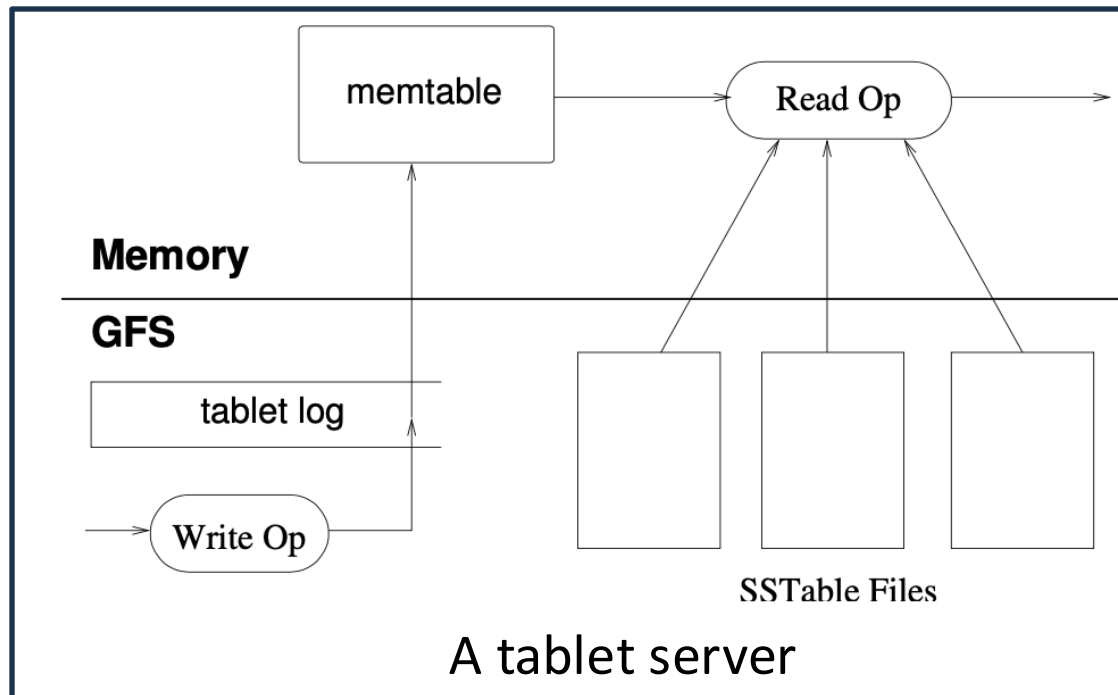
Updates to tablets are batched and applied periodically:

- **1. Tablet Log:** Updates are first recorded in a **commit log** (write-ahead log) in GFS to ensure recoverability.



# SSTable(2)

- **2. Memtable:** The **most recent updates** are stored in the **memtable**: an in-memory sorted structure.
- The memtable stores the latest committed data to enable quick access to recent changes before they are written to GFS.



# BigTable – Compactions

- Memtable & Minor Compaction
  - Updates stored in memtable; grows with writes.
  - When full, memtable is frozen and converted to an SSTable in GFS.
  - Reduces memory use and minimizes commit log data needed for recovery.
  - Reads/writes continue during compaction.
- Merging Compaction
  - Periodically merges multiple SSTables and the current memtable.
  - Reduces the number of SSTables to optimize read performance.

# BigTable - Monitoring

- BigTable uses Chubby for monitoring in a rather interesting way
- Maintains a directory in Chubby containing files representing each of the available tablet servers.
- When a new tablet server comes along, it creates a new file in this directory and more importantly, it obtains an **exclusive lock** on this file.
- Masters periodically check the status of the tablet server locks and if a lock is lost by a server, then master can infer a problem with the server.
- The tablet server can also surrender its lock if the machine is needed (reassigned) for a different purpose.

# Summary of Design Choices

<i>Element</i>	<i>Design choice</i>	<i>Rationale</i>	<i>Trade-offs</i>
<i>Chubby</i>	Combined lock and file abstraction	Multipurpose, for example supporting elections	Need to understand and differentiate between different facets
	Whole-file reading and writing	Very efficient for small files	Inappropriate for large files
	Client caching with strict consistency	Deterministic semantics	Overhead of maintaining strict consistency
<i>Bigtable</i>	The use of a table abstraction	Supports structured data efficiently	Less expressive than a relational database
	The use of a centralized master	As above, master has a global view; simpler to implement	Single point of failure; possible bottleneck
	Separation of control and data flows	High-performance data access with minimal master involvement	-
	Emphasis on monitoring and load balancing	Ability to support very large numbers of parallel clients	Overhead associated with maintaining global states



# Some Cloud Reflection

- Google is much more than just a search engine.
- Similar to Amazon, the cloud was a side-effect of building applications on a massive scale.
- This “side-effect” infrastructure has allowed them to build more applications that were previously not possible.
- This also opened the door for others (as cloud customers).

# Additional Reading

- CDKB: Chapter 21
- Lots of material at Google Labs:
  - The Anatomy of a Large-Scale Hypertextual Web Search Engine
  - The Google File System
  - The Chubby Lock Service for Loosely-Coupled Distributed Systems
  - Paxos Made Live – An Engineering Perspective
  - MapReduce: Simplified Data Processing on Large Clusters
  - Bigtable: A Distributed Storage System for Structured Data
  - Interpreting the Data: Parallel Analysis with Sawzall

<http://research.google.com/pubs/papers.html>

# Expected Learning Outcomes

## At the end of these sessions:

- You should have an appreciation of the challenges of **web search** and the provision of **cloud services**
- You should understand the **problem of scalability** as it applies to the above problems and the trade-offs with other dimensions (reliability, performance, openness)
- You should understand the architecture of **Google Infrastructure** and how such a **middleware solution** can support the business
- You should also be aware of the specific techniques used within Google relating to:
  - ProtocolBuffers
  - GFS & MapReduce
  - Chubby
  - BigTable(and why they are designed the way they are and how this contributes to delivering the goals of Google)