

Interrupts

Dr Andrew Scott


a.scott@lancaster.ac.uk

1

Polling or Busy Waiting

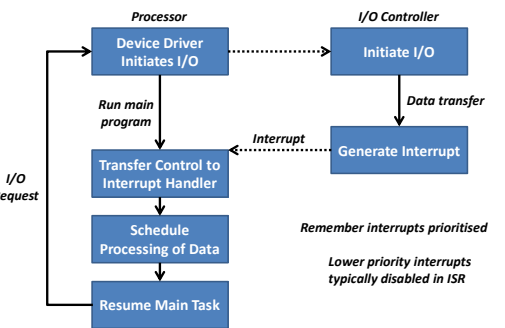
- Code ‘spins’ on busy flag* awaiting status change
 - Very wasteful of processor time
- Very difficult to interleave useful code
 - Need automated mechanism
- Also known as
Synchronous I/O

* Remember: a flag is just a status bit on device
e.g. data available, operation complete, ...



2

Interrupt Driven/ Asynchronous I/O



Remember interrupts prioritised
Lower priority interrupts typically disabled in ISR

3

Interrupt Service Routine (ISR)

Note we push a copy of all registers so C handler has access to all system state

Some interrupt systems by default

- Disable interrupts - typically do
- Push all registers - easy but slow
- Push few, if any - fast

Interrupt

Memory

Interrupt Vector (holds address of ISR)

0x0008

0x2000

Address 0x2000

One of the few places we need to use assembler

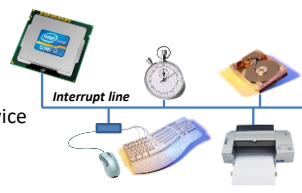
```
isr:  pushal  # all registers
      call  C_handler
      popal  # cf. pushal
      sti   # re-enable interrupts*
      iret  # restart main program
```

* iret instruction on many processors re-enables interrupts, so wouldn't need sti

4

Interrupts

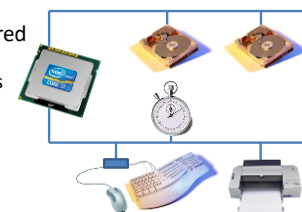
- If devices share interrupt/ interrupt line
 - Know at least one device changed status
 - ...better, but no idea which (generated interrupt)
 - Back to polling
- After interrupt
 - Must poll each device for data ready/ state change



5

Multiple Interrupt Lines

- Narrows search when subsequently polling
- Groups similar devices
- More support required on processor
 - Multiple lines/ wires
 - Cost, complexity, ...
 - Multiple ISRs



6

Multiple Interrupt Service Routines

Interrupt Vectors

Address	ISR Routine
0x2000	isr: push \$1 ... iret
0x2020	isr: push \$2 ... iret
0x2040	isr: push \$3 # the value 3 pushal # all registers call C_handler # Our C code popal # cf. pushal addl \$8,%esp # 'pop' \$3 sti iret

Note:
Instead of having a separate C function for each interrupt we can *push* a number identifying the specific interrupt, giving:
`C_handler(3, ...)`
Also note that we must *remove/pop* this number once we've finished ...the *addl*

7

Pushal Instruction and C Handler

- Pushes x86 registers in following order

```
struct cpu_registers {  
    uint32_t edi;  
    uint32_t esi;  
    uint32_t ebp;  
    uint32_t esp;  
    uint32_t ebx;  
    uint32_t edx;  
    uint32_t ecx;  
    uint32_t eax;  
};
```

ISR: `push $3`
`pushal`
`call C_handler`
`...`

C Handler:
`void C_handler(// Interrupt handler, called from assembler ISR`
`struct cpu_registers regs, uint32_t irq,`
`uint32_t instr_ptr, uint32_t cs, uint32_t cpu_flags`
`) {`

`Always pushed by x86 processor`

`}`

Interrupt number we pushed (\$3)

8

Programmable Interrupt Controller

- Processor can't have separate wire to every device
 - PC now uses Advanced PIC (Intel 82093 APIC)
- Gives priority to most important devices
 - Orders generated interrupts

PIC 1

- IRQ0: Programmable Timer
- IRQ1: Keyboard
- IRQ3: Serial/ COM2
- IRQ4: Serial/ COM1
- IRQ5: LPT2 Printer
- IRQ6: Floppy Disk
- IRQ7: LPT1 Printer/ Spurious

PIC 2

- IRQ8: Real Time Clock (RTC)
- IRQ9: Not used
- IRQ10: Not used
- IRQ11: Not used
- IRQ12: PS/2 mouse
- IRQ13: Floating-Point exception
- IRQ14: ATA Hard Disk
- IRQ15: ATA Hard Disk

10

NT Interrupt Request (IRQ) Levels


- Interrupts often disabled on entry to ISR
 - But generally allow high-priority IRQ to interrupt handling of lower-priority IRQ

IRQL	
31	High
30	Power Failure
29	Inter-processor signalling
28	Clock
	Device <i>n</i>
	Device 1
2	Dispatch/ Deferred Procedure Call (DPC) software interrupts
1	Asynchronous Procedure Call level (APC) software interrupts
0	Low (normal thread execution)

13

Nested Interrupts

- Interrupts may themselves be interrupted
 - Processing becomes nested → complicates code + timing
 - Typically group interrupts into priority levels (IRQL)
 - Defer processing of interrupts with lower *Int. Request Level* (IRQL)



14

Generalising the Mechanism

- Interrupts
 - Interrupt Request signal
 - Generated by device attached to specific wire/ line
 - Message based
 - Requested in 'message' over intelligent bus, e.g. PCI Bus
- Exceptions (*can be generic term encompassing others*)
 - Generated in response to internal processor problem
- Software Interrupts
 - Initiated by executing special Interrupt/ Trap instruction

15

```
const char *
x86Exception[ ] = {
    "Division by zero",           // 0 F -
    "Single step",               // 1 FT -
    "Non-Maskable Interrupt (NMI)", // 2 I -
    "Breakpoint",                // 3 T -
    "INTO executed with overflow flag set", // 4 T -
    "BOUND executed with index outside array bounds", // 5 F -
    "Invalid Op-code",           // 6 F -
    "Device/ Co-processor not available", // 7 F -
    "Double fault - no handler for exception", // 8 A 0
    "Co-processor Segment overrun (pre i486)", // 9 F -
    "Invalid Task State Segment (TSS)", // 10 F E
    "Segment not present",        // 11 F E
    "Stack/ Segment Fault",       // 12 F E
    "General Protection Fault (GPF)", // 13 F E
    "Page Fault",                // 14 F E
    "Exception 0x0F - reserved by Intel", // 15 -
    "x87 Floating Point exception", // 16 F -
    "Alignment check (Ring 3: user mode)", // 17 F 0
    "Machine check (internal processor error)", // 18 A ?
    "SIMD Floating Point exception", // 19 F -
}
```

x86 Exception List

16

When Exceptions go Wrong

- Interrupts MUST be correctly handled
- If no interrupt service routine set for interrupt
 - Processor raises **Double Fault** interrupt
 - Unknown state can be hard to recover from
 - Blue Screen of Death
- If double fault isn't handled
 - Processor **Triple Faults**
 - Causes machine restart

17
