



SCC.211 Operating Systems


Lecture 4 – Semaphores

Amit Chopra


School of Computing & Communications, Lancaster University, UK

Objectives


- Examples
- Definition
- Difference between locks and semaphores




Example: Vaccination by Locking




**Lock allows only one in at a time:
underutilization of booths**




Example: Vaccination by Semaphore




Vaccination by Semaphores



Vaccination by Semaphores



Vaccination by Semaphores




```
public class Vaccination {  
    public static void main(String[] args) {  
        Semaphore v = new Semaphore(3);  
        Thread[] threads = new Thread[20];  
        for(Thread th: threads) {  
            th = new Thread(new Person(v));  
            th.start();  
        }  
    }  
}
```

```
class Person implements Runnable {
```

```
    Semaphore v;
```

```
    Person(Semaphore v) {
```

```
        this.v=v;
```

```
    }
```

```
    public void run(){
```

```
        v.acquire();
```

```
        System.out.println("Getting vaccinated!");
```

```
        v.release();
```

```
    }
```

```
}
```

(Omitted exception handling)

Selling tickets via Semaphore

```
public class Vaccination
    public static void main(String[] args)
        Semaphore tickets = new Semaphore(10);

        Thread[] threads = new Thread[20];
        for(Thread th: threads)
            th = new Thread(new Buyer(tickets));
            th.start();
```

Omitted curly braces to save space

```
class Buyer implements Runnable
    Semaphore t;

    Buyer(Semaphore t)
        this.t=t;
```

Problem: Buyers who can't get a ticket blocked indefinitely

```
public void run()
    acquire();
    System.out.println(this + " got ticket 😊");
```

Selling tickets via Semaphore

```
public class Vaccination
    public static void main(String[] args)
        Semaphore tickets = new Semaphore(10);

        Thread[] threads = new Thread[20];
        for(Thread th: threads)
            th = new Thread(new Buyer(tickets));
            th.start();
```

```
class Buyer implements Runnable
    Semaphore t;
```

```
    Buyer(Semaphore t)
        this.t=t;
```

```
    public void run()
        if(t.tryAcquire())
            System.out.println(this + " got ticket 😊");
        else System.out.println(this + " out of luck 😞");
```

tryAcquire returns immediately, with a permit if available (returns true) and without if not (returns false).

What happens if


- Number of buyers and initial number of tickets are equal?
- There are zero tickets initially?
- Number of tickets is greater than the number of buyers?

Locks only provide mutual exclusion

- One thread executes in the critical section at a time

May want more threads using resources concurrently

- Produce/consumer problem
 - Need one thread to run after the other
 - Don't wish to operate in lockstep



A **synchronization primitive** implemented via a **counter** with a **nonnegative initial value**, and supports **two operations** that are **guaranteed to be atomic**:

acquire

If counter is positive, then decrement counter by 1 and let calling thread continue. If counter is 0, block the thread *on the semaphore*.

release

Increment counter by 1. If there are blocked threads on the semaphore, unblock some such thread.

acquire aka **wait** aka **P** aka **down**

release aka **signal** aka **V** or **up**

Note: Counter value is never negative!

```
public class Bank_account {  
    private int bal = 0;  
    private Semaphore mutex = new Semaphore(1)  
  
    public void Bank_account(int start_balance) {  
        bal = start_balance;  
    }  
    public void update(int amount) {  
        mutex.acquire();  
        bal = bal + amount;  
        mutex.release();  
    }  
}  
...  
Bank_account b = new Bank_account(0);
```

Standard initial value
of 1 is used for
mutual exclusion

‘Bracket’ critical
section with
acquire & release

How does this work?

1

Assume processes P1 and P2 arrive around the same time
(with P1 just ahead)

2

P1 calls **acquire**. Semaphore value is 1.
So value is decremented to 0 and P1 enters critical section without blocking

3

P2 calls **acquire**. Semaphore value is 0.
So P2 blocks on the semaphore

4

P1 calls **release**.
So value is incremented and P2 is unblocked

5

P2 runs whenever scheduled

```
public class Bank_account {  
    private int bal = 0;  
    private Semaphore sem = new Semaphore(3);
```

```
    public void Bank_account(int start_balance) {  
        bal = start_balance;  
    }
```

```
    public void update(int amount) {  
        sem.acquire();  
        bal = bal + amount;  
        sem.release();  
    }
```

Can now allow 3
threads to execute in
critical section

‘Bracket’ critical
section with
acquire & release calls

Incorrect because atomicity is violated!
Three threads potentially manipulating
balance at same time

```
...  
Bank_account b = new Bank_account(100);
```

Internally managed by thread system

The thread system records the threads that are blocked on a semaphore. E.g., it may place them in a queue for the semaphore.

When semaphore is incremented, the system picks a thread from the queue to run.

The thread runs whenever it is scheduled!

Each Java object **O** has an associated *monitor* **M**, with two conceptual components, a lock **L** and wait set **W**

- **L** ensures at most one thread in any code synchronized on **O**
 - Recall: *synchronized method() {body} in O* equiv. to.
method(){synchronized(O){body}}
- **W**: *maintains threads blocked on M.*
 - **O.wait()** suspends calling thread and adds it to **W**
 - Can only be called from synchronized blocks
 - Releases **L**
 - **O.notify()** takes a thread from **W** and makes it runnable.
 - Can only be called from synchronized blocks
 - Thread must obtain **L** before resuming
 - (**O.notifyall()** makes all waiting threads runnable)

Semaphores in Java (Potentially Buggy)

```
public class Semaphore {  
    private int count = 0;  
    public Semaphore(int init_val)  
    {  
        count = init_val;           // Should check it's >= 0  
    }  
    public synchronized void acquire()  
    {  
        if(count > 0) count = count--;  
        else wait();                //Go on this semaphore's wait queue  
    }  
    public synchronized void release()  
    {  
        count = count++;  
        notify();                    // Make a waiting thread runnable  
    }  
}
```

```
public class Semaphore {  
    private int count = 0;  
  
    public Semaphore(int init_val)  
    {  
        count = init_val;  
    }  
}
```

```
    public synchronized void acquire()  
    {  
        while() {  
            if (count > 0) {count = count--; break;}  
            else wait();  
        }  
    }
```

```
    public synchronized void release()  
    {  
        count = count + 1;  
        notify();  
    }  
}
```

Fixed: Recheck counter condition
because other threads might
have used acquired the semaphore
between waking up and proceeding

- Semaphores generalizes over the notion of mutual exclusion
 - Semaphore: n permits available for threads to acquire concurrently
 - If $n=1$, we get mutual exclusion
- Think of semaphore as a counter with two operations.
 - Acquire: If $\text{counter} > 0$, decrement counter, thread can acquire a permit. Else thread blocks.
 - Release: Increment counter. Unblock some thread.
- Can be implemented via Java's synchronization facilities
- Two Uses
 - Controlling access to reusable resources (vaccination booths)
 - Counting resources down (tickets)
- More complex uses coming!