Unix File-System

Dr Andrew Scott

a.scott@lancaster.ac.uk

1

Typical Unix *inode* Structure

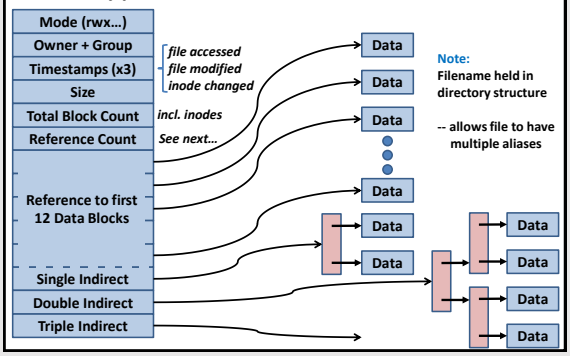

2

Reference Count and Hard Links

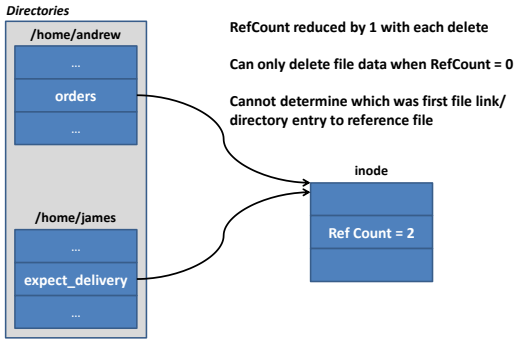

3

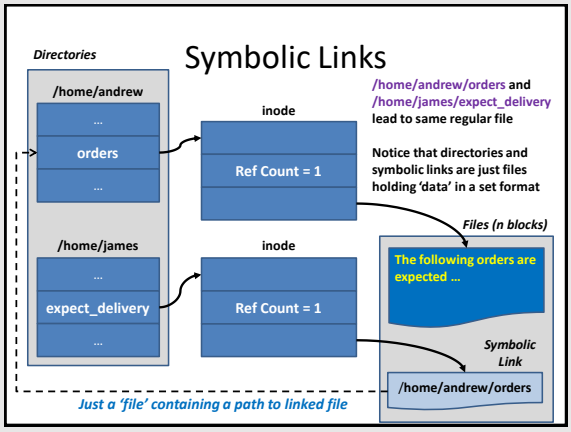
Symbolic Links
Directories
/home/andrew
...
orders
...
/home/james
...
expect_delivery
...
inode
Ref Count = 1
inode
Ref Count = 1
/home/andrew/orders and
/home/james/expect_delivery
lead to same regular file
Notice that directories and
symbolic links are just files
holding 'data' in a set format
Files (n blocks)
The following orders are
expected ...
Symbolic
Link
/home/andrew/orders
Just a 'file' containing a path to linked file

4

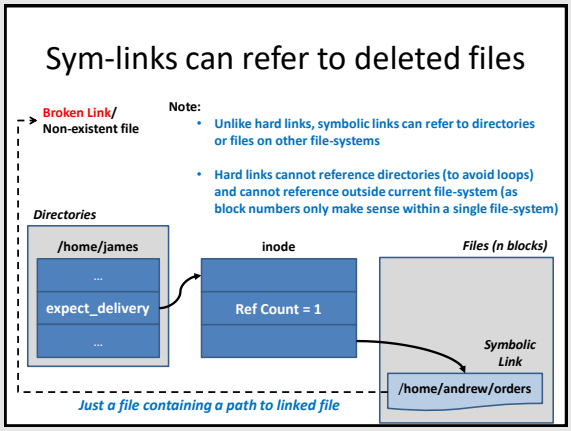
Sym-links can refer to deleted files
Broken Link/
Non-existent file
Note:
• Unlike hard links, symbolic links can refer to directories or files on other file-systems
• Hard links cannot reference directories (to avoid loops) and cannot reference outside current file-system (as block numbers only make sense within a single file-system)
Directories
/home/james
...
expect_delivery
...
inode
Ref Count = 1
Files (n blocks)
Symbolic
Link
/home/andrew/orders
Just a file containing a path to linked file

5

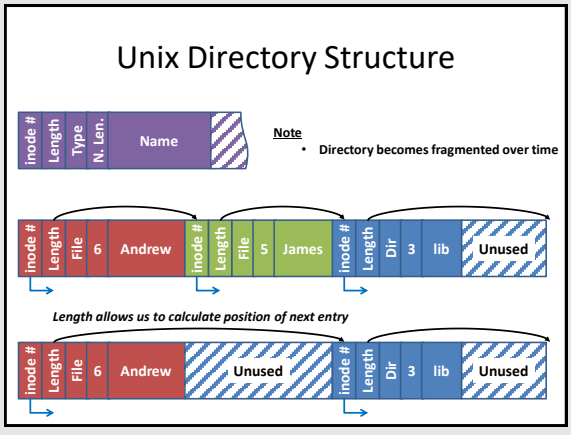
Unix Directory Structure
inode #
Length
Type
N. Len.
Name
Note
• Directory becomes fragmented over time
inode #
Length
File
6
Andrew
inode #
Length
File
5
James
inode #
Length
Dir
3
lib
Unused
Length allows us to calculate position of next entry
inode #
Length
File
6
Andrew
Unused
inode #
Length
Dir
3
lib
Unused

6

File Lookup: */home/acs/211*

- Often
  - Search from last directory entry read
    - Processes typically scan directories sequentially
  - Use filename caches to speed up access

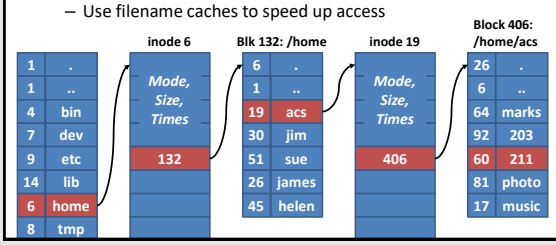

7

Block Size and Efficiency/ Waste

- Values taken off 1.2GB 4.3BSD Unix file-system
  - Wastage as percentage of disk space
  - Larger block sizes allow for larger files
    - More links per index block
    - More data in each referenced data block

Space Used (Mbyte)	%age Waste	Structure
775.2	0.0	(control value) Just the raw data
828.7	6.9	Data & inodes with **512** byte blocks
866.5	11.8	Data & inodes with **1024** byte blocks
948.5	22.4	Data & inodes with **2048** byte blocks
1128.3	45.6	Data & inodes with **4096** byte blocks

8

Block Fragments

- Efficiency best when block and file sizes matched, but
  - Small blocks tend to suit small files
  - Large blocks tend to suit large files
- Tended to move to larger blocks as file sizes increased
  - Media files, etc.
- Some file-systems can allocate fragments of blocks
  - i.e. 4 x 1K fragments from a single 4K disk block
- Blocks must be multiple of disk sector size

9

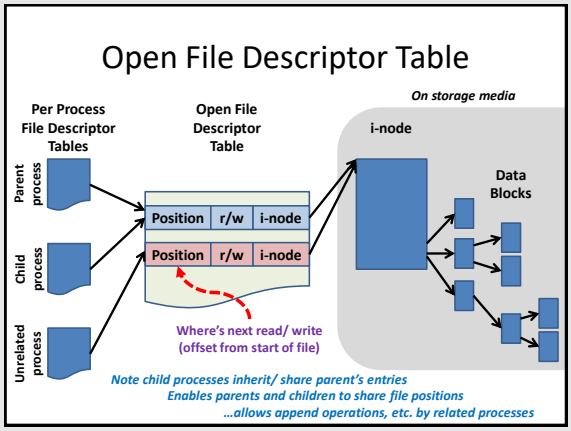
Open File Descriptor Table
On storage media
Per Process File Descriptor Tables
Open File Descriptor Table
i-node
Data Blocks
Parent process
Child process
Unrelated process
Position
r/w
i-node
Position
r/w
i-node
Where's next read/ write (offset from start of file)
Note child processes inherit/ share parent's entries
Enables parents and children to share file positions
...allows append operations, etc. by related processes

10