

Part II

COMPUTING AND COMMUNICATIONS [2 Hours]

SCC.211 Operating Systems

Candidates are asked to answer THREE questions from FOUR; each question is worth a total of 20 marks.

Use a separate answer book for each question.

An appendix for question 3 is included at the end of the paper


Question 1

1.a

- i. Define the distinction between concurrency and parallelism, and describe an example of a program which is concurrent but *not* parallel. [3 marks]
- ii. Provide *two* motivations for concurrency [2 marks]

1.b

- i. Provide *three* reasons why concurrent programmers are discouraged from relying on the scheduler to handle concurrency. [3 marks]
- ii. Define the two main aspects of concurrency management required to achieve synchronization, and provide examples for each within the context of computing. [4 marks]



1.c The above diagram shows a scenario where Master thread M creates n slave threads S_n to perform computation. This computation is a random number generator, returning values between 1 and 10, residing within a loop. M also sits within a loop, and once every t seconds will unblock and call three slave threads at random to return their current random value generated. Values returned by the three slave threads are then summed by the Master.

- i. In this scenario explain what are the potential risks with respect to synchronization. [2 marks]

Question 1 continues on next page...

Question 1 continued...

- ii. Write an implementation for this scenario: You are free to use either Java-based or C-based pseudo-code and any concurrent programming primitives as you see fit. A precise description of the semantics of the concurrent primitives should be provided, which argues how they ensure non-determinism. Code should be as precise and complete as possible.

[6 marks]

Total 20 marks

Question 2

2.a

- i. Define what is a kernel level thread and a user level thread. **[2 marks]**
- ii. Provide *two* advantages of user threads over kernel threads. **[2 marks]**
- iii. Provide *one* disadvantage of user threads compared to kernel threads. **[1 mark]**

2.b The following snippet of code provides multiple autonomous threads the ability to upload and download data into a database connected via a network. Locks are used to restrict access to the network and database from other threads during upload and download.

```
Object networkLock = new Object();
Object databaseLock = new Object();

public void SysDownload( ) {
    lock(networkLock);
    lock(databaseLock);

    downloadData( );

    unlock(networkLock);
    unlock(databaseLock);
}

public void sysUpload( ) {
    lock(databaseLock);
    lock (networkLock);

    uploadData( );

    unlock(databaseLock);
    unlock(networkLock);
}
```

- i. Explain the problems with this code, and why do they occur. **[3 marks]**
- ii. Rewrite the above code snippet to provide a thread-safe synchronized solution. **[2 marks]**

Question 2 continues on next page...

Question 2 continued...

2.c The code in 2.b.ii has now been redeveloped so that the system cannot upload data unless it has first been downloaded and modified. The data modification is performed using a `modify()` method. It is possible for the system to contain up to five pending download requests, with subsequent uploads resulting in data removal from the database.

- i. In this scenario semaphores represent one means to achieve synchronization; explain what the methods `wait()` and `signal()` represent. **[1 mark]**
- ii. Provide a pseudo-code implementation for this scenario, using semaphores to achieve synchronization. Semaphore objects should be represented with syntax similar to:

```
Semaphore s = new Semaphore(x);
```

Where x indicates the semaphore counter value. The semaphore method and its semantics should also be included, which argues how it ensures non-determinism.

Your code should be as precise and complete as possible. **[7 marks]**

- iii. Discuss whether blocking or spin locks would be more effective in the context of the described scenario, and explain why.

[2 marks]

Total 20 marks

Question 3

3.a Two alternative implementations of a function forming part of a memory allocation scheme are provided in an appendix at the end of this paper.

- i. Given the code for Algorithm A and Algorithm B in the appendix; explaining your reasoning clearly identify the memory allocation scheme each is implementing.

[4 marks]

- ii. If the conditional in Algorithm A is changed to the following, which memory allocation scheme would be implemented?

```
if ( free && ( foundBlock == NULL ) && ( size >= requestSize ) ) {
```

[2 marks]

3.b A Round-Robin scheduler is presented with the following processes

Process	1	2	3	4	5
Arrival Time (ms)	0	5	25	30	35
Burst Time	20	30	10	25	5

- i. Explain the terms:

- Burst Time
- Average Waiting Time
- Average Turnaround Time
- Scheduling Quantum

[4 marks]

- ii. Showing all working and being careful to justify your answer in terms of average and total waiting and turnaround times, would it be better to have a scheduling quantum of 10 or 15ms for this scenario?

[10 marks]

Total 20 marks

Question 4

4.a Modern multi-tasking operating systems must manage possibly many different processes and both fairly and safely share resources between them.

- i. Outline what we mean by a *cooperative scheduler* and a *pre-emptive scheduler* being careful to highlight the differences in terms of fairness and efficiency, and explain when and how each type of scheduler is invoked.

[8 marks]

- ii. What do we mean by a *Deferred Procedure Call* (DPC), what purpose do they serve, and why this type of mechanism is common in a modern operating system.

[3 marks]

4.b To test the efficiency of a resource allocation scheme you look at frame allocation over a short execution period. The current configuration allocates *four frames* to process under a *First In First Out* (FIFO) scheme and the process experiences *ten page faults*.

Given the page reference string { 6, 3, 4, 5, 6, 3, 8, 6, 3, 4, 5, 8 }, would there be any advantage in restricting the number of frames to *three* under either a FIFO or *Least Recently Used* (LRU) allocation scheme? You should assume that no pages are loaded at the start and, must show all working and justify your answer.

[9 marks]

Total 20 marks

Appendix: Code for question 3

Algorithm A

```
/*
 *   Algorithm A
 *
 *   Find block for request of size requestSize
 *   Returns pointer to most appropriate block or NULL if no suitable block found
 */
FreeBlock *
getBlock ( uint64_t requestSize ) {

    /*
     *   Record most appropriate block found so far
     *   and its size
     */
    FreeBlock * foundBlock = NULL;
    uint64_t    foundSize  = MAXBLOCKSIZE;

    /*
     *   Iterate through list of free memory blocks (FreeList)
     */
    FreeBlock * ptr = FreeList;

    while ( ptr != NULL ) {
        uint64_t size = ptr -> blockSize;

        if ( ( size < foundSize ) && ( size >= requestSize ) ) {
            foundBlock = ptr;    // Keep a record of the block we've just found
            foundSize  = size;    // ...and its size
        }

        ptr = ptr -> next;      // Move to next block in FreeList
    }

    return foundBlock;
}
```

Algorithm B appears on next page...

Code for question 3 continued, Algorithm B

```
/*
 *   Algorithm B
 *
 *   Find block for request of size requestSize
 *   Returns pointer to most appropriate block or NULL if no suitable block found
 */
FreeBlock *
getBlock ( uint64_t requestSize ) {

    /*
     *   Record most appropriate block found so far
     *   and its size
     */
    FreeBlock * foundBlock = NULL;
    uint64_t    foundSize  = 0;

    /*
     *   Iterate through list of free memory blocks (FreeList)
     */
    FreeBlock * ptr = FreeList;

    while ( ptr != NULL ) {
        uint64_t size = ptr -> blockSize;

        if ( ( size > foundSize ) && ( size >= requestSize ) ) {
            foundBlock = ptr;    // Keep a record of the block we've just found
            foundSize  = size;    // ...and its size
        }

        ptr = ptr -> next;      // Move to next block in FreeList
    }

    return foundBlock;
}
```

--- End of Paper ---